

23/10

Digital Images: Formation, Transformations and Filters

Giacomo Boracchi

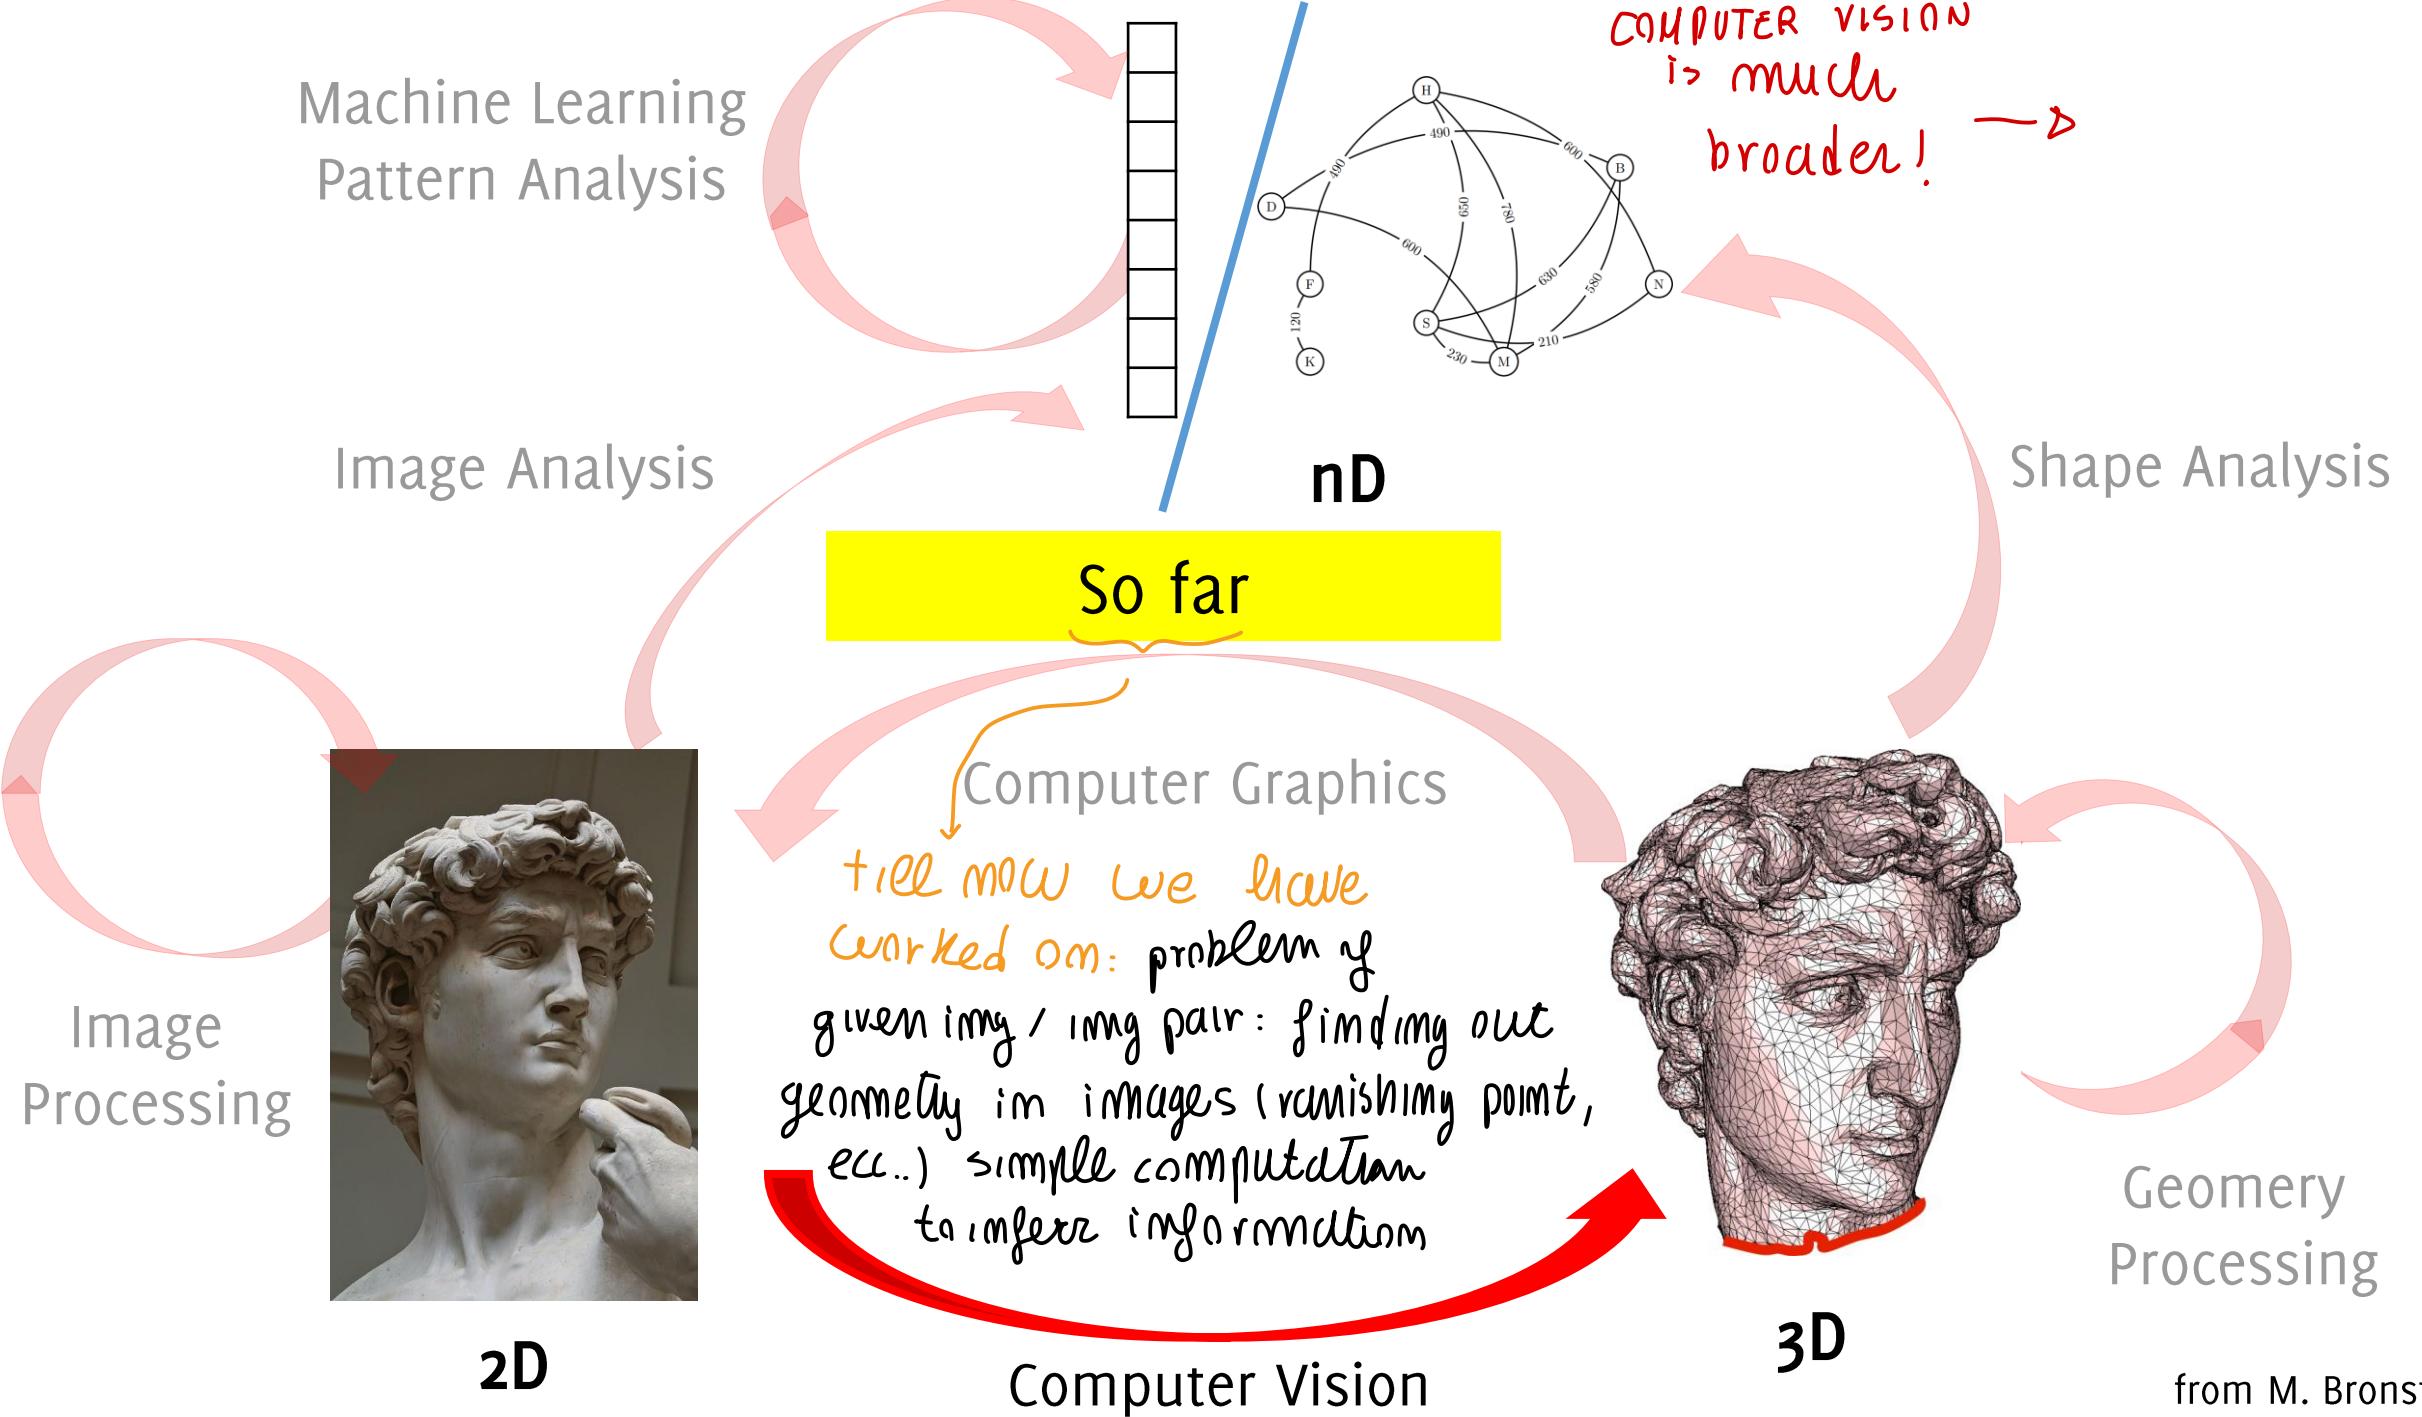
Image Analysis and Computer Vision

Politecnico di Milano

October 23rd 2024

Book: GW, chapter 3

The Broad Landscape of Computer Vision and Pattern Recognition



*all operations
that leave
within an
image... taking
an img as INPUT and OUTPUT
related to that image
itself*

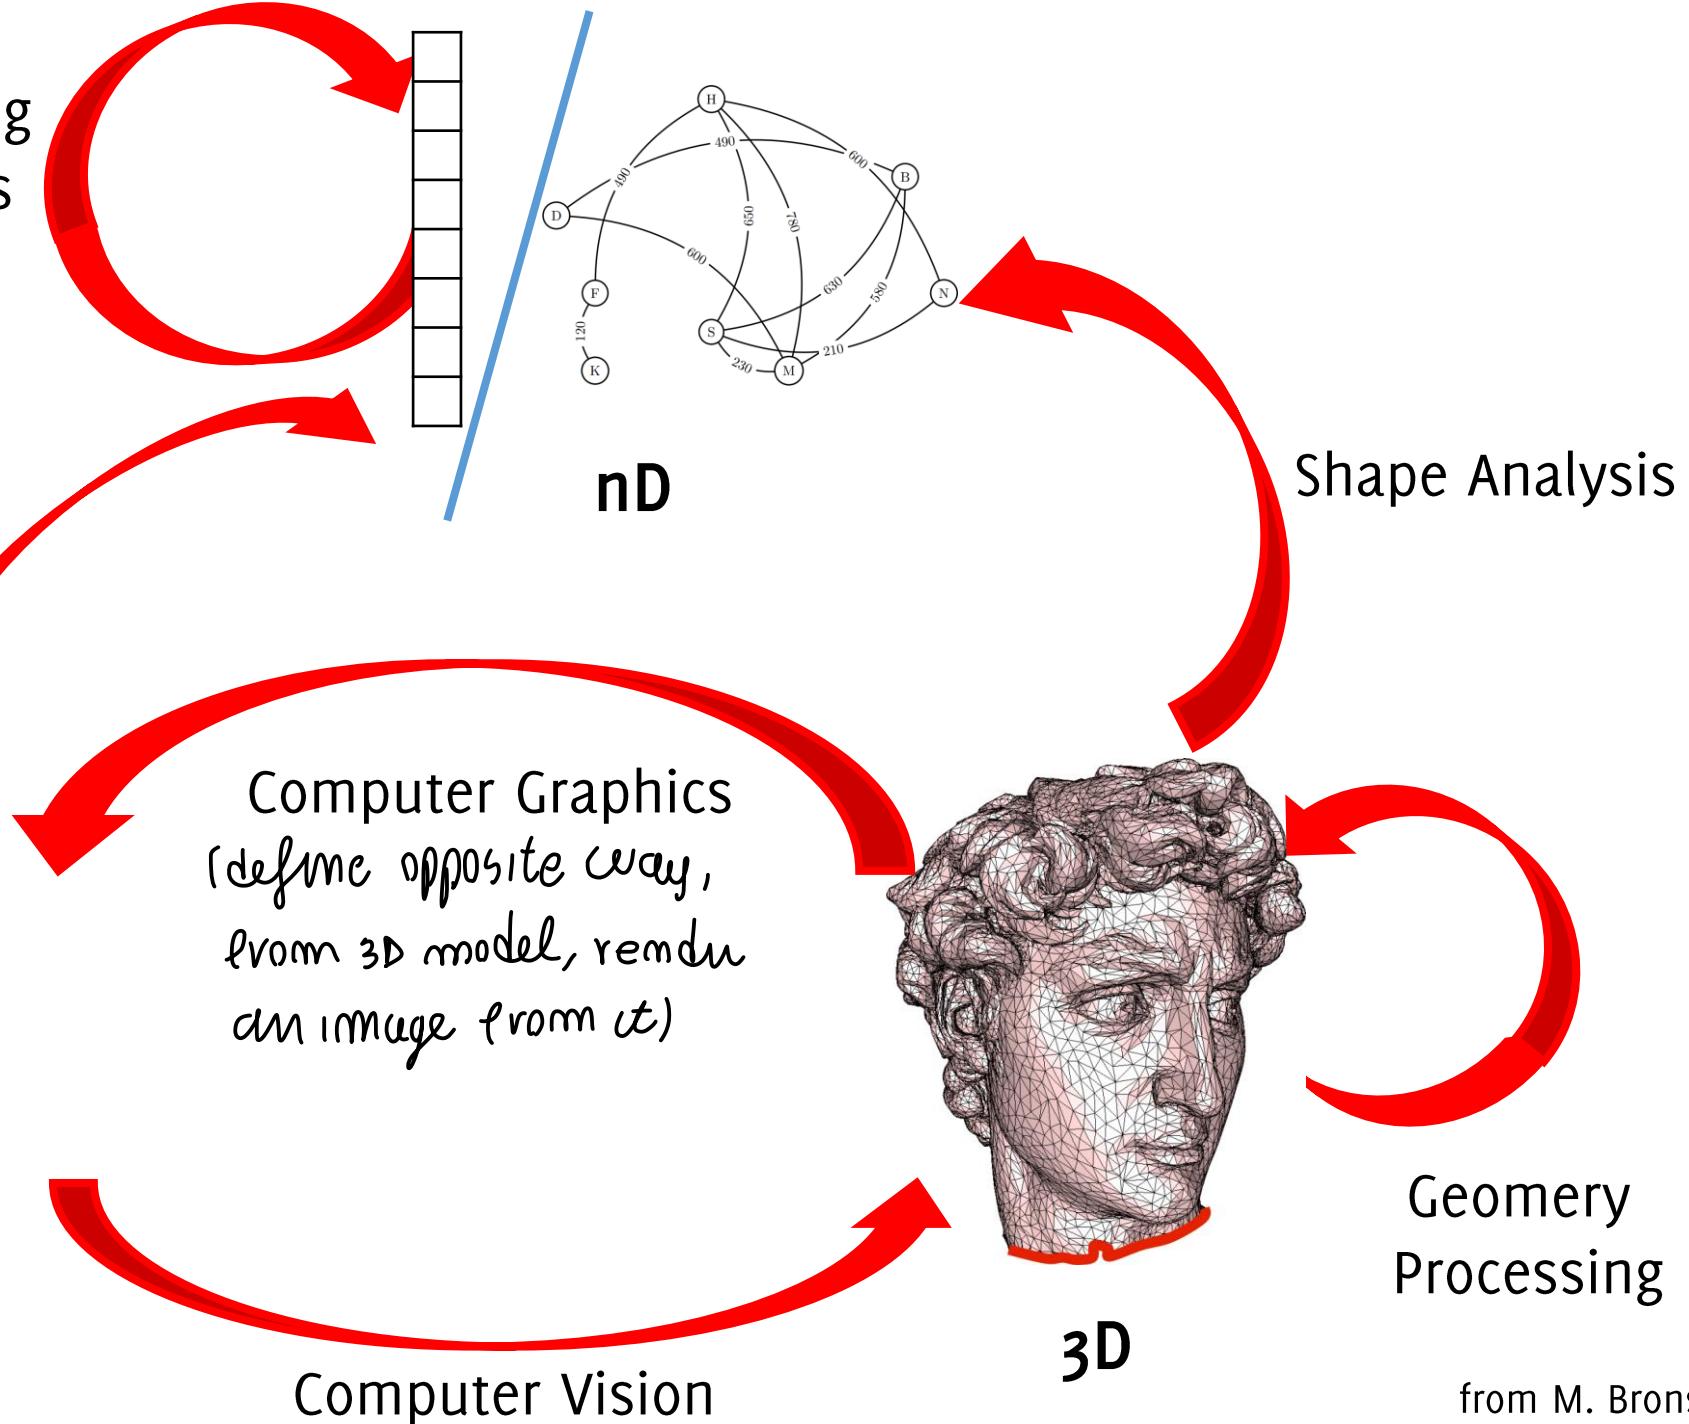
Machine Learning
Pattern Analysis

Image Analysis

Image Processing



2D



Machine Learning
Pattern Analysis

Image Analysis

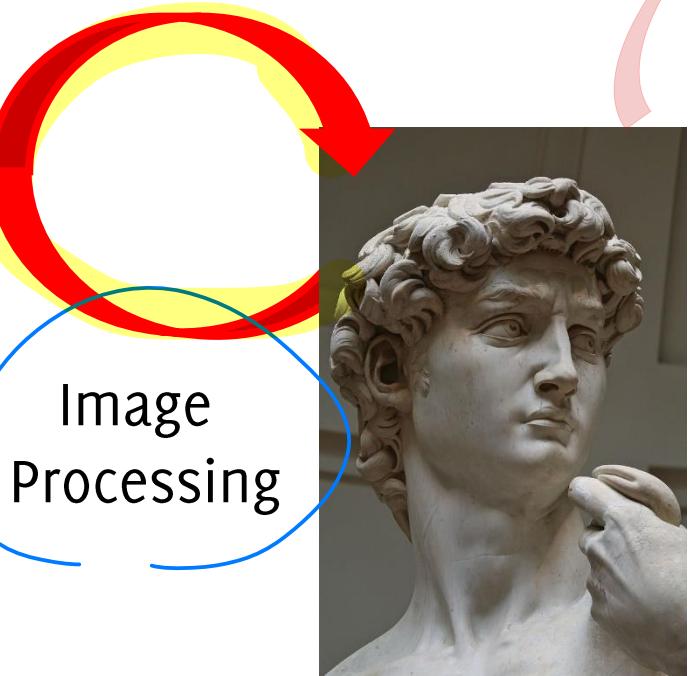


Image
Processing

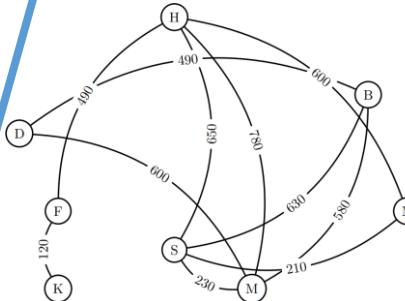
2D

For the next few lectures

Computer Graphics

Computer Vision

nD



3D

visual interpretation
of the image, moving from
data (3D mesh) to graph
analysis done for
semantic understanding
the image

Shape Analysis



Geometry
Processing

from M. Bronstein

In particular, we will see:

Restoration & Inverse Problems



Dabov, K., Foi, A., Katkovnik, V., & Egiazarian, K. Image denoising by sparse 3-D transform-domain collaborative filtering. *TIP* 2007

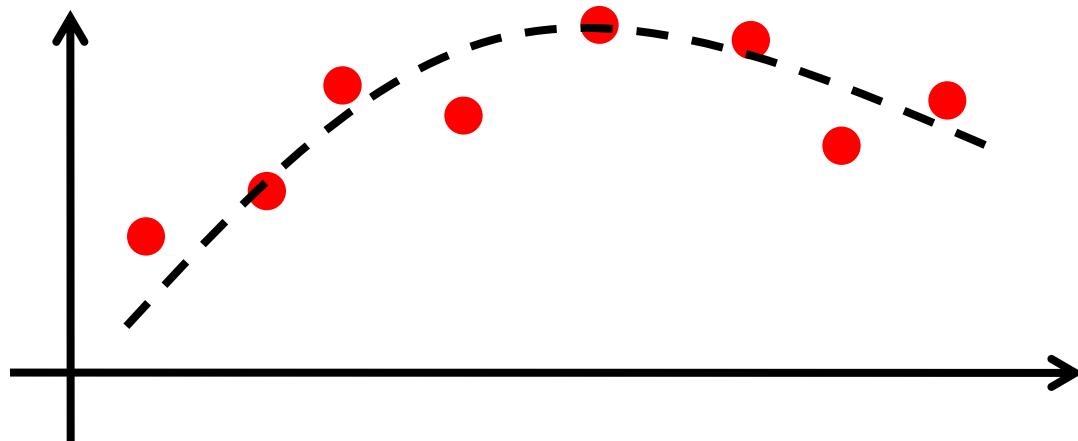
Restoration & Inverse Problems

↓
suppression
of noise in
images



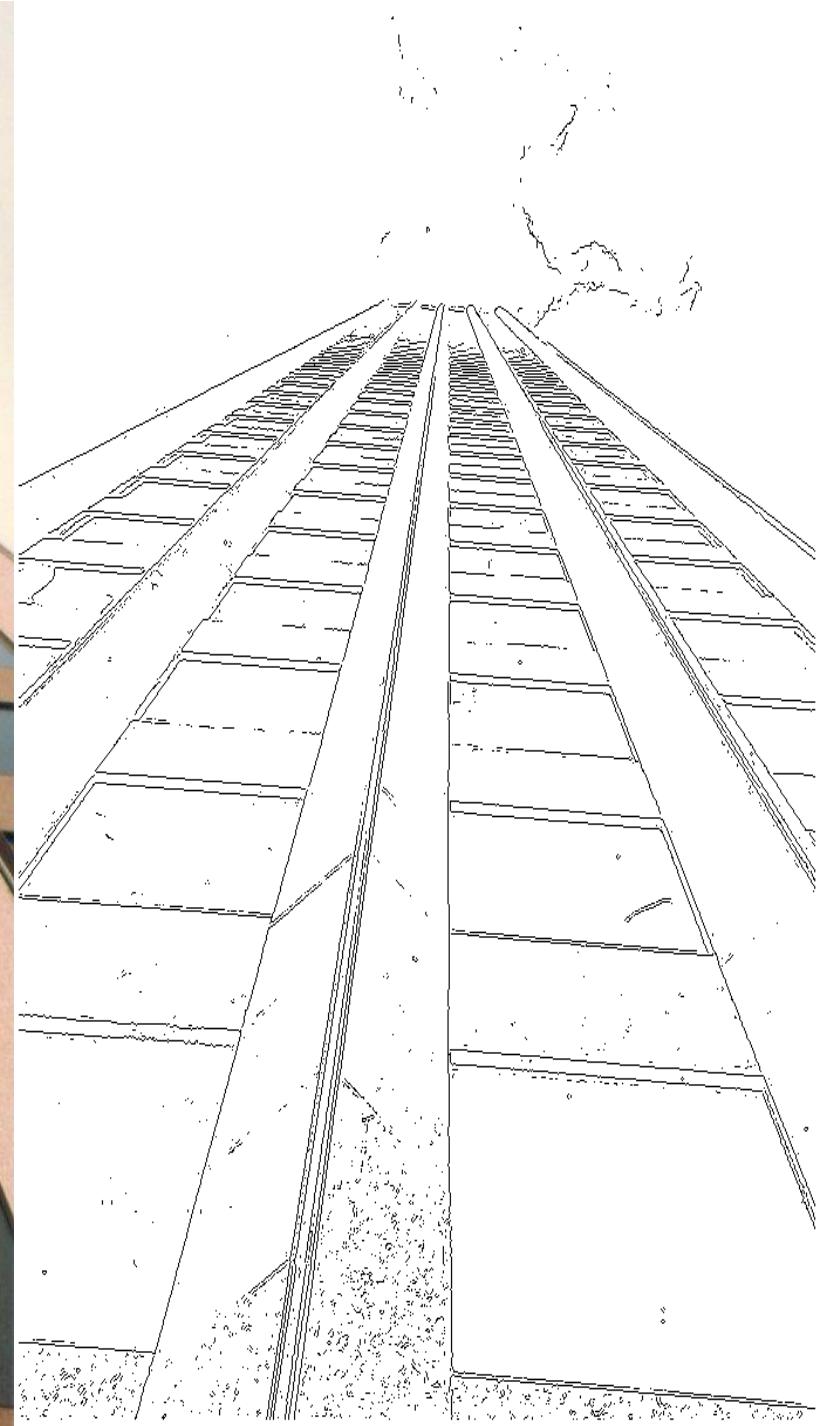
Filtering / denoising

each point with its neighborhood
is used to suppress noise, gathering points
and making estimate
~ estimate noise free data



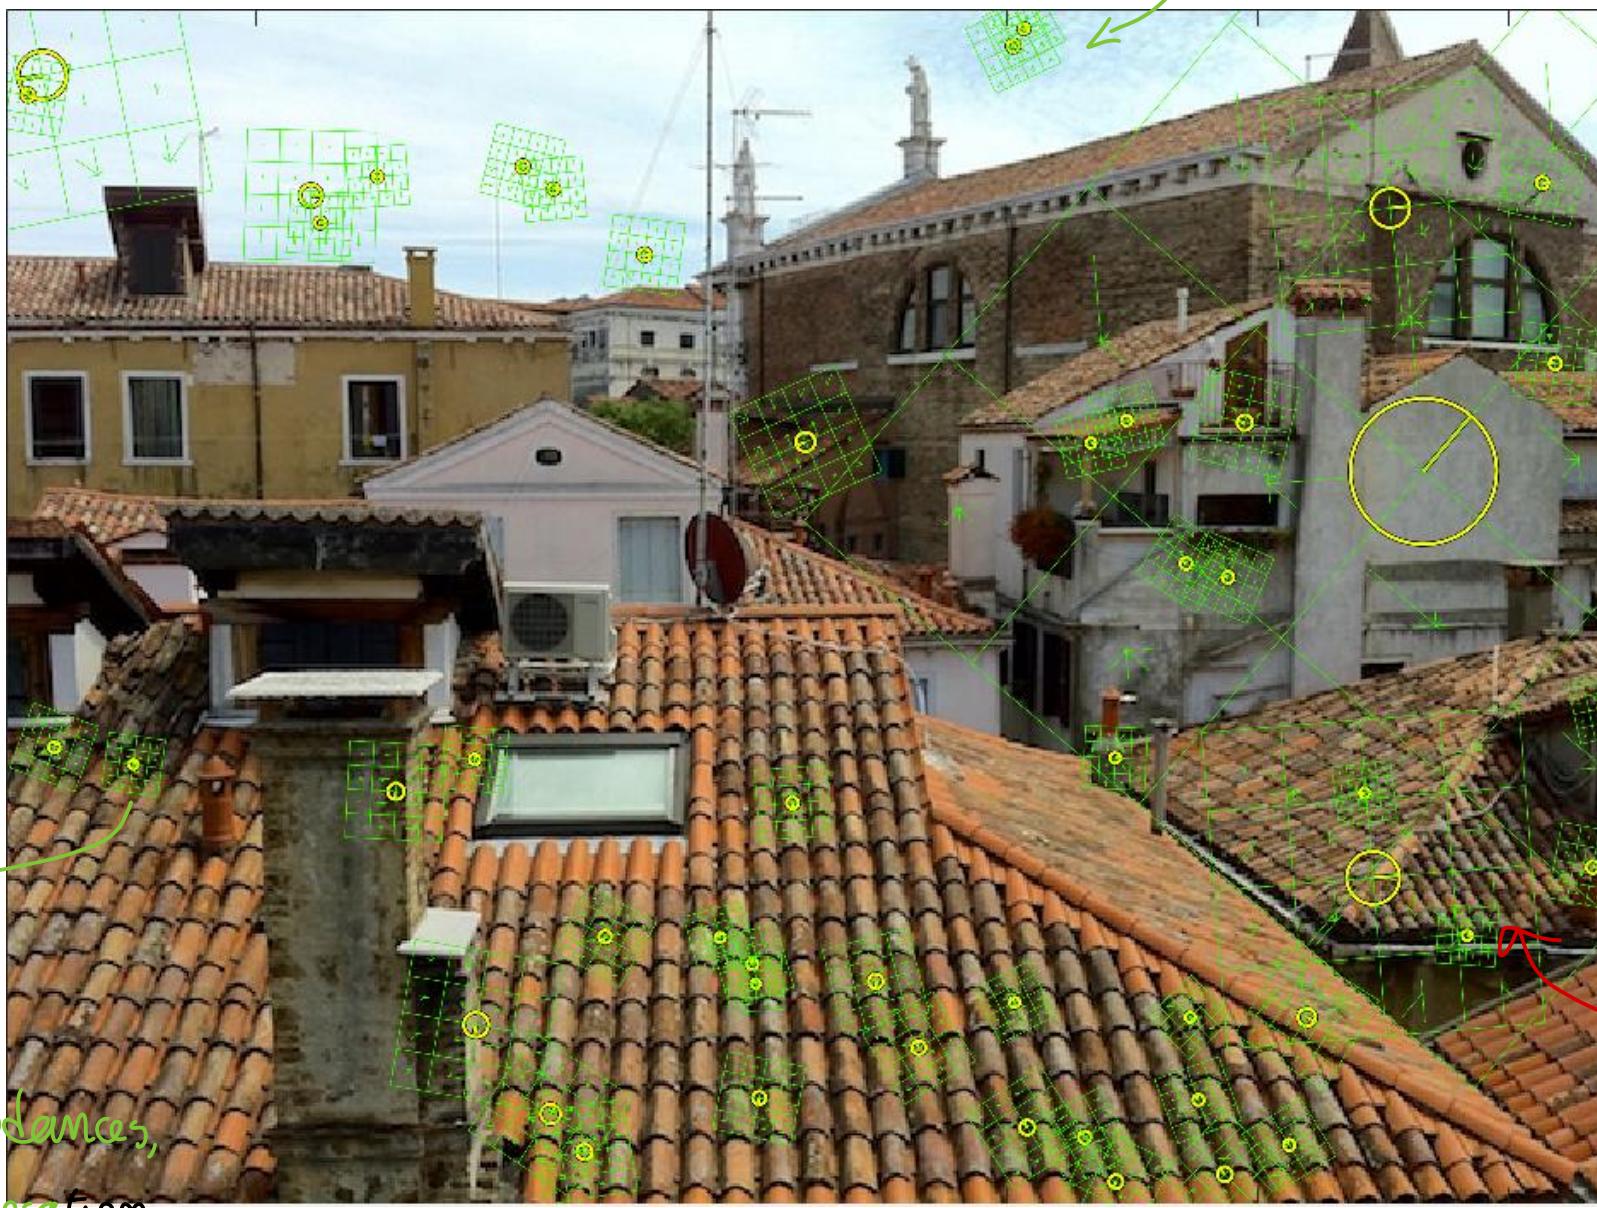
Edge detection

given img, how
to extract
boundary of the
img, retrieving a binary
image (which include
only boundary pixels)



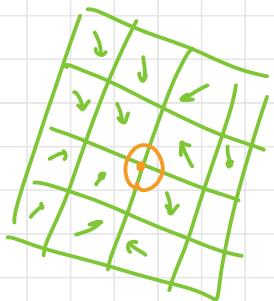
Salient Point / Feature Extraction

fundamental ingredient
for Reconstruction!



(*) h_x
candidate
points matched
on same img
on different
view! like
CIRCLES as
key point
for reconstruct
there are vectors
describing content
of image there,
that vector allow
to create correspondances,
to extract same location

↓
we assume
to establish
point - to - point
correspondence
between img
pairs! them
intersect
to establish
this correspond.
is ~~done~~
automatically
(*)



this arrow representation
allow to extract same
location between
two images of the same scene
from different view
↓
by similar desciption

Photometric Image Formation

↑
Let's discuss how
images are generated

Colour Filter Array

Colour Filter Arrays

[camera sensor model...]
↓ but how digital images are taken
behind the photosensor

Typical **photosensors** detect light intensity with little or no wavelength specificity, and therefore **cannot separate colour information**.

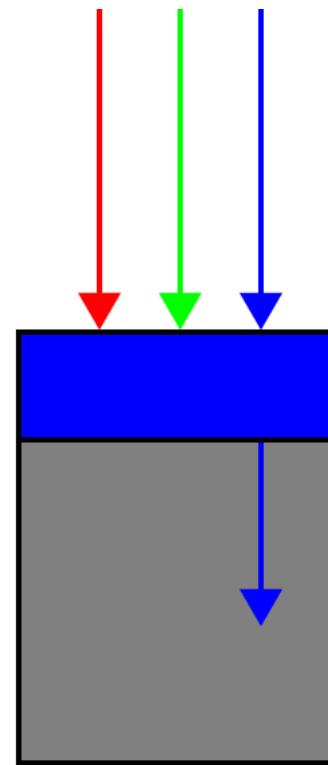
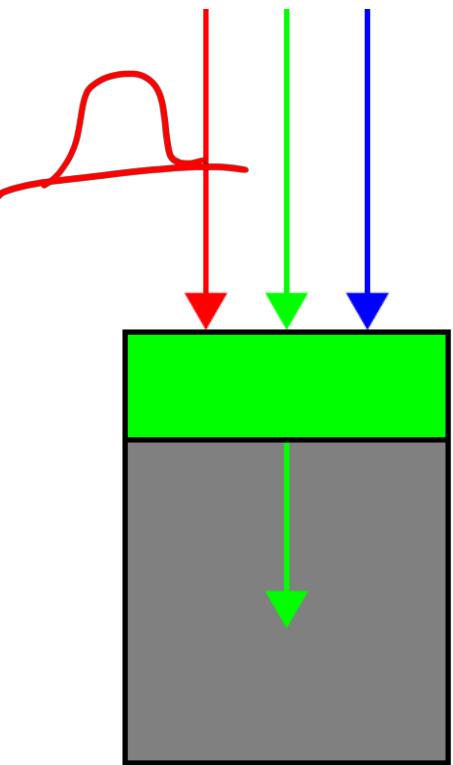
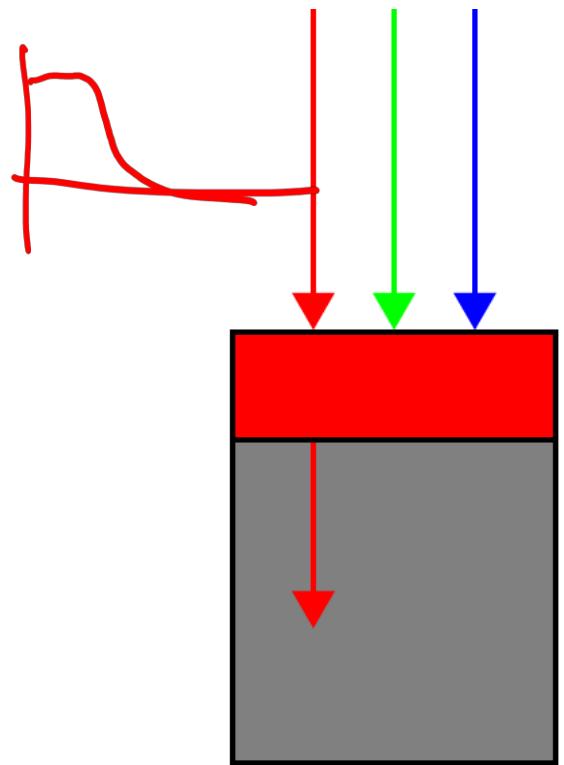
Colour Filters Array (CFA) are used to filter the light by wavelength range.

Separate filtered intensities include information about the colour of light.

For example, the Bayer filter gives information about the intensity of light in red, green, and blue (RGB) wavelength regions

Colour Filter Arrays

on top of each sensor (like CMOS) in a big ARRAY
made of photo detector

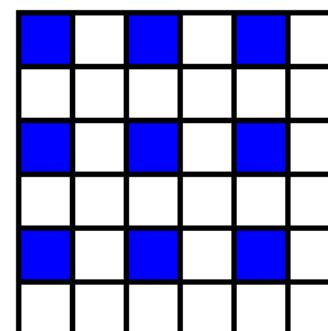
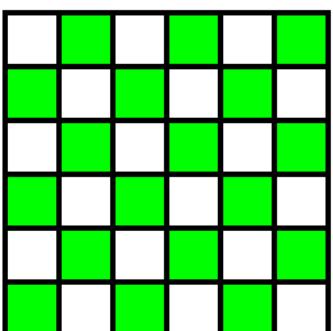
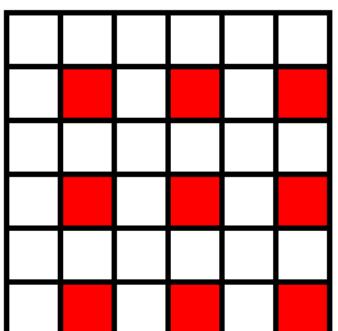


Incoming light

Filter layer

Sensor array

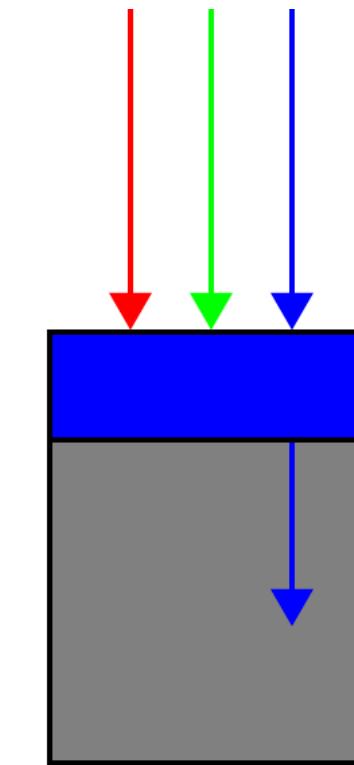
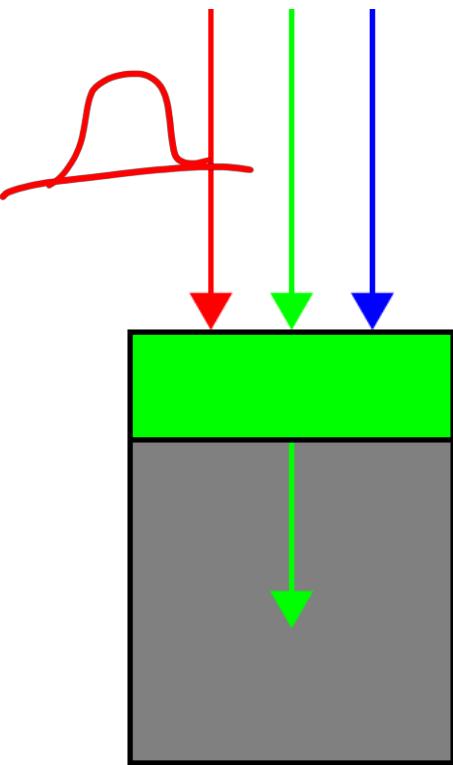
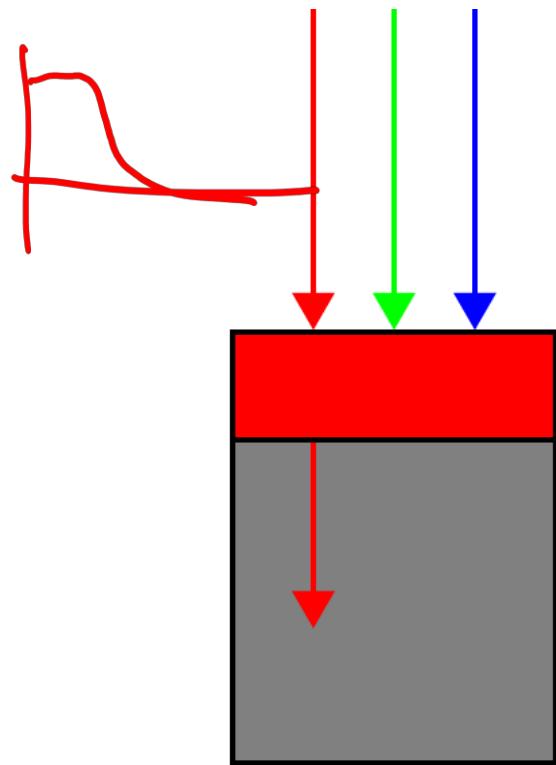
Resulting pattern



By en:User:Churnett - Own workThis W3C-unspecified vector image was created with Inkscape., CC BY-SA 3.0,
<https://commons.wikimedia.org/w/index.php?curid=1496872>

sensors
measuring
the amount
of light
brightness, NOT
color information
itself ... to retrieve
color information
RGB values, we
place on top of
each receptor
a COLOR FILTER
array

Colour Filter Arrays



NATURAL light comes with RGB information (cell)



Incoming light

color filter array takes

only a certain wavelengths,
different colors

"
different wave lengths,
let pass only specific part of color light

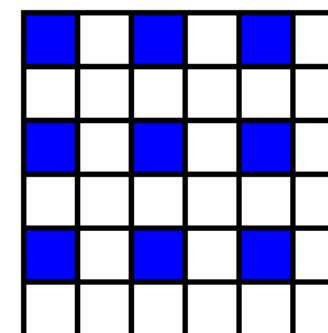
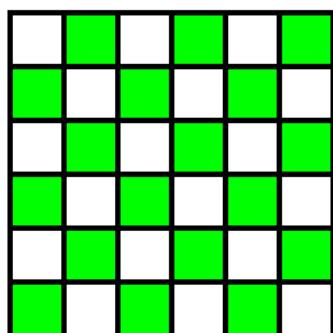
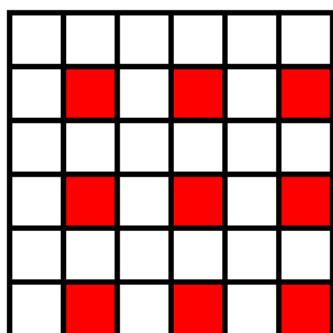
Filter layer

let pass only specific part of color light

Sensor array

In same pixel you have a single

value between R, G or B not
couple of colors



Resulting pattern

By en:User:Cburnett - Own work This W3C-unspecified vector image was created with Inkscape., CC BY-SA 3.0, <https://commons.wikimedia.org/w/index.php?curid=1496872>

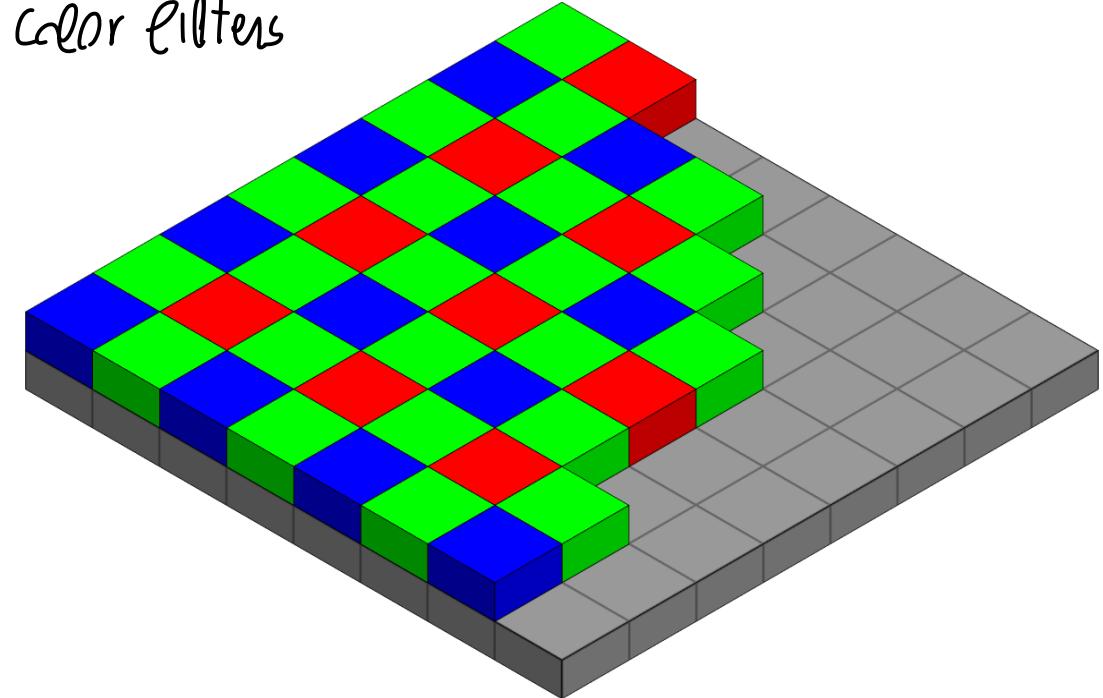
Bayer Pattern

standard layout for color filters

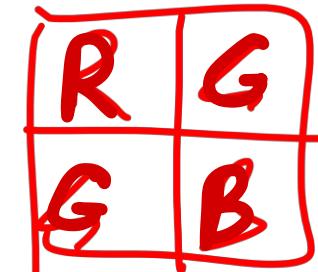
For example, the Bayer filter (RGGB) gives information about the intensity of light in red, green, and blue wavelength regions.

- Green colour is sampled twice

There are many different patterns, including RYYB which gives a better response in low-light conditions



By Cburnett - Own work, CC BY-SA 3.0,
<https://commons.wikimedia.org/w/index.php?curid=1496858>



measurements
correspond to

The raw output of digital camera

different colors

=>

(RGB depends on location)

Every pixel of the array is only sensitive to a single colour. This is what the sensor with CFA acquires ↓

brightness of that color

dark,
no intensity
of light

This is the scene

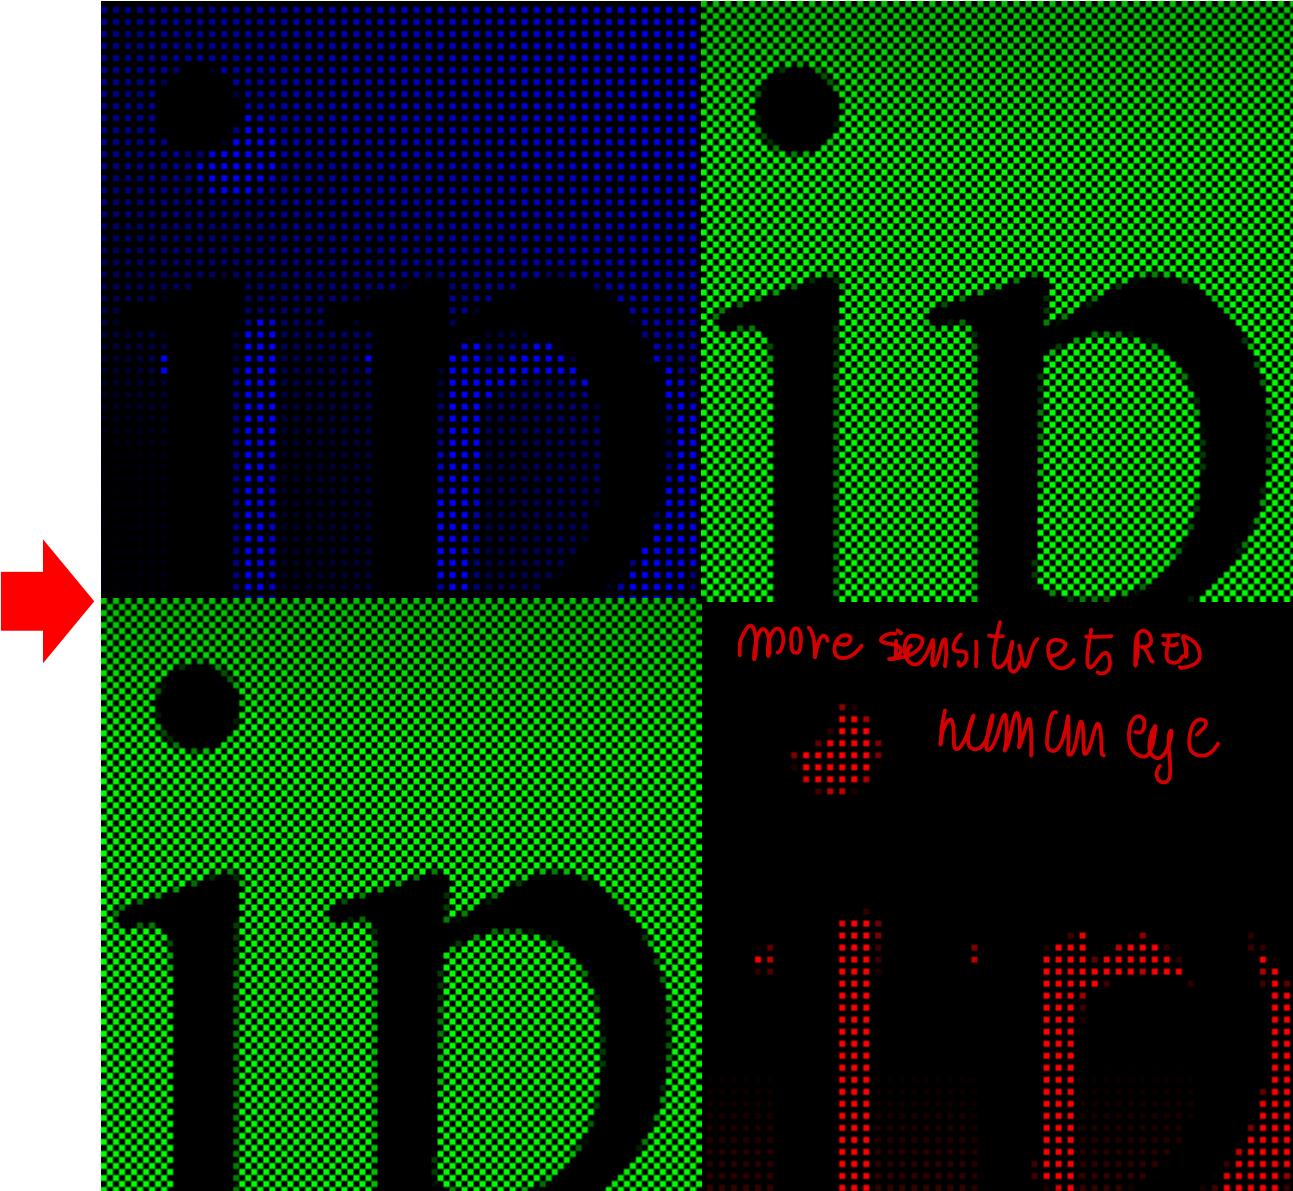


bright
area is
RGB values
spread
very SATURATED
= fully intensity



The raw output of digital camera

correspond to 4 images pixels



→ given that 4 images ... →
how to reconstruct
the image ?

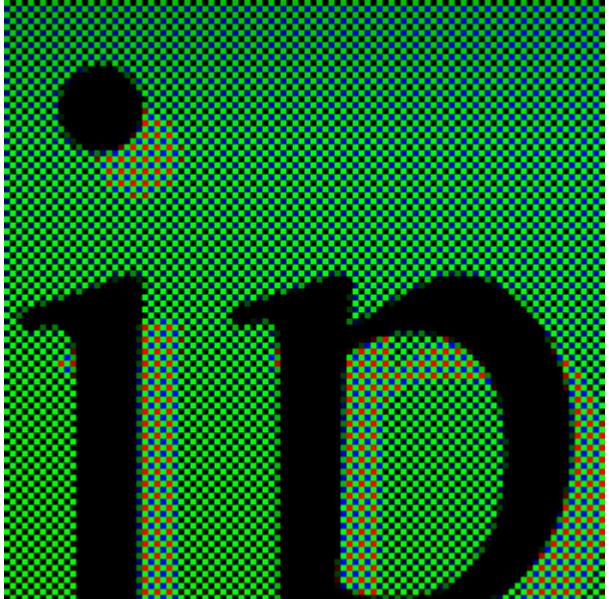
Demosaicing

From the mosaic image, you want to retrieve the image... "different colors" the interpolation

Demosaicing, a.k.a. CFA interpolation or Colour Reconstruction

Algorithm to reconstruct a full colour image (3 colours per pixel) from the incomplete colour output from an image sensor (CFA).

This is a **multivariate regression problem**

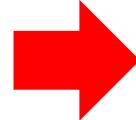


10MPx



90MPx
still here!

demosaicing



but RGB value
for each pixel



Demosaicing

can yield to a weird artifact.. bad quality
of image of Bad demosaicing

Issues:

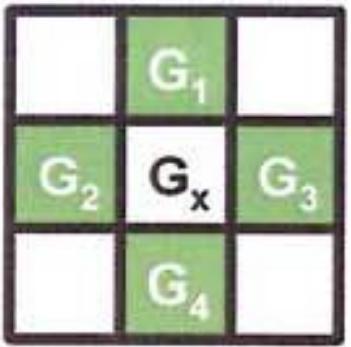
- In Bayer pattern each pixel is sensitive to a single colour, while in the image each pixel portrays a mixture of 3 primary colours

Desiderata: *What you want to achieve...*

- Avoid colour artefacts (you want to be efficient for using later)
- Maximum preservation of the image resolution
- Low complexity or efficient in-camera hardware implementation
- Amiability to analysis for accurate noise reduction

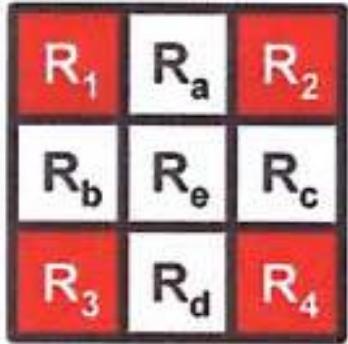
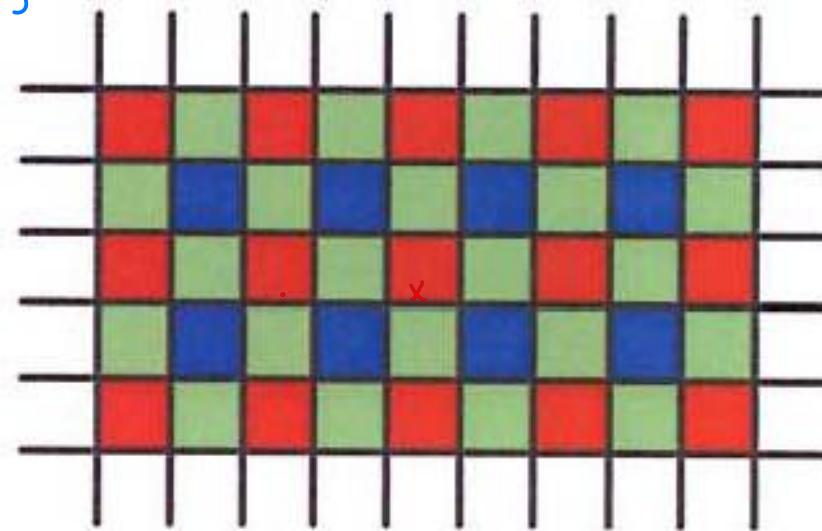
Example of Demosaicing by bilinear interpolation

if G missing
you have 4
pixels G
around...
you average
the ones
around



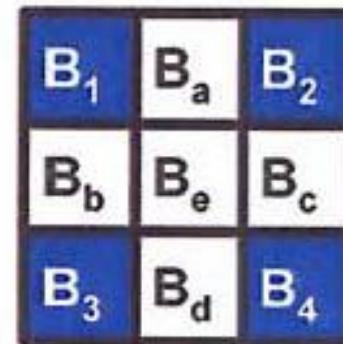
$$G_x = (G_1 + G_2 + G_3 + G_4) / 4$$

simple solution is
interpolating
missing pixel



$$R_e = (R_1 + R_2 + R_3 + R_4) / 4$$

$$\begin{aligned} R_a &= (R_1 + R_2) / 2 \\ R_b &= (R_1 + R_3) / 2 \\ R_c &= (R_2 + R_4) / 2 \\ R_d &= (R_3 + R_4) / 2 \end{aligned}$$



$$B_e = (B_1 + B_2 + B_3 + B_4) / 4$$

$$\begin{aligned} B_a &= (B_1 + B_2) / 2 \\ B_b &= (B_1 + B_3) / 2 \\ B_c &= (B_2 + B_4) / 2 \\ B_d &= (B_3 + B_4) / 2 \end{aligned}$$

More sophisticated channel-wise interpolation include bicubic/spline interpolation, Lanczos resampling

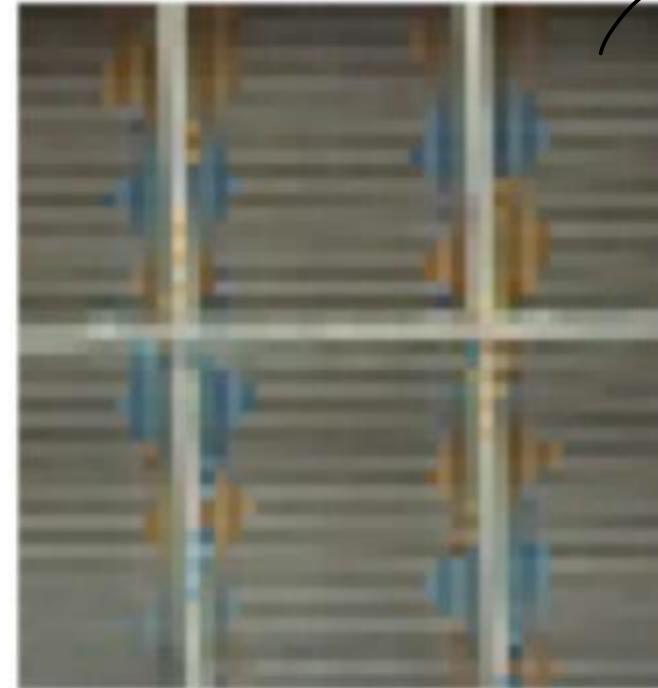
Demosaicing

with this bilinear regression, you get
bad results...

Color-independent algorithms typically present artifacts in regions containing edges and textures



Zipper effects are unnatural changes of intensities over a number of neighboring pixels, manifesting as an “on-off” pattern in regions around edges



False colors are spurious colors which are not present in the original image scene [...] They appear as sudden hue changes due to inconsistency among the three color planes and usually around fine image details and edges

False Colors

↳
this
happens when
higher contrast
occurs...

↳
Only on
RGB channels,
intensity is preserved
in primary difference
in R, G channel on
White/black pixel should
have some difference

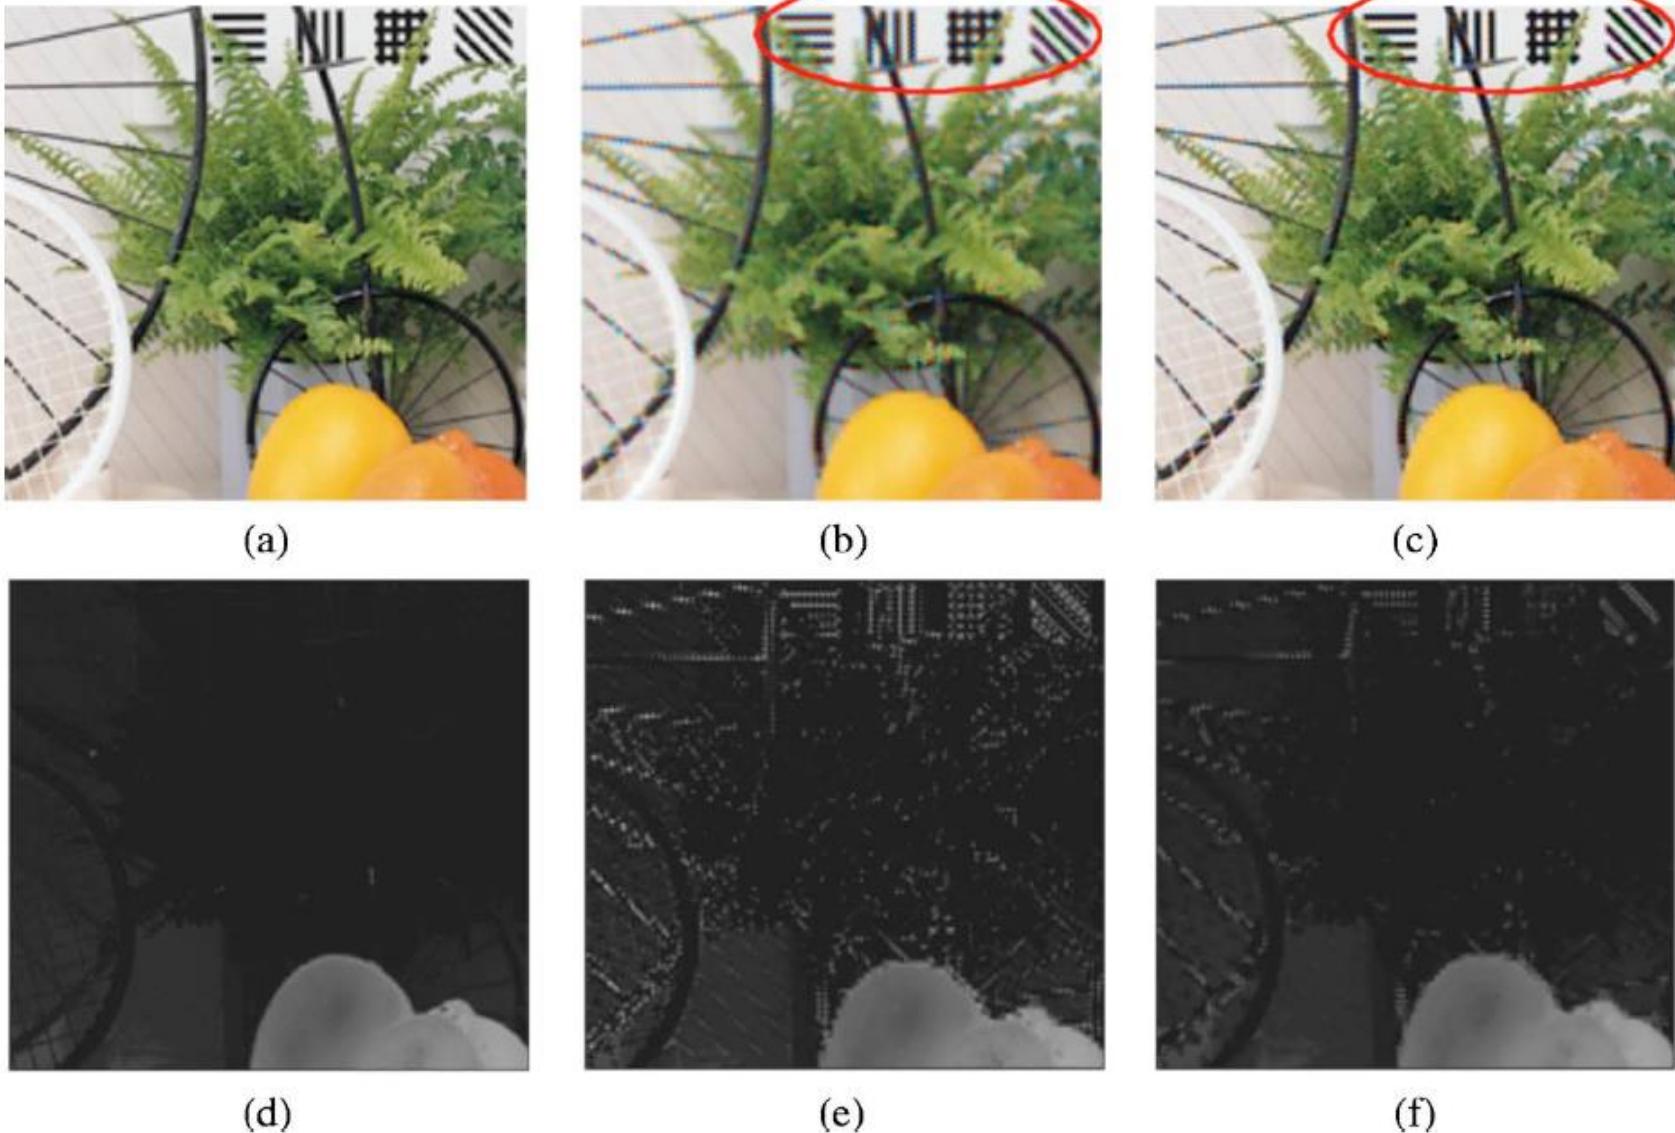


Fig. 10 Original image region (from image 27 in Fig. 15) and its demosaicked results obtained by (b) bilinear interpolation and (c) Freeman's method. The corresponding color difference planes (green minus red) are shown in (d), (e), and (f), respectively.

Demosaicing

Examples of priors to be exploited to improve demosaicing quality

- Channel-wise similarity / consistency (colour differences, colour ratio)
- Spatial correlation, the structure of images
- Spectral correlation

Post-processing can be employed to suppress typical demosaicing artifacts

← demosaicing
is hard, NOT
many KNOWN
Results, mainly
commercial
results.

When we work channel-wise...

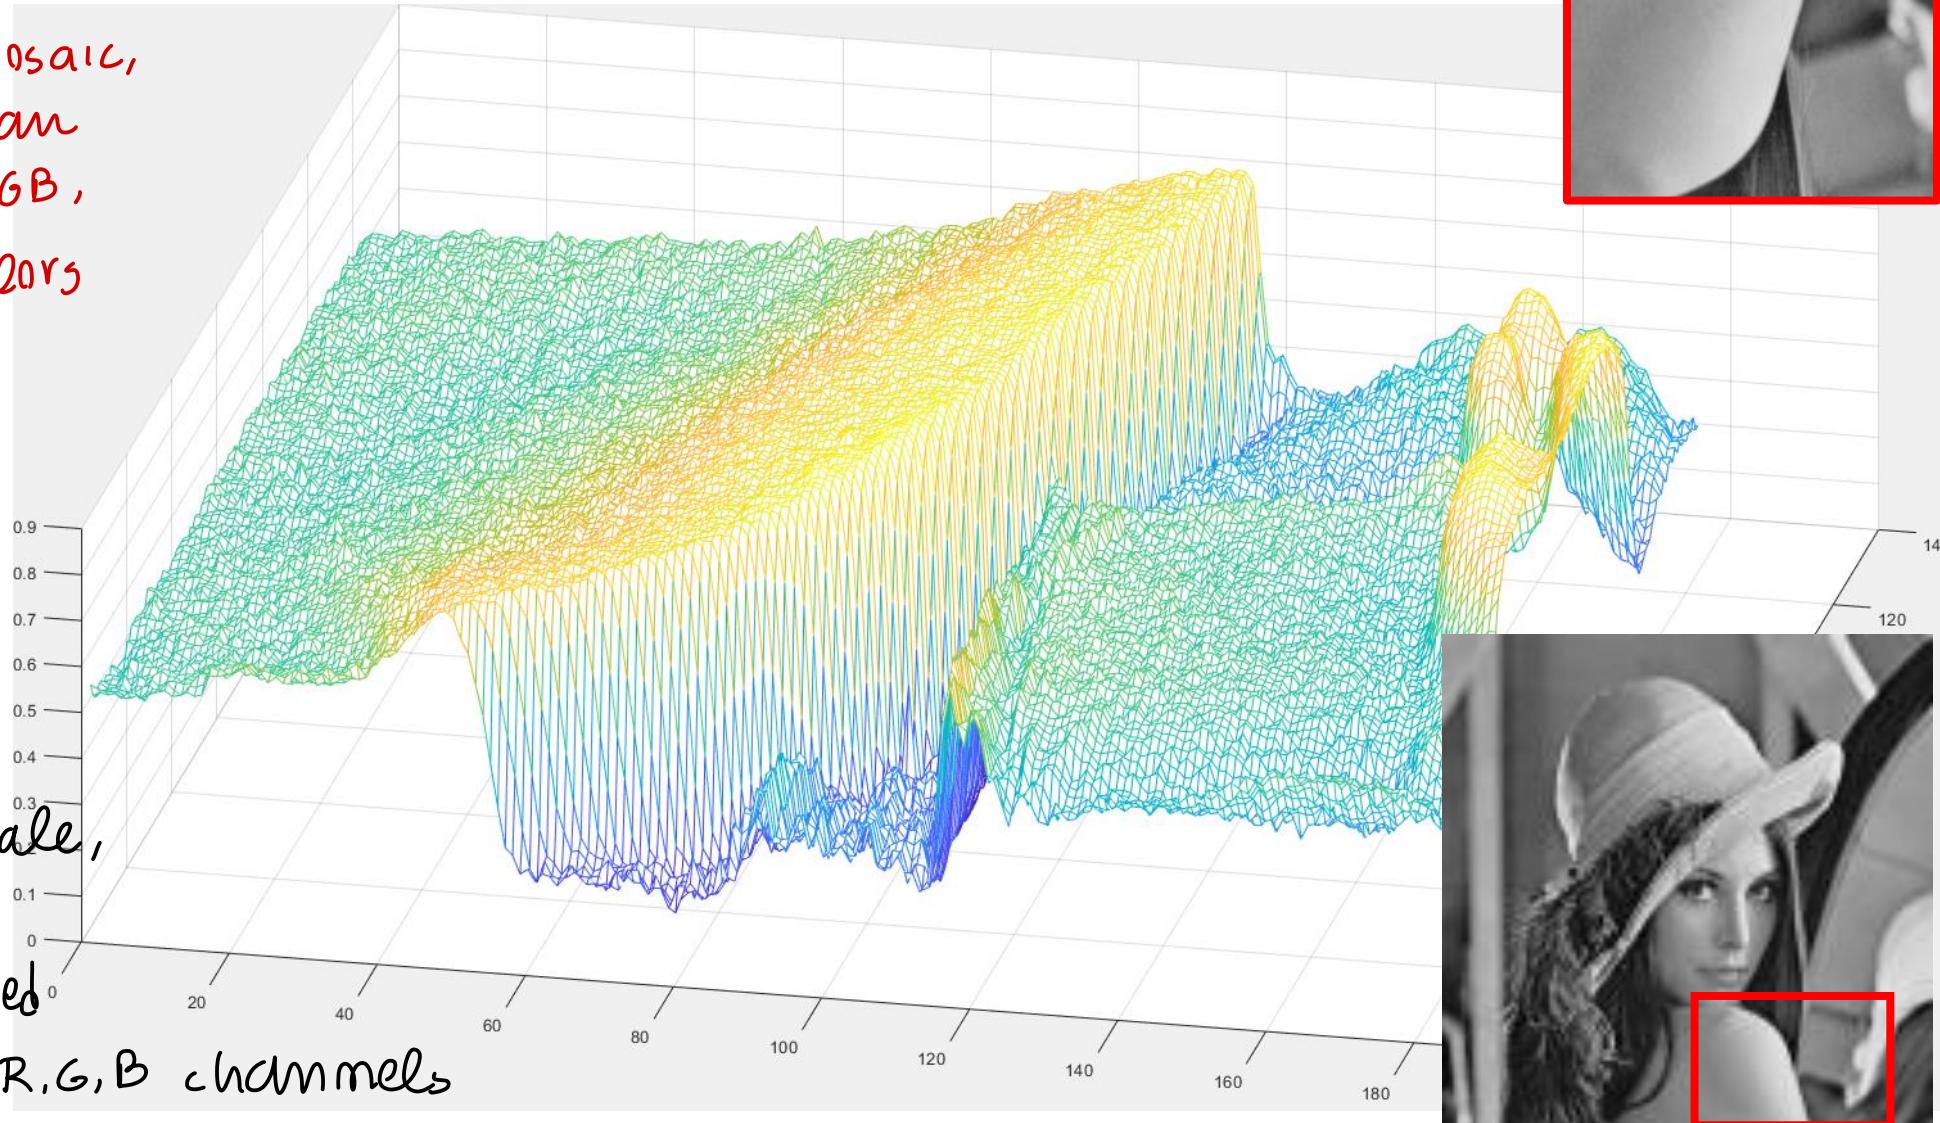
Think of an image as a 2D, real-valued function



after demosaic,
you have an
image RGB,
with 3 colors
per pixel!



for now we
ignore
color info,
We focus
on grayscale,
anything
is replicated
on the R,G,B channels



Think of an image as a 2D, real-valued function

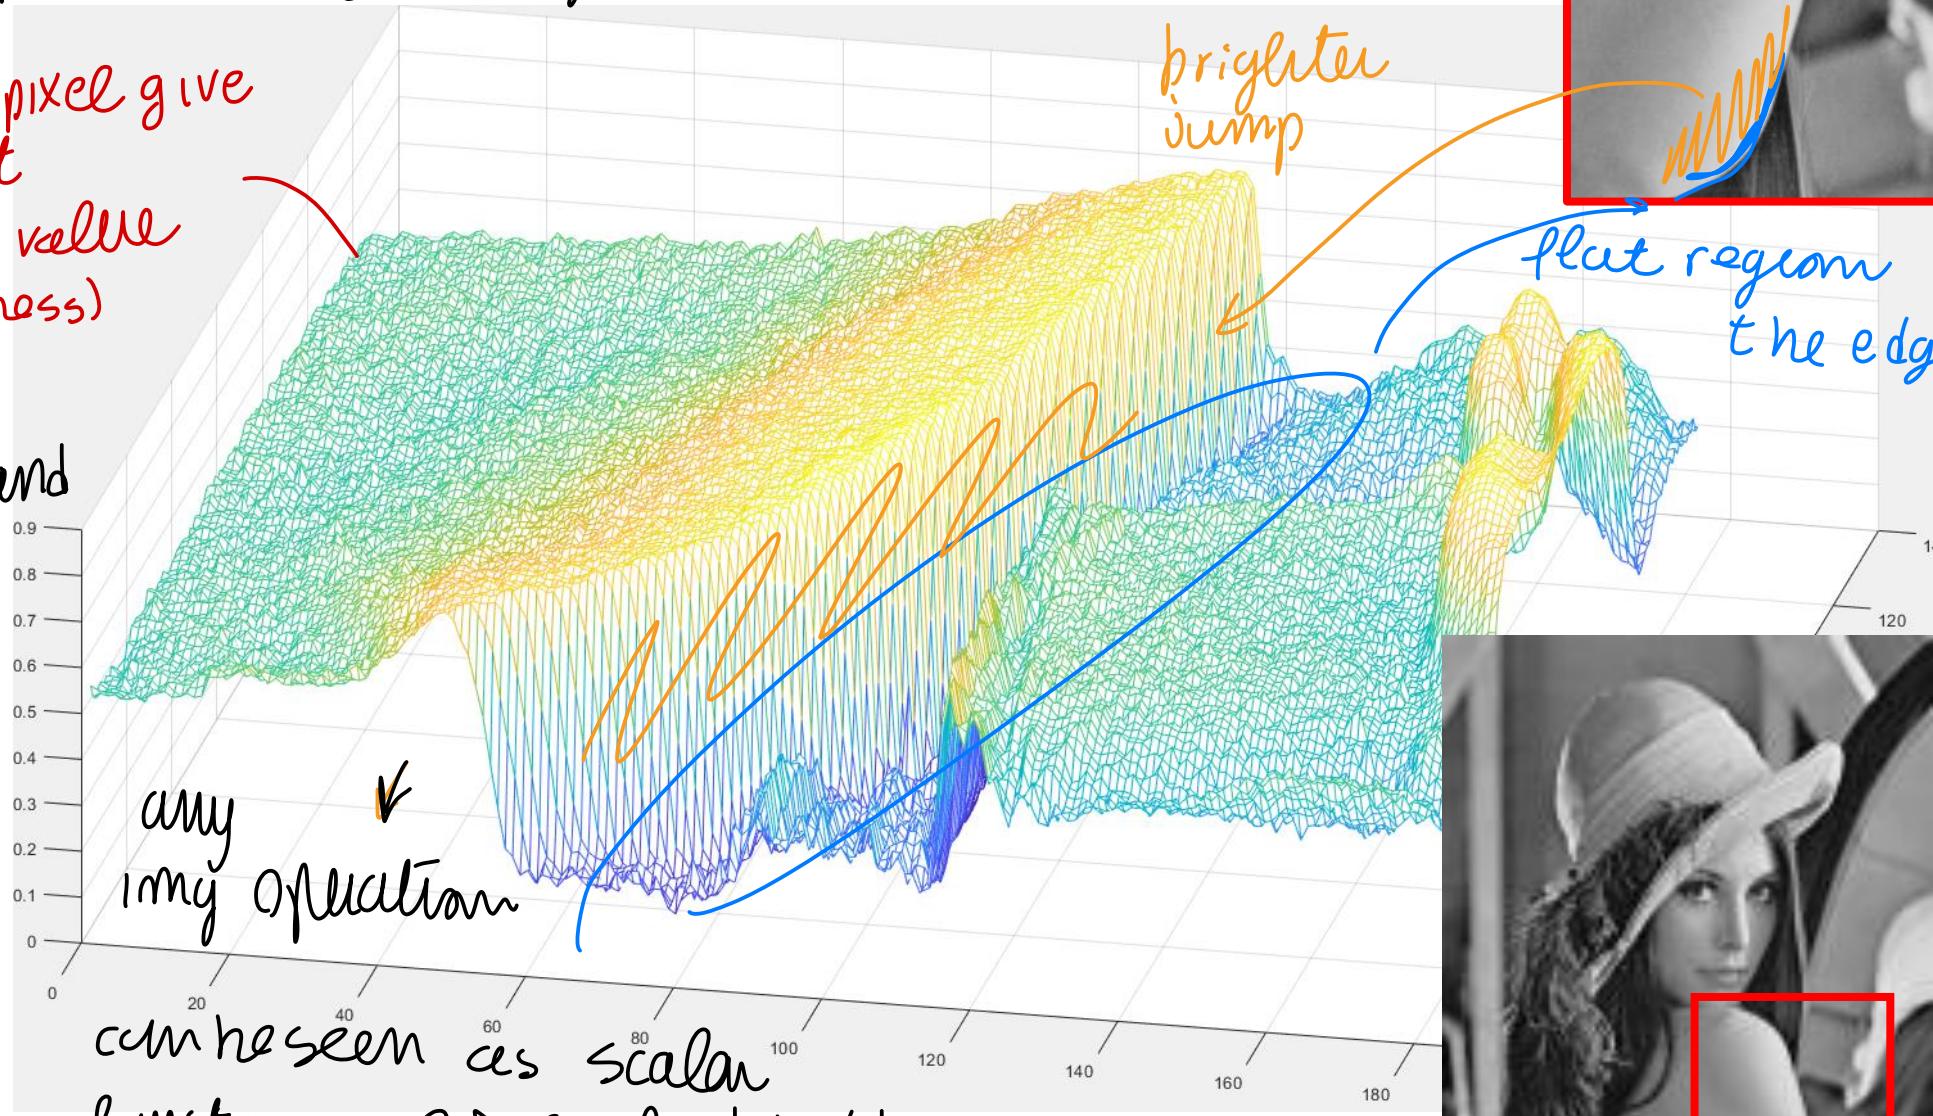
↓ we skip color RGB, go to Gray Scale ↑

each pixel give
the light
intensity value
(brightness)

image as a
2D surface, and
each pixel
has
intensity

any
img operation

can be seen as scalar
function in 2D applied to this...



Spatial-Domain Methods in Image Processing

we see 2 types of Transform:

A survey of most important operations in image processing:

- Spatial Intensity Transformations (change intensity of image)
- Spatial Local Transformations: convolution *useful for smoothing and differentiating*

Spatial transformations (intensity or local) are direct manipulation of pixel intensities. Relevant examples of convolutional filters:

- Smoothing Filters (denoising) *< simplest form of noise suppression*
- Differentiating Filters (edge detector) *< by convolution we can compute derivatives that give rise to edge detection*

Bibliography

“Digital Image Processing”, 4th Edition Rafael C. Gonzalez, Richard E. Woods, Pearson 2017

Intensity Transformations

Transformations that operate
on each single pixels of an image

Intensity Transformations

In general, these can be written as

$$G(r, c) = T[I(r, c)]$$

intensity transformation, take
single pixel as
input, give
single intensity

$r = \text{pixel row}$
 $c = \text{pixel column}$

Where

- I is the input image to be transformed, as a 2D array with scalar entries
- G is the output
- T is a function, for instance
 - $T: \mathbb{R}^3 \rightarrow \mathbb{R}$ (e.g. colour to grayscale conversion) (when original image is RGB)
 - $T: \mathbb{R}^3 \rightarrow \mathbb{R}^3$ (e.g. changing the colour encoding) colour conversion
 - $T: \mathbb{R} \rightarrow \mathbb{R}$ (many channel-wise intensity transformation)
↑ for each channel

T operates independently on each single pixel.

If the input image is RGB triplet, then

$$T: \mathbb{R}^3 \rightarrow \mathbb{R}$$

RGB → Grayscale Conversion

simply intensity information



A linear transformation of pixel intensities $T: \mathbb{R}^3 \rightarrow \mathbb{R}$

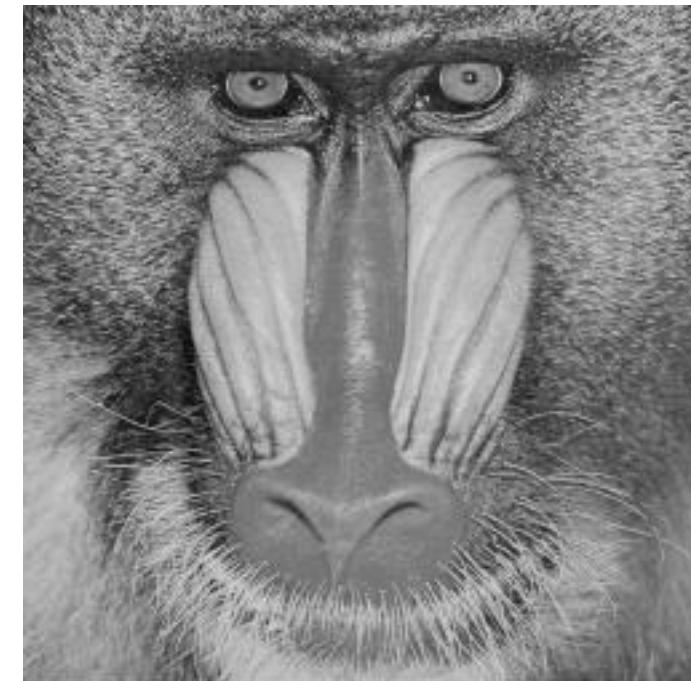
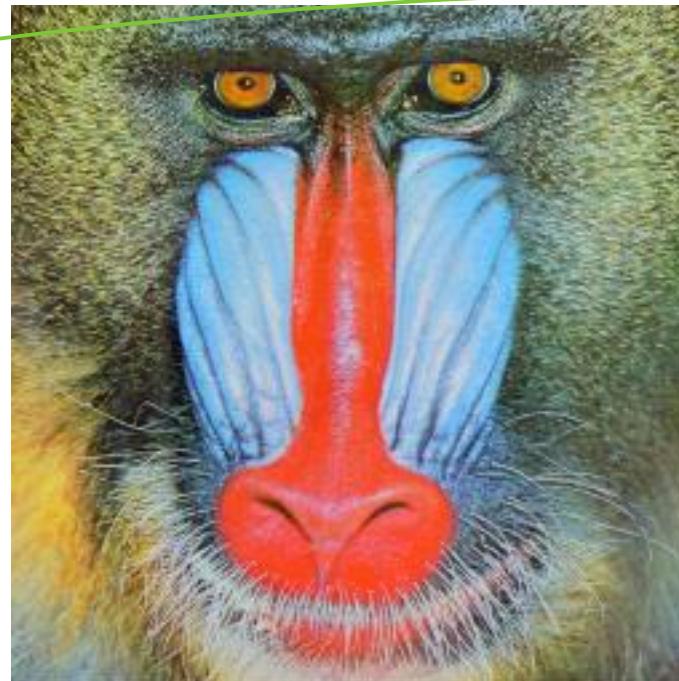
$$Gray(r, c) = [0.299, 0.587, 0.114] * [R(r, c), G(r, c), B(r, c)]'$$

which corresponds to a linear combination of the 3 channels

weighted average

$$Gray(r, c) = 0.299 * R(r, c) + 0.587 * G(r, c) + 0.114 * B(r, c)$$

more G than B, as standard Grayscale compression, because of human sensitivity



standard in any SW for computer vision

YCbCr color space



Color space conversion $T: \mathbb{R}^3 \rightarrow \mathbb{R}^3$ to map RGB to YCbCr

- Y is the luma signal, similar to grayscale
- Cb and Cr are the chroma components

red amount

blue amount

Luminance = how bright the pixel is

Human eye is less sensitive to color changes than luminance variations.

Thus, *very popular in my processing → being human eye is more sensitive in light intensity change rather than color*

- Y can be stored / transmitted at high resolution
- Cb and Cr can be subsampled, compressed, or otherwise treated separately for improved system efficiency

higher compression

in Cb,Cr,

because eye less sensitive

(e.g. in JPEG compression the chromatic components are encoded at a coarser level than luminance)

*+ to reduce size, handle 3 channels differently,
in YCbCr you compress less Y to avoid artifacts*

RGB → YCbCr this is a LINEAR transformation, very popular

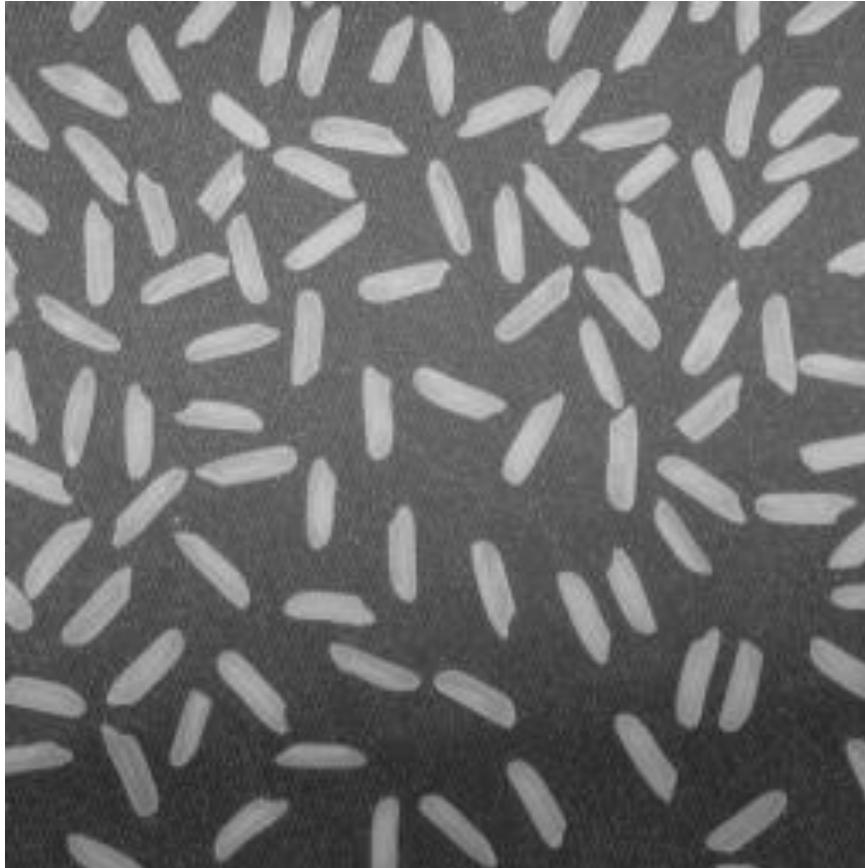
There are many variants \rightarrow this case is a matrix multiplication with proper configuration

$$\begin{bmatrix} Y' \\ P_B \\ P_R \end{bmatrix} = \begin{bmatrix} K_R & K_G & K_B \\ -\frac{1}{2} \cdot \frac{K_R}{1-K_B} & -\frac{1}{2} \cdot \frac{K_G}{1-K_B} & \frac{1}{2} \\ \frac{1}{2} & -\frac{1}{2} \cdot \frac{K_G}{1-K_R} & -\frac{1}{2} \cdot \frac{K_B}{1-K_R} \end{bmatrix} \begin{bmatrix} R' \\ G' \\ B' \end{bmatrix}$$

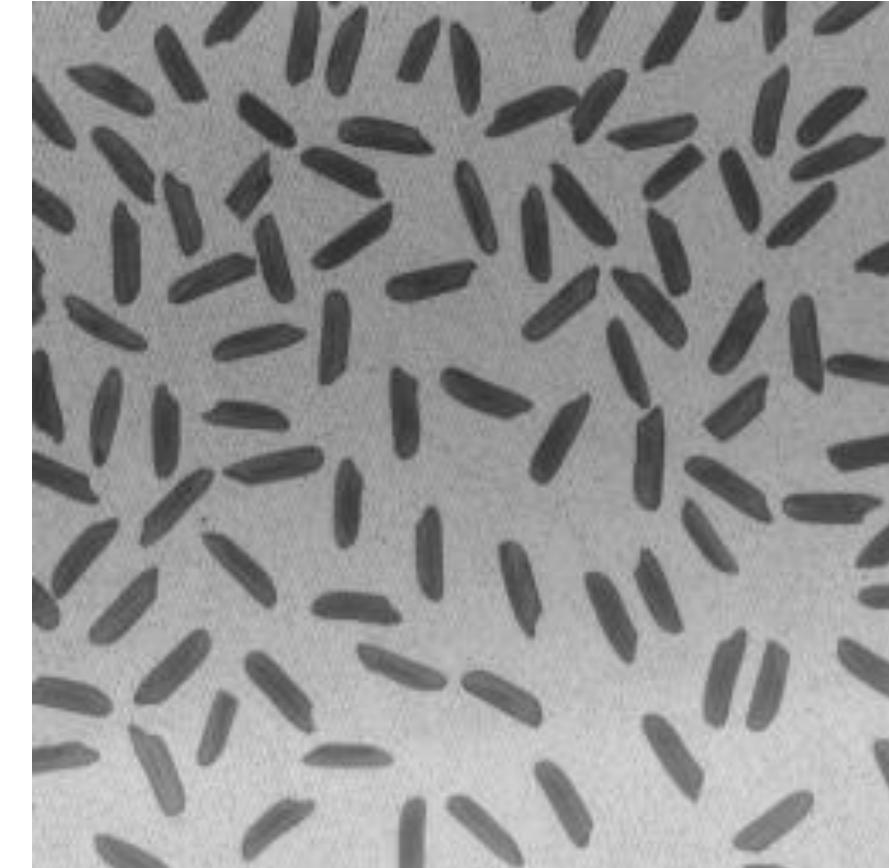
Where ' denotes the intensities are in the [0,1] range and where K_R , K_G , and K_B are ordinarily derived from the definition of the corresponding RGB space, and required to satisfy $K_R + K_G + K_B = 1$

Example of Intensity Transformation

original

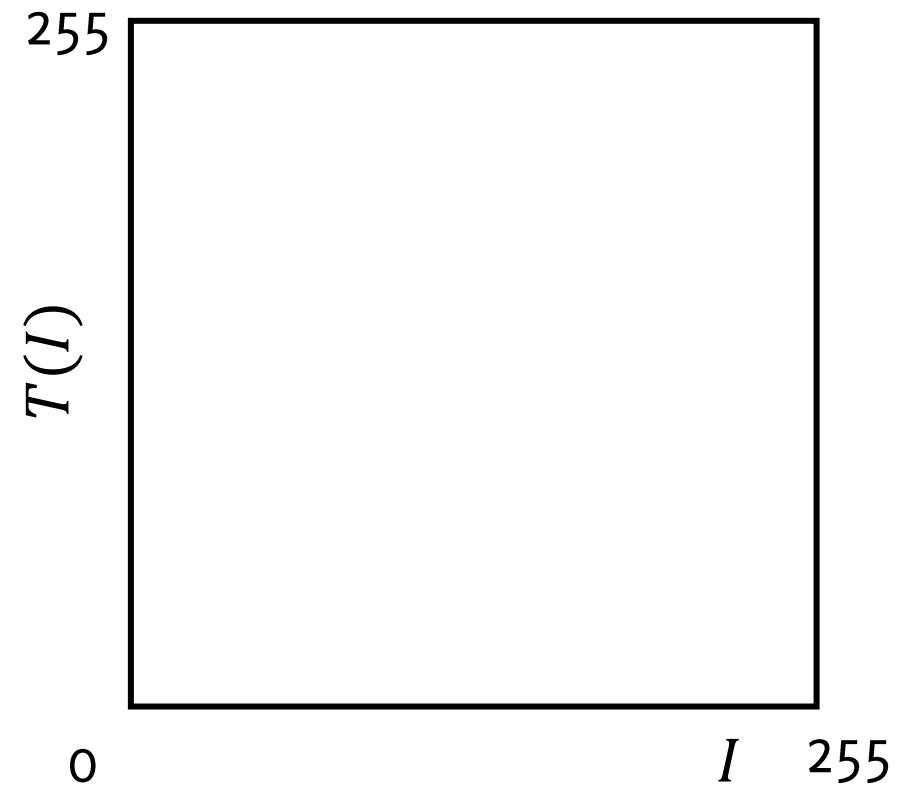


transformed



Negative Transformation $\mathbb{R} \rightarrow \mathbb{R}$

Simple transformation that maps black to white and white to black, and accordingly all the intensity levels in between



Negative Transformation

flip color bright \rightarrow dark
dark \rightarrow bright

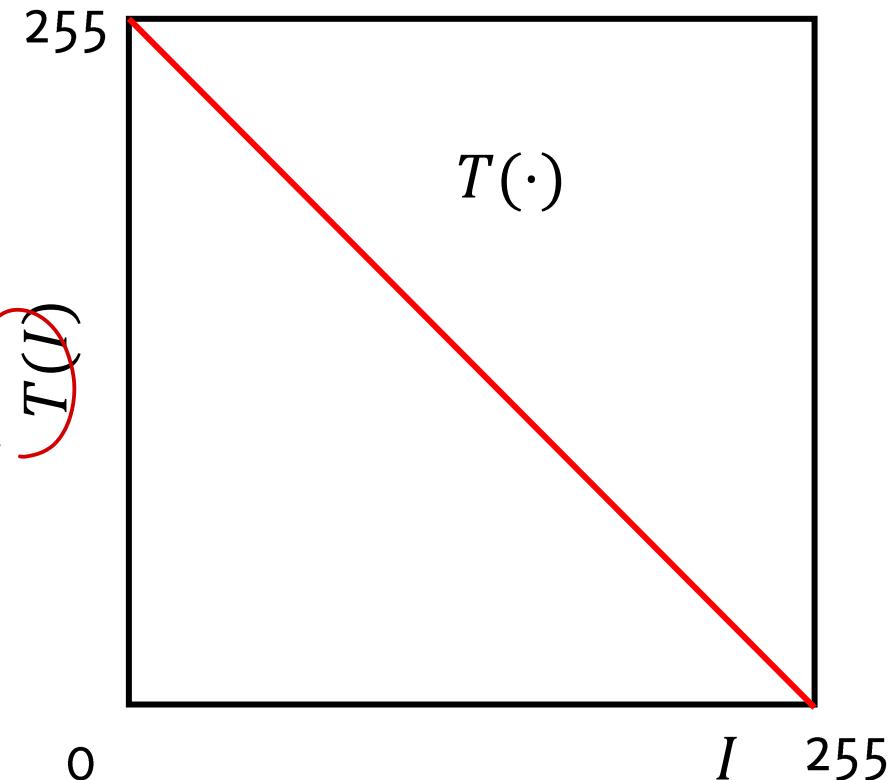
Negative Transformation in equations

$$I(r, c) \rightarrow 255 - I(r, c)$$

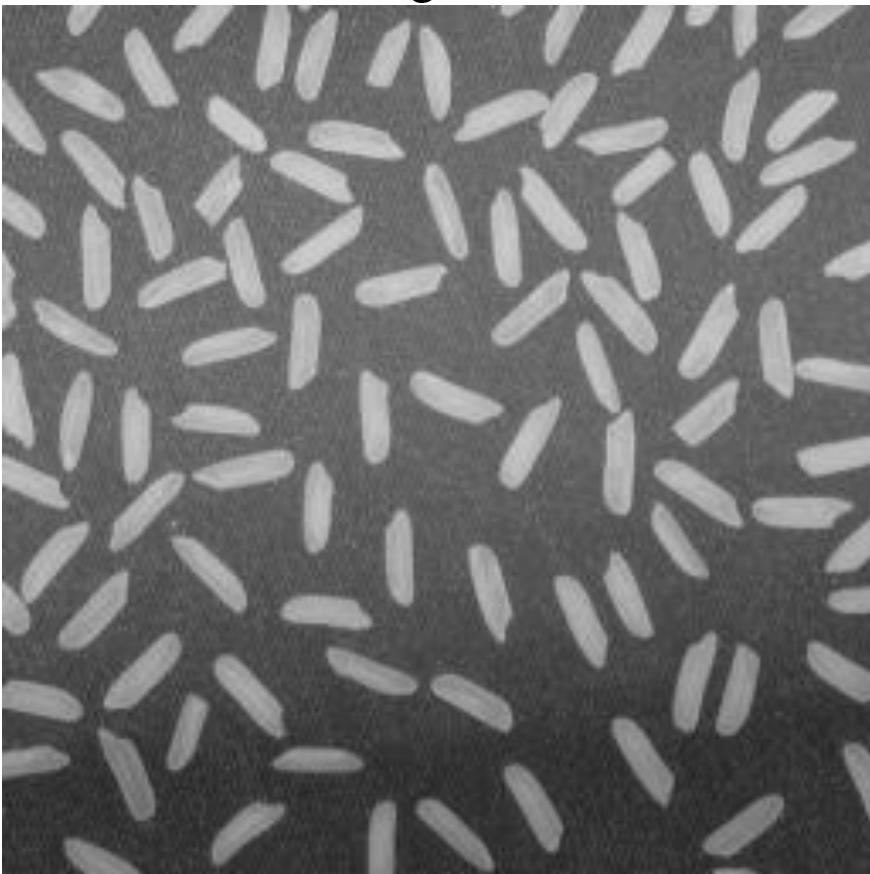
\downarrow brings intensity to opposite!

This is a linear transformation of intensities

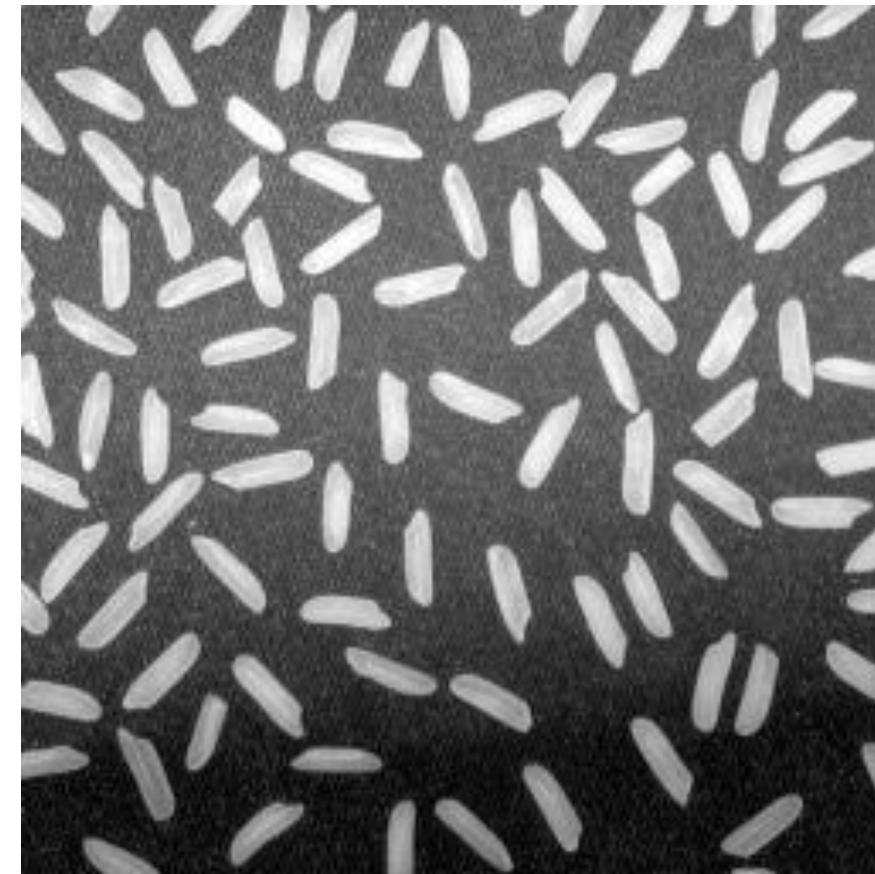
output on each pixel



original



transformed



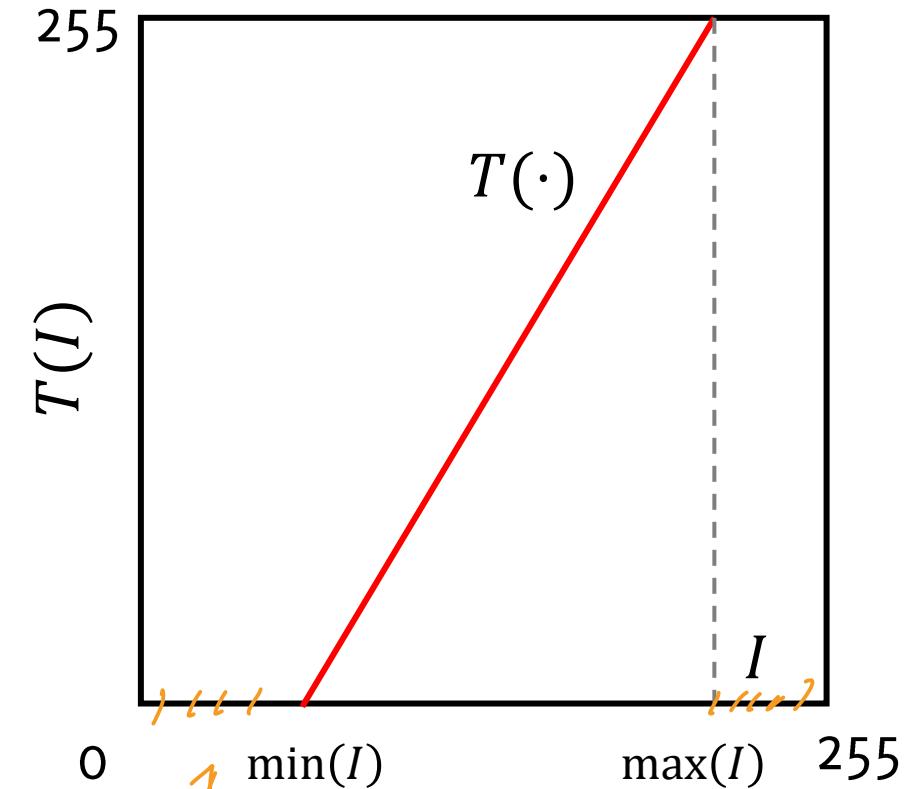
Intensity Rescaling

What does this transformation do?

rescaling
to bring
image light
in $[0 \div 1]$ (or $0 \div 255$)

← this is
stretching
the intensity
range

this one instead,
this take values between
 $\min(I)$ to $\max(I)$ and
map it linearly to $0 \div 255$



NO INTENSITY here!

Intensity Rescaling

In some cases images are conveniently mapped in the [0,255] range, covering such that

- $\min(T(I)) = 0$
- $\max(T(I)) = 255$

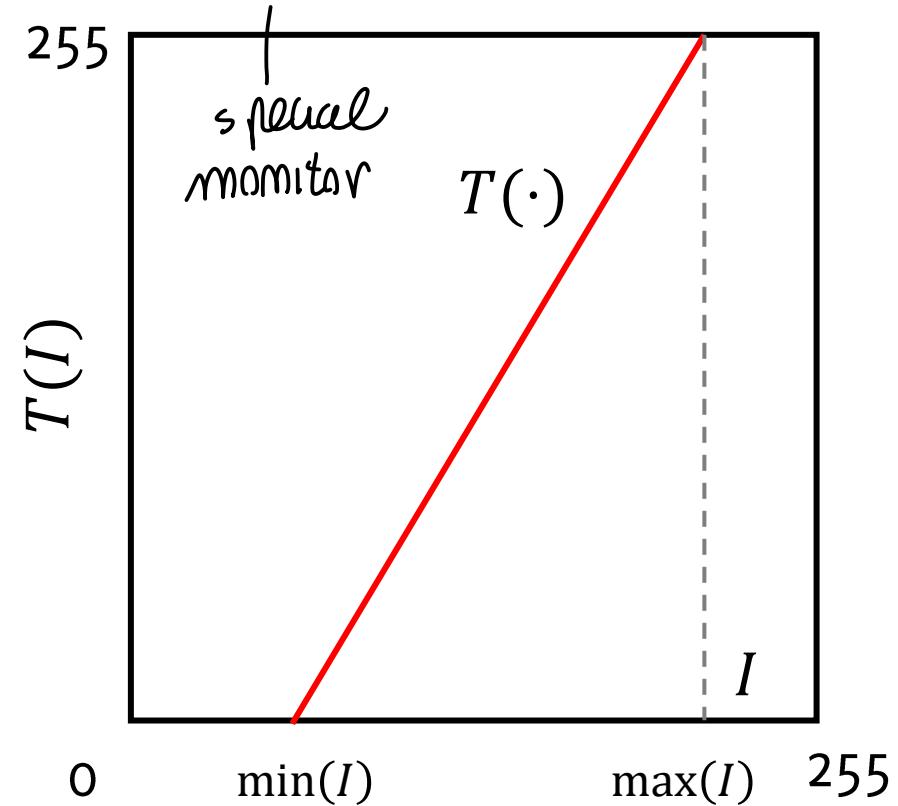
$$I(r, c) \rightarrow \left\lfloor 255 * \frac{I(r, c) - \min(I)}{\max(I) - \min(I)} \right\rfloor$$

This is a linear transformation of intensities

linear rescaling of all intensity to be sure that image take the entire range of intensities

you may want to rescale, to have all visualization range!
↳ meaning of all intensity to see better the intensity difference!

isn't always needed, isn't for images where you want to keep values, like in XRAY is important to preserve # absorbed radiation! → **IMPORTANT INFORMATION**

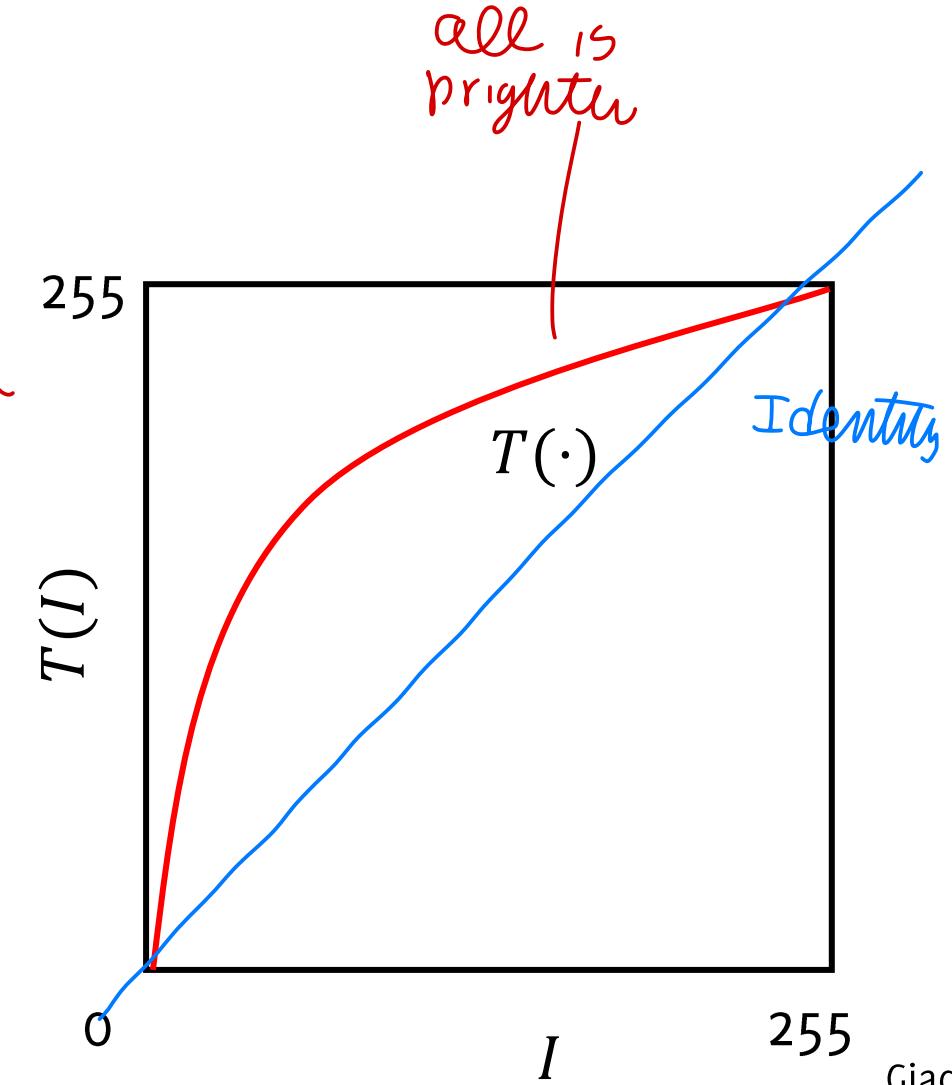


Gray-level mapping

A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately

not main purpose
↓
you increase
the CONTRAST
for low
intensity , the ones in $0 \div 20$ get
expanded a lot ! dark regions are
increased!

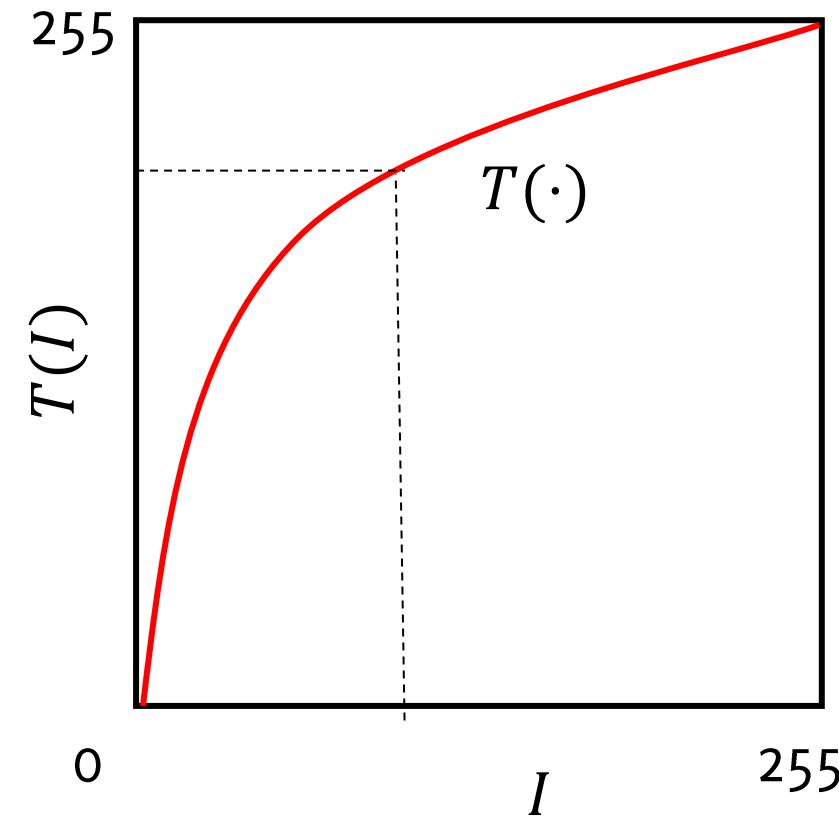
it makes image
brighter on
average !
| always above $T(I) = I$



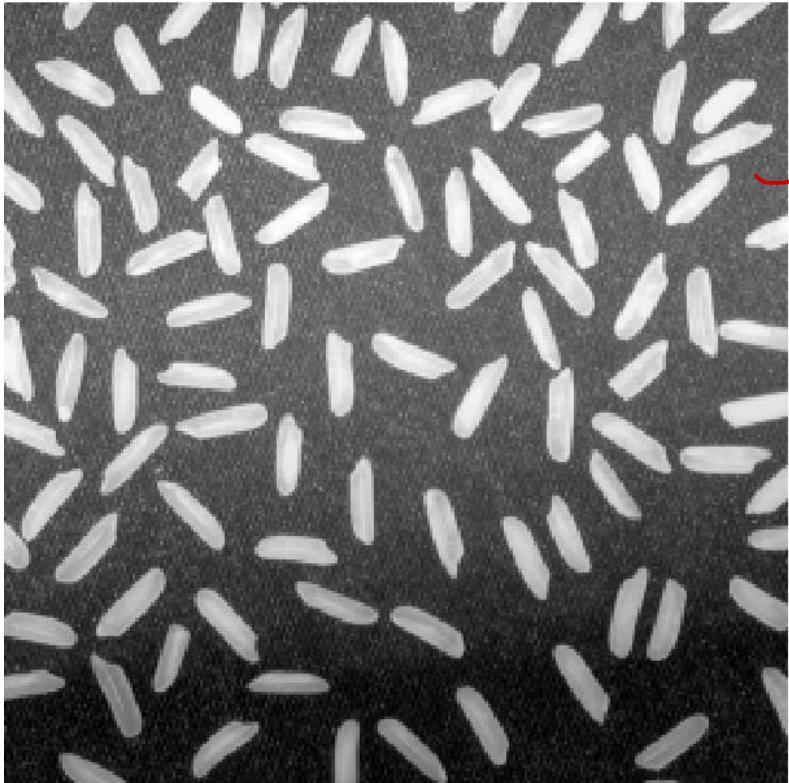
Gray-level mapping

A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately

What does this T do?

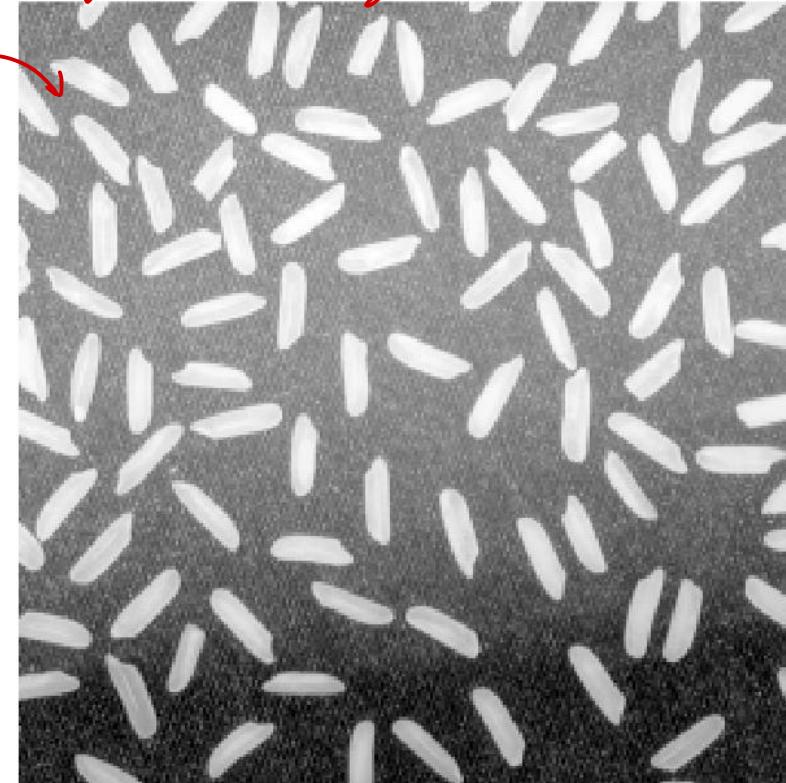


Contrast increases in dark, decreases in bright



Input I

more contrast in dark regions while
decreasing light regions contrast



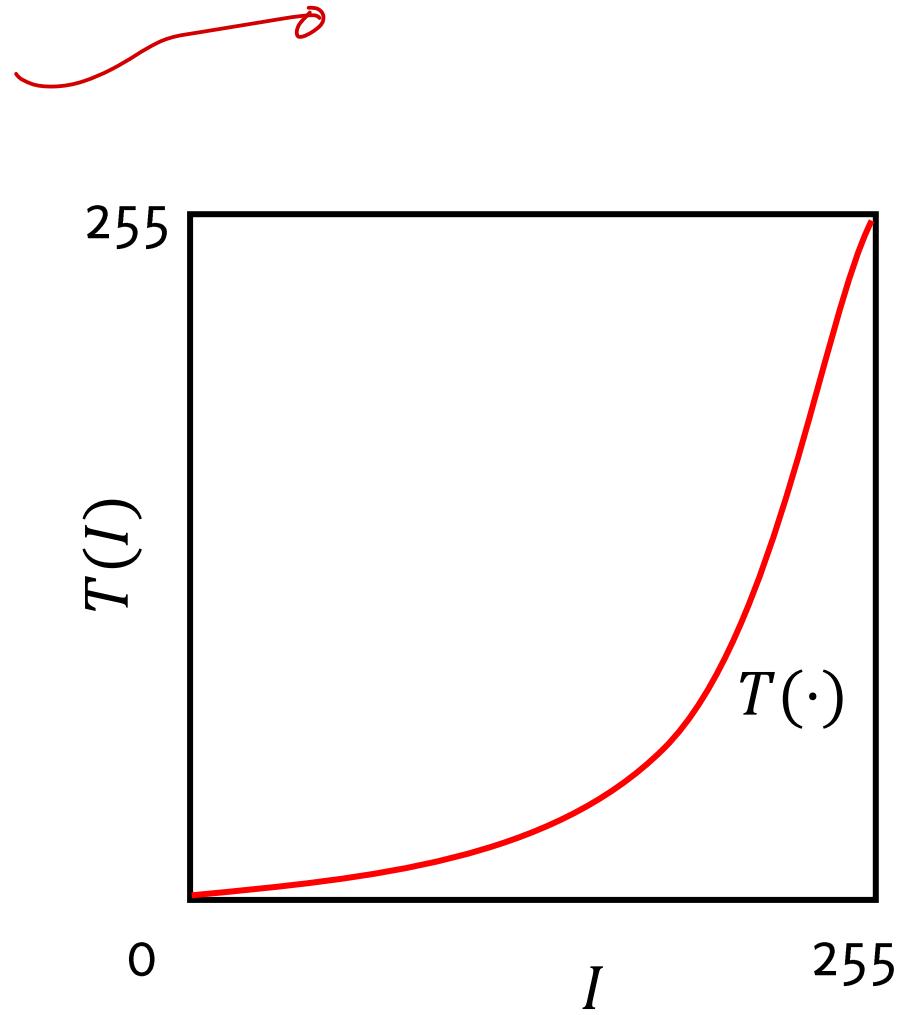
Output $G = T(I)$

Gray-level mapping

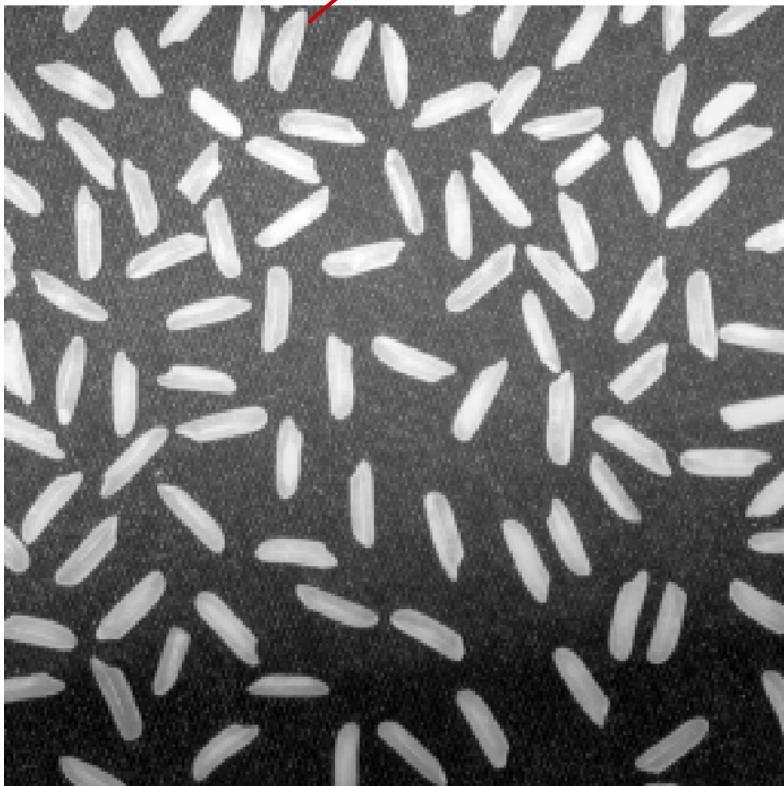
A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately

What does this T do?

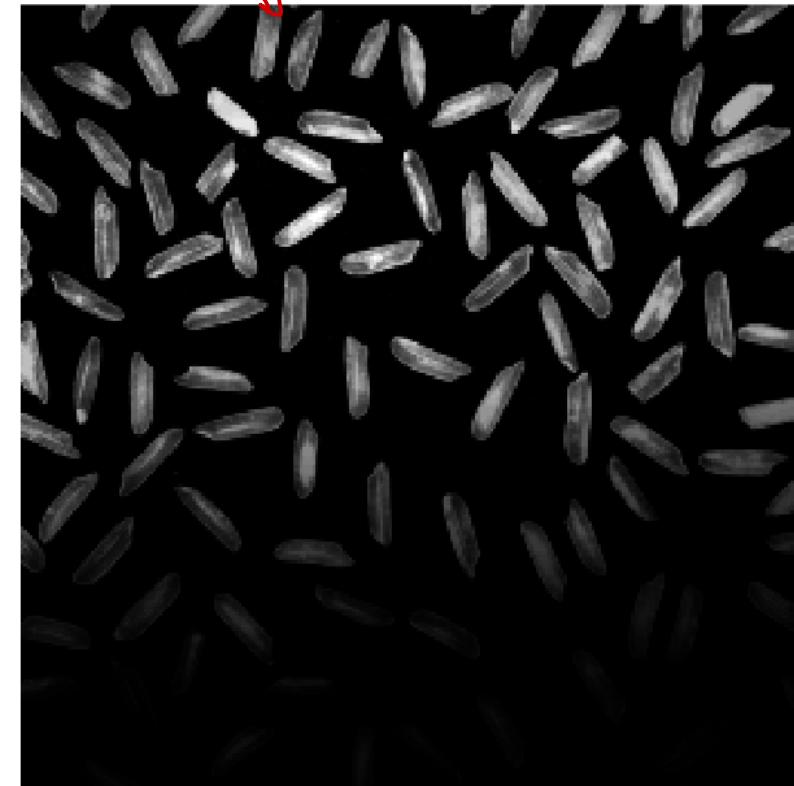
opposite of before



Contrast increases in bright, decreases in dark



Input I



more visible contrast
in bright region

Gamma Correction

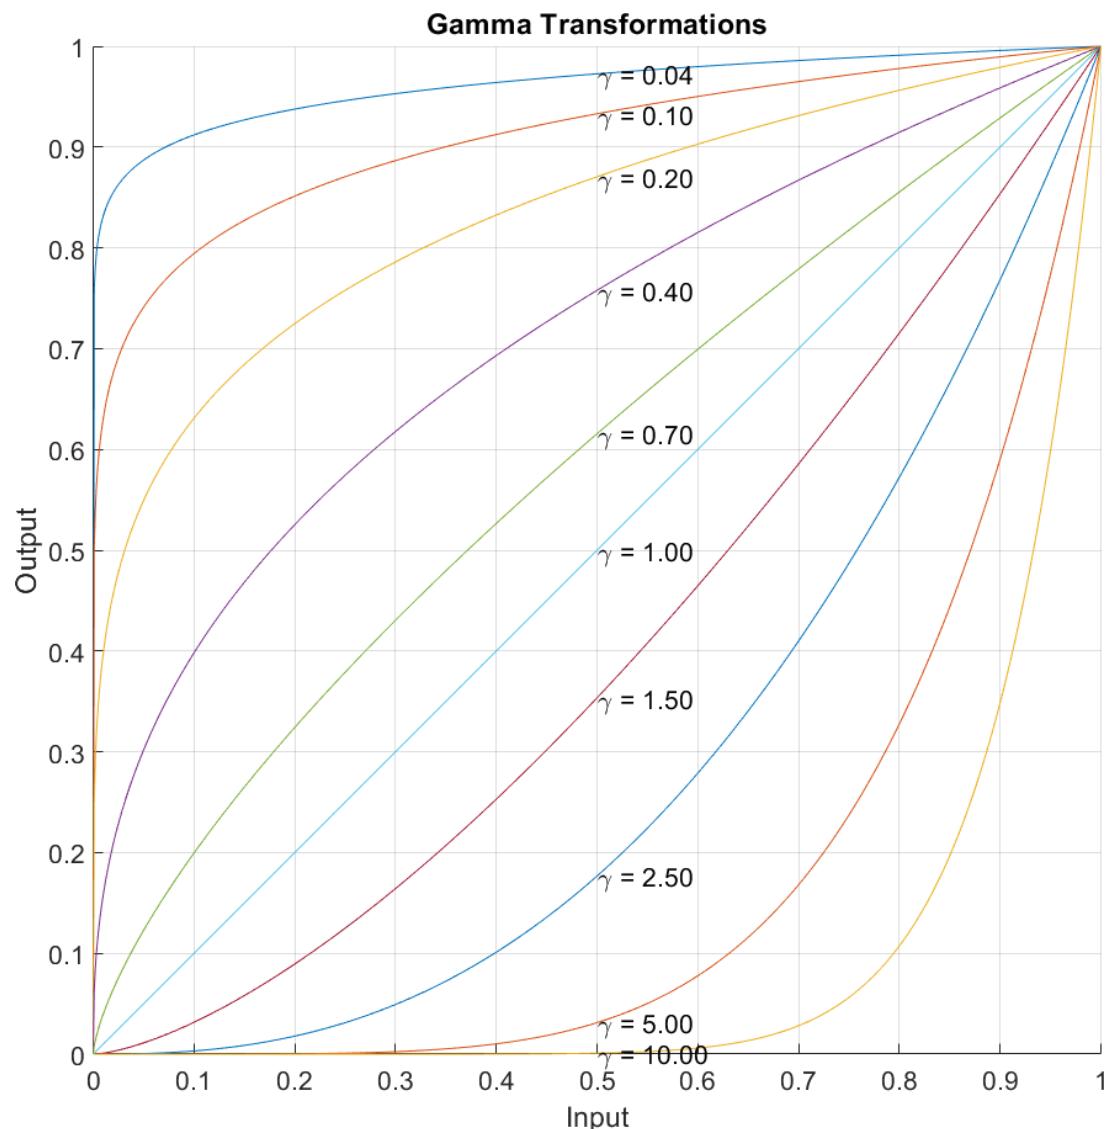
↓
that kind of
transformation are of this
correction type

Power-low transformation that
can be written as

$$G(r, c) = I(r, c)^\gamma$$

↓
take pixel intensity I
and raise to γ power

- { $\gamma = 1$: identity
- $\gamma > 1$: decrease dark
- $\gamma < 1$: increase dark



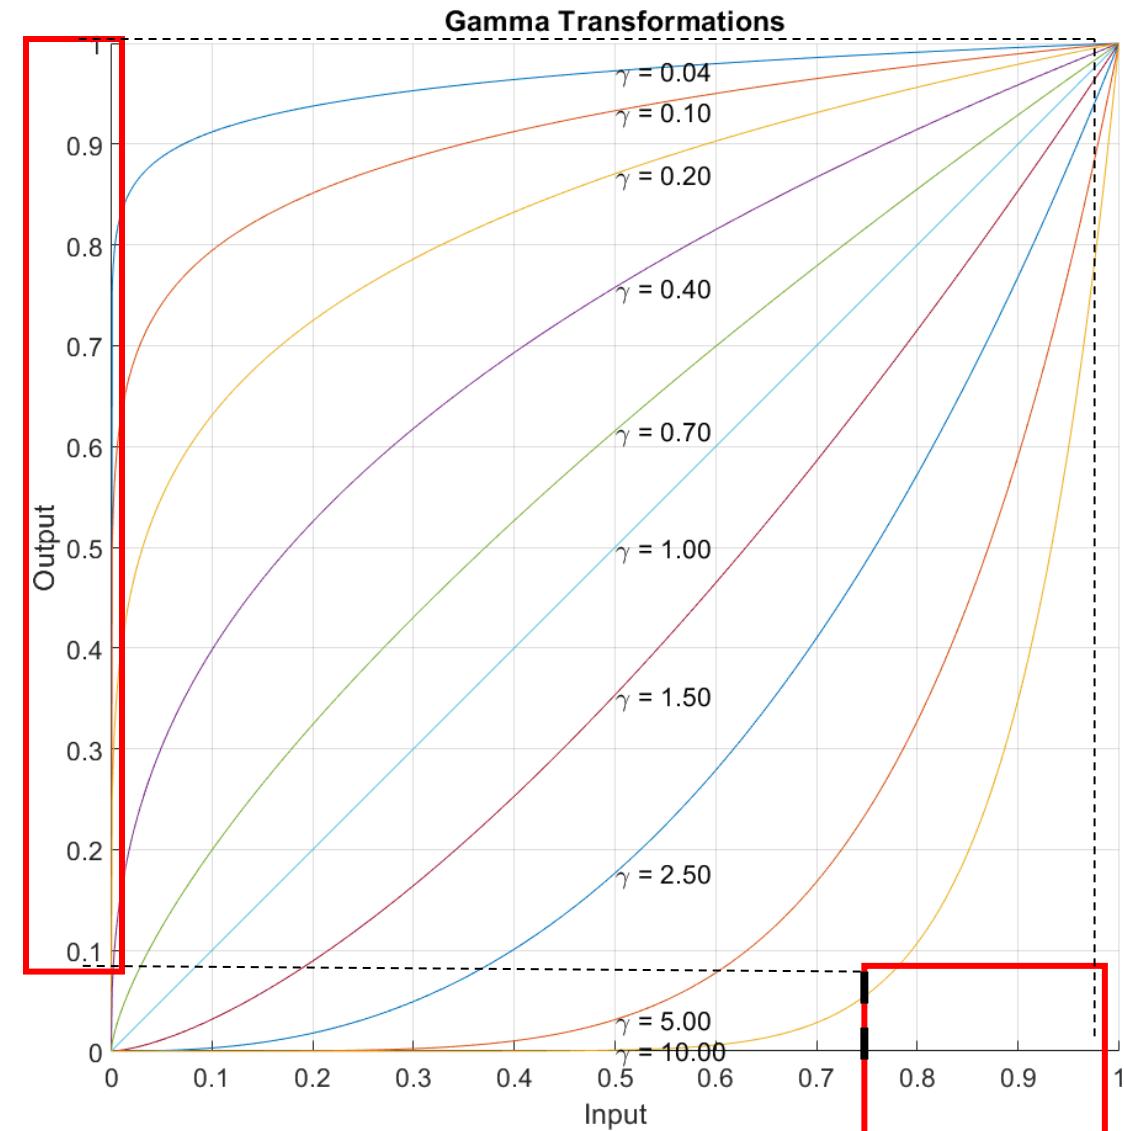
Gamma Correction

Power-low transformation that can be written as

$$G(r, c) = I(r, c)^\gamma$$

Contrast Enhancement:

- Low values of γ stretch the intensity range at high-values



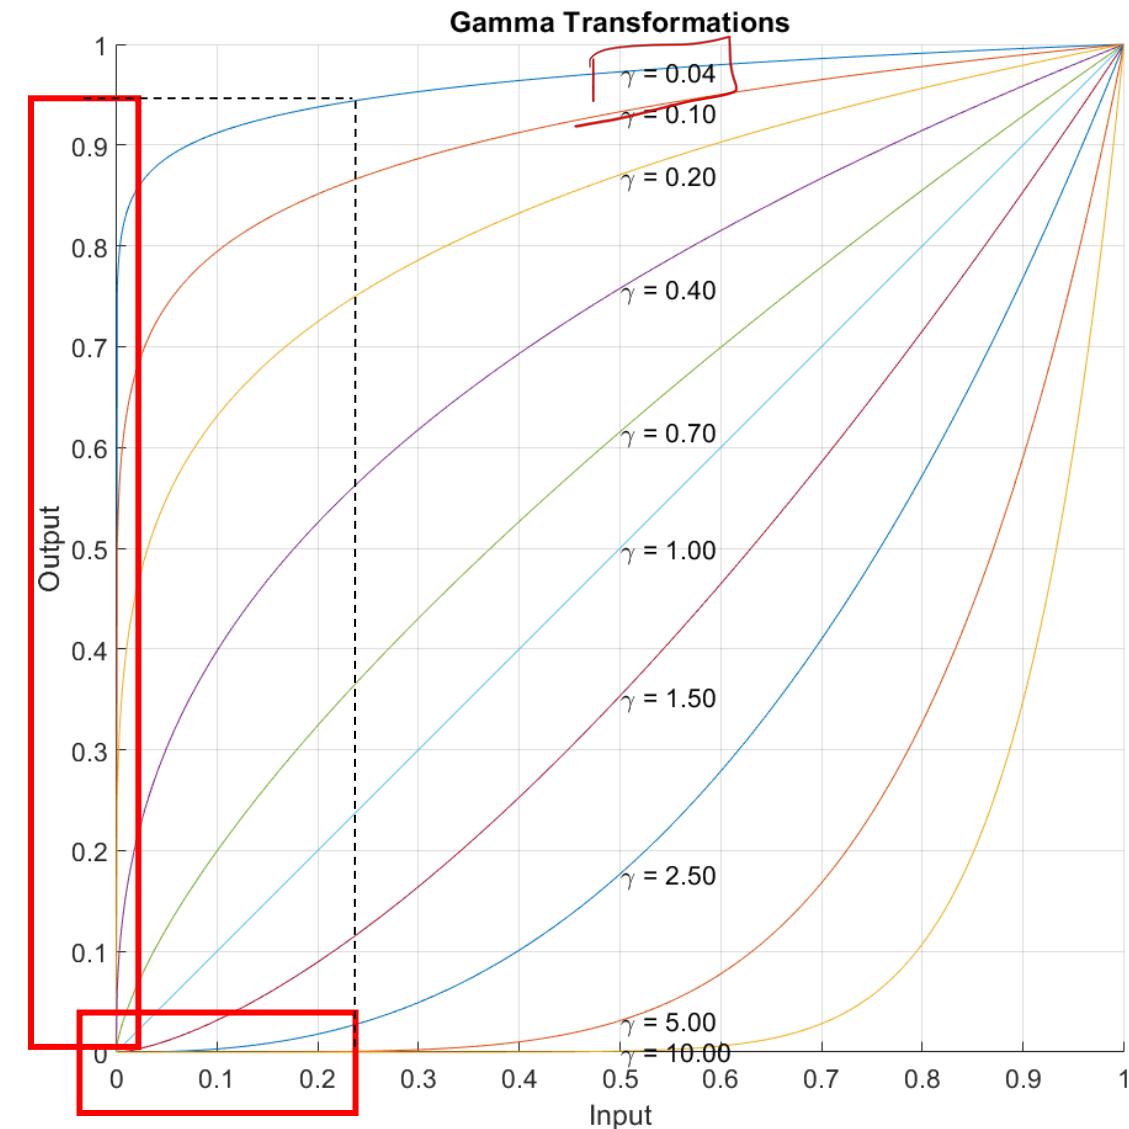
Gamma Correction

Power-low transformation that can be written as

$$G(r, c) = I(r, c)^\gamma$$

Contrast Enhancement:

- Low values of γ stretch the intensity range at high-values
- High values of γ stretch the intensity range at low values

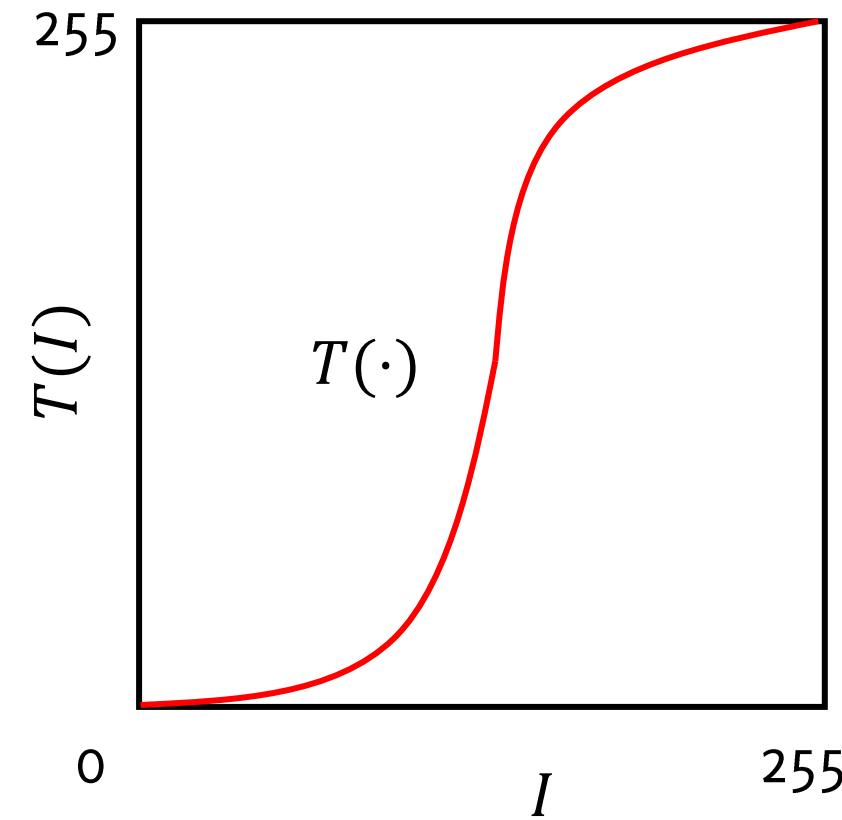


Gray Level Mapping

A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately

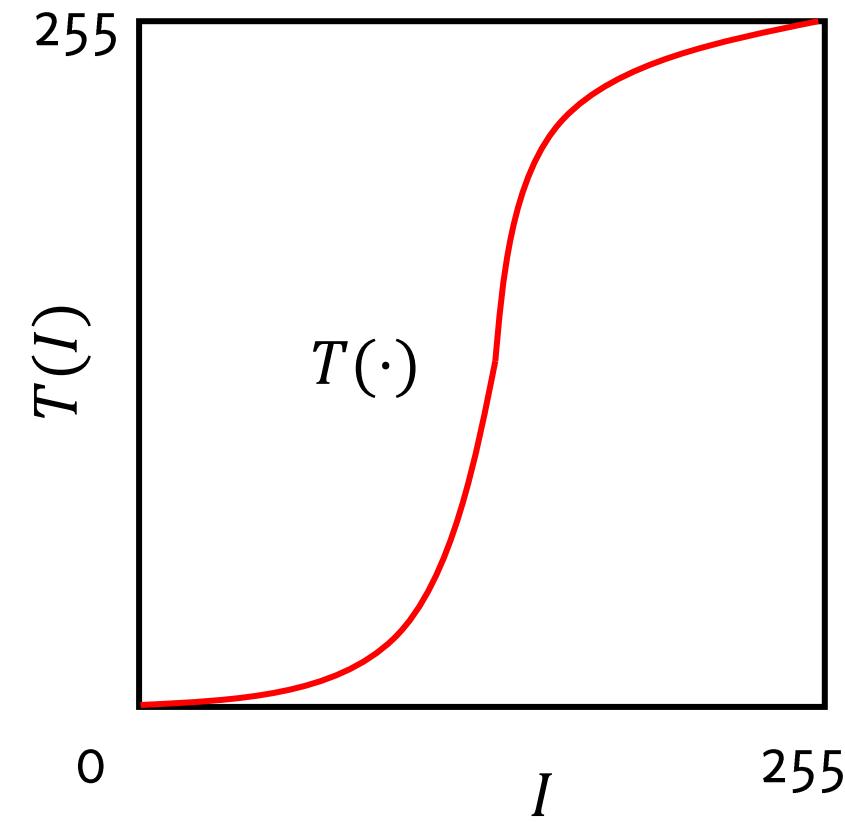
What does this T do?

↓
↑ increases CONTRAST
in a specific range of
intensity!



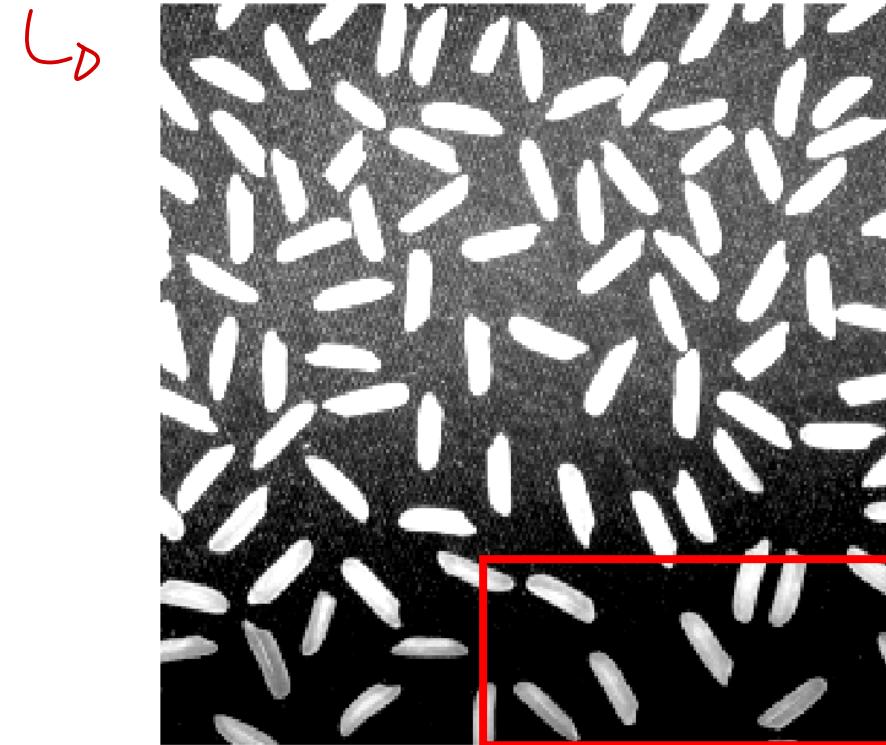
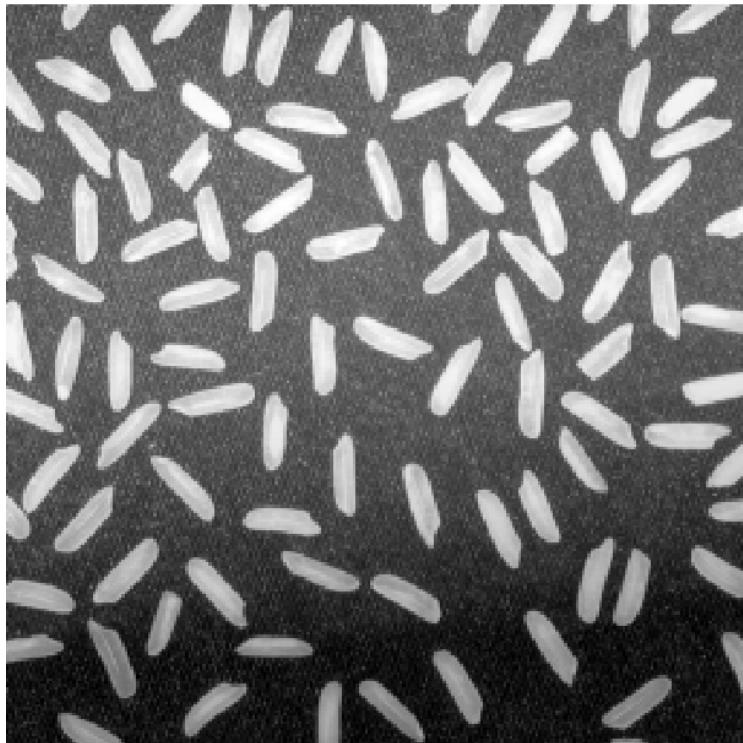
Contrast Stretching

A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately



Contrast Stretching

saturate math bright and dark regions ...



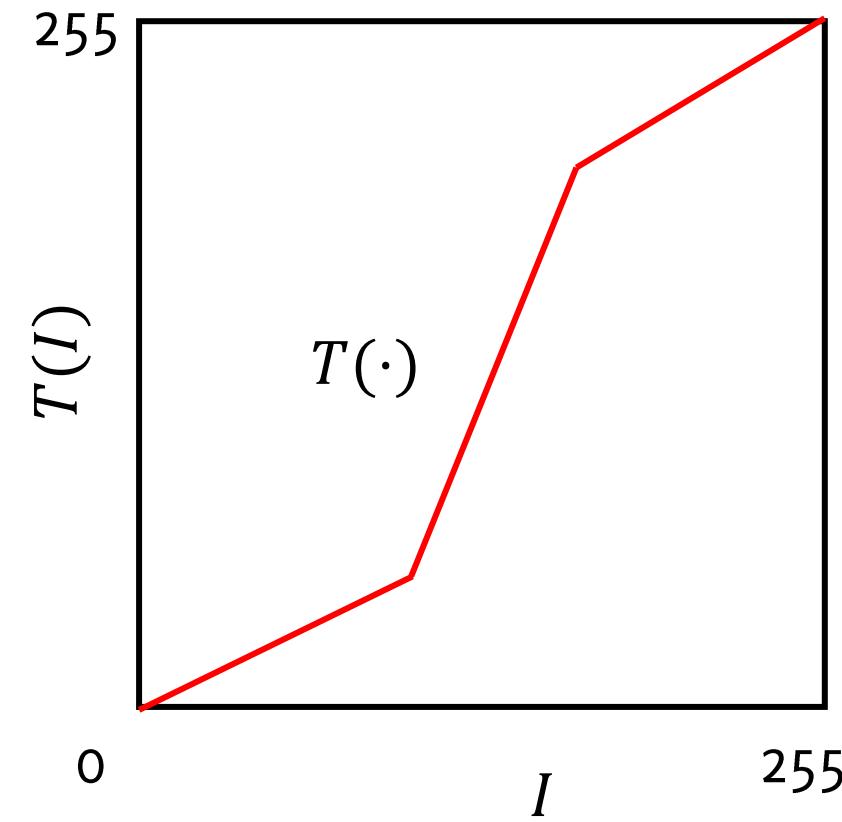
Contrast stretching: increases the constant at values in the middle of intensity range, decreases contrast at bright and dark regions.

It is implemented by piecewise or parametric transformations

Contrast Stretching

(also piecewise is possible)

Can be defined by piecewise linear mapping...

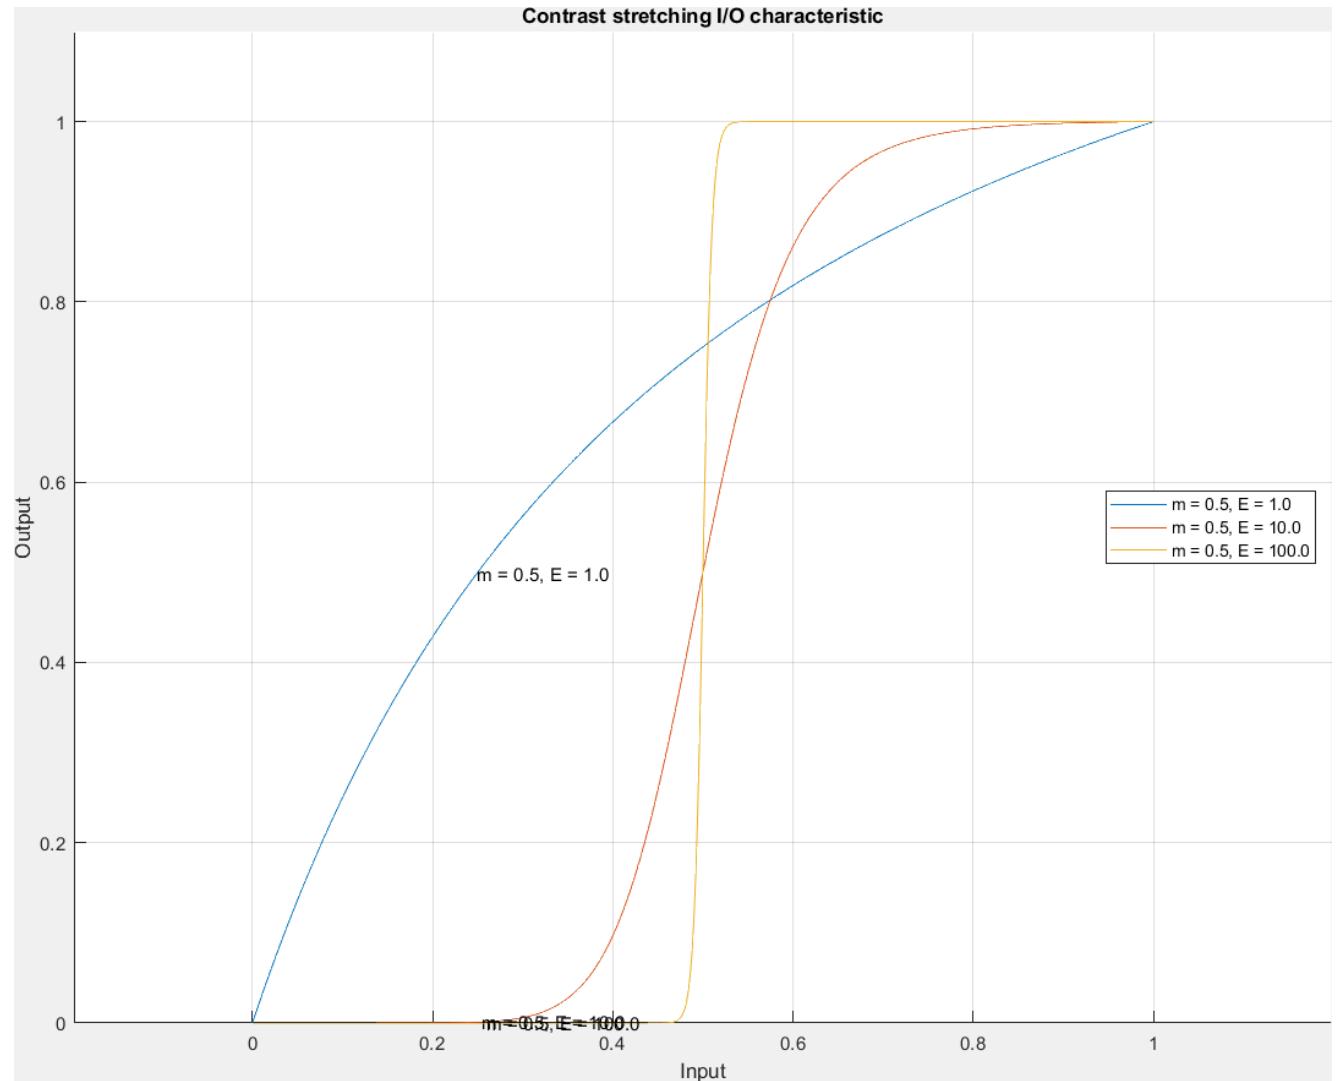


Contrast Stretching

And there are also analytical expressions

$$I(r, c) \rightarrow \frac{1 + m^e}{\left(1 + \left(\frac{m}{I(r, c) + \epsilon}\right)\right)^e}$$

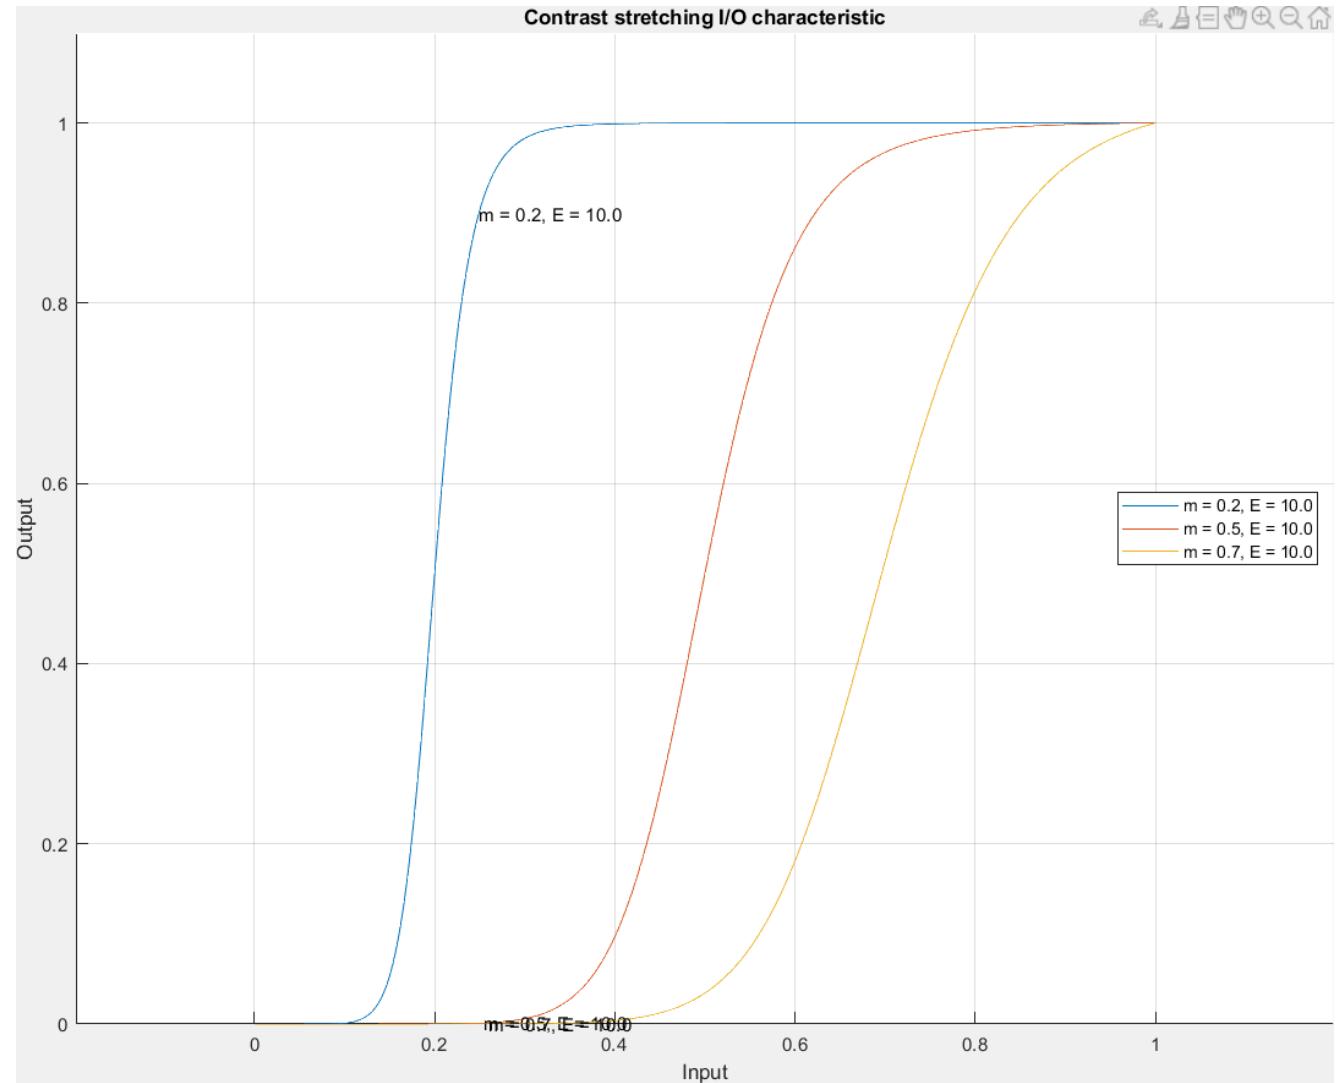
by analytical
desuptum



Contrast Stretching

And there are also analytical expressions

$$I(r, c) \rightarrow \frac{1 + m^e}{\left(1 + \left(\frac{m}{I(r, c) + \epsilon}\right)\right)^e}$$

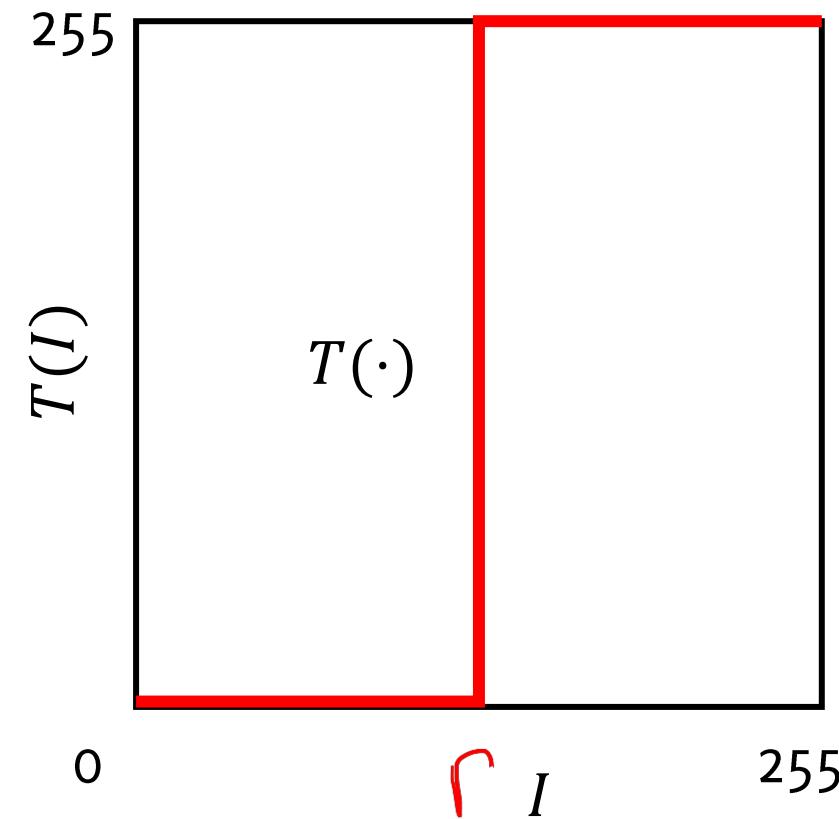


Gray-level mapping

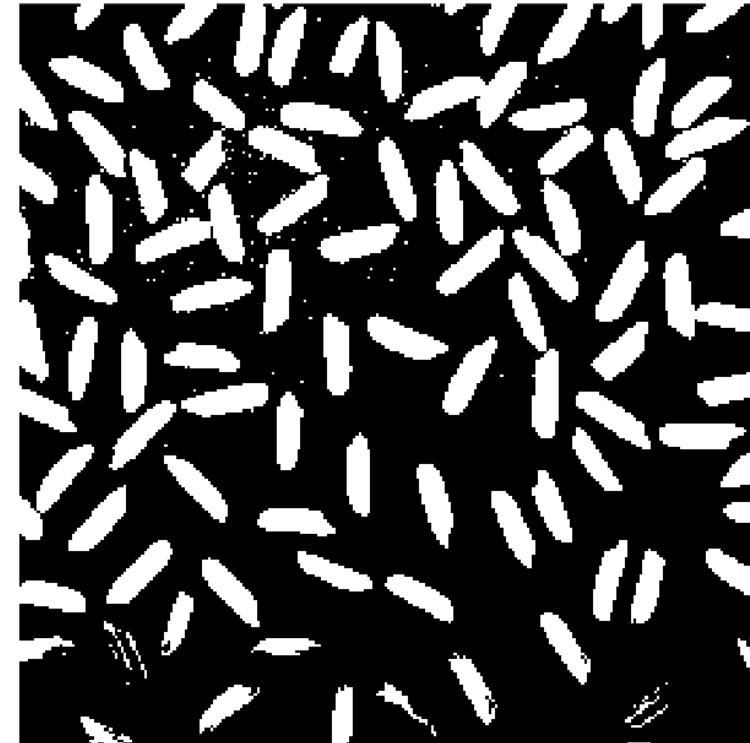
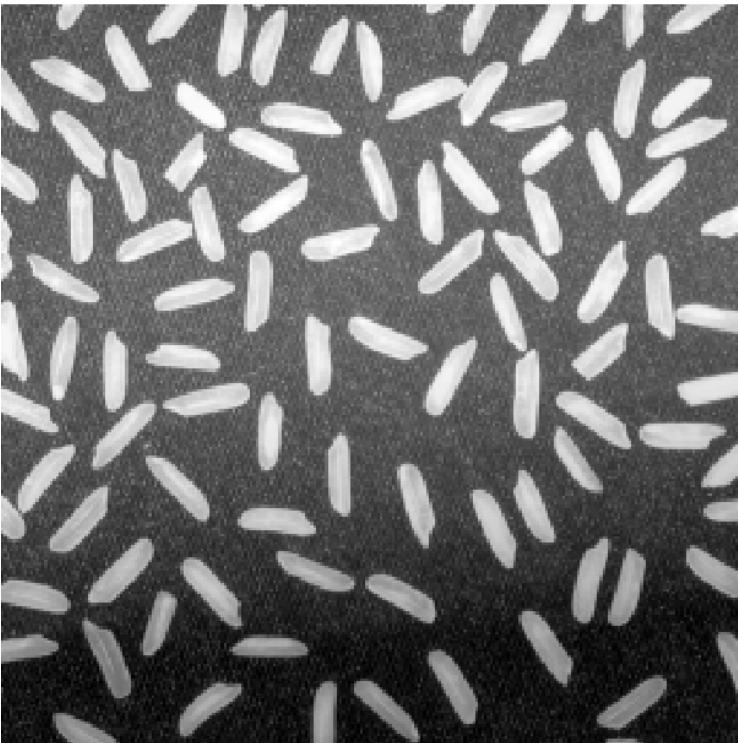
A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately

What does this T do?

↪
binary image
 $\rightarrow 0 \div 1$ by
 γ + threshold
to change
brightness



Thresholding



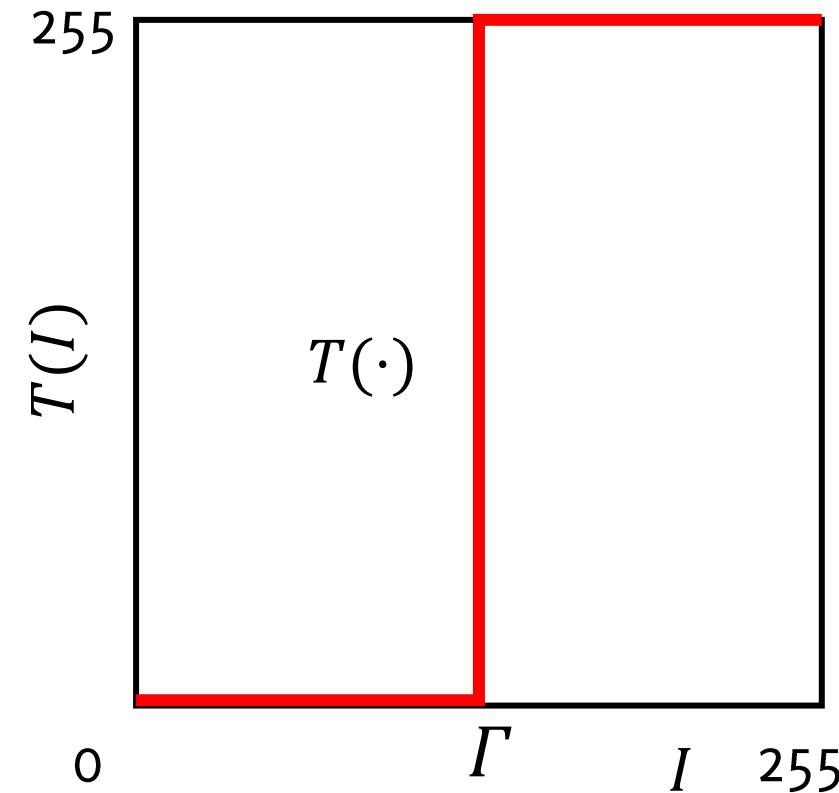
Thresholding binarizes images



Thresholding

A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately

$$T(I(r, c)) = \begin{cases} 255, & \text{if } I(r, c) \geq \Gamma \\ 0, & \text{if } I(r, c) < \Gamma \end{cases}$$

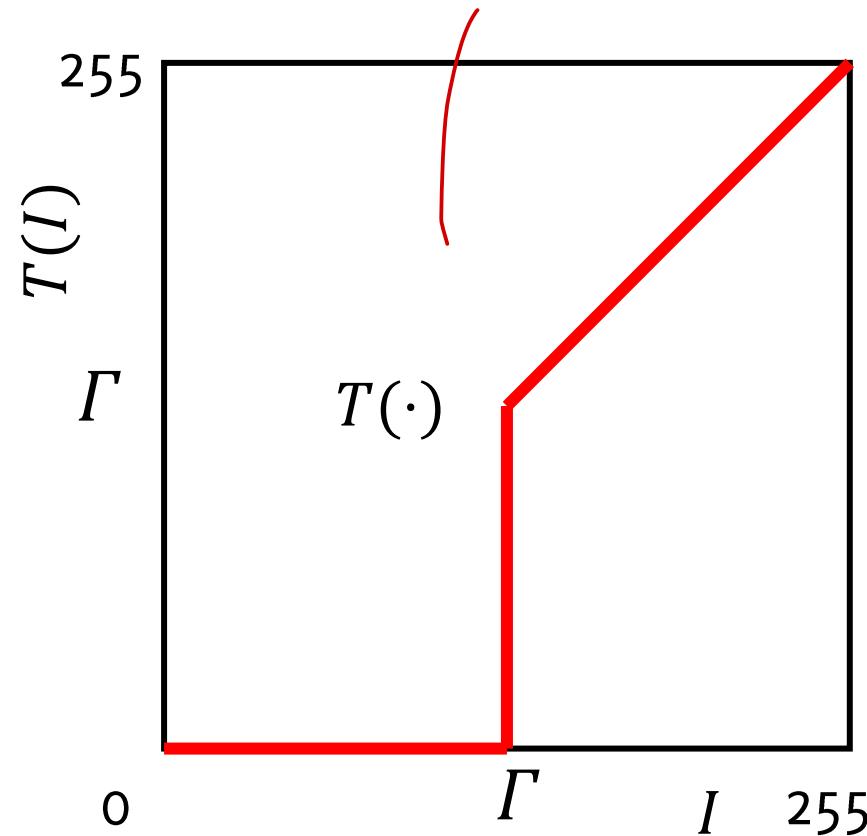


Thresholding

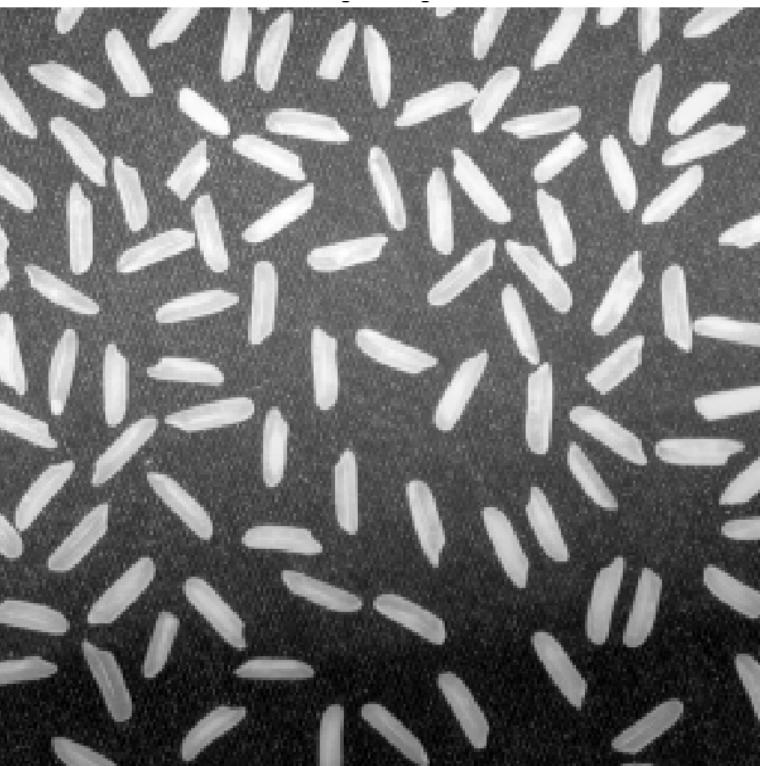
A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately

What does this T do?

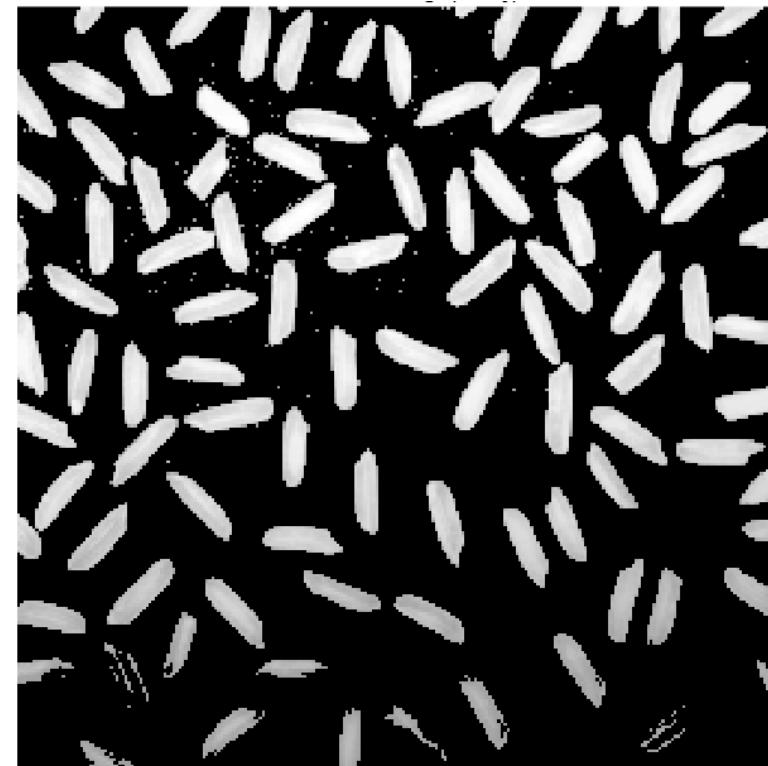
this extract only
pixels higher than a
certain threshold!



Thresholding



only some bright
parts are taken



simple
image processing
operations ...
by segmentation

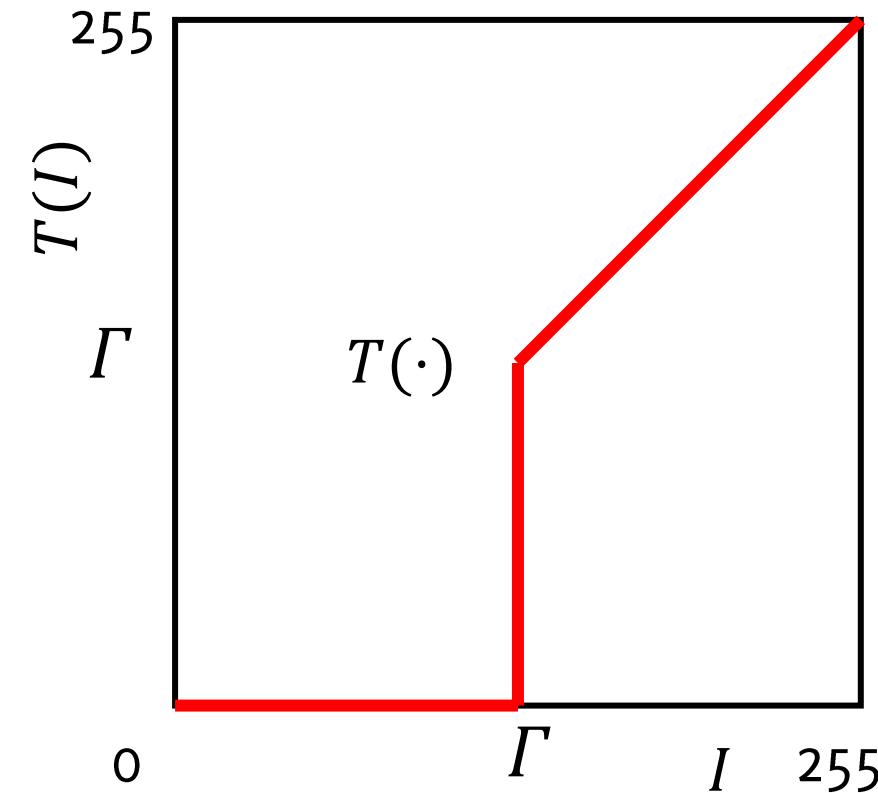
Thresholding

A transformation $T: \mathbb{R} \rightarrow \mathbb{R}$ that operates on gray-scale images or on each color-plane separately

$$T(I(r, c)) = \begin{cases} T(I(r, c)), & \text{if } I(r, c) \geq \Gamma \\ 0, & \text{if } I(r, c) < \Gamma \end{cases}$$

as used in ARTIFICIAL
Neural Networks ReLU activation
with $\Gamma=0$

This simple operation is one of the most frequently used to add nonlinearities in CNN: the ReLU Layers are like these with $\Gamma = 0$ (values are allowed to be negative)



← that transformation act on each pixel, it is fine
for any data, NOT only for images!

↓ there are better functions
taught for images

Local (Spatial) Transformations: Correlation and Convolution

Local (Spatial) Transformation

In general, these can be written as

$$G(r, c) = T_U[I](r, c)$$

Where

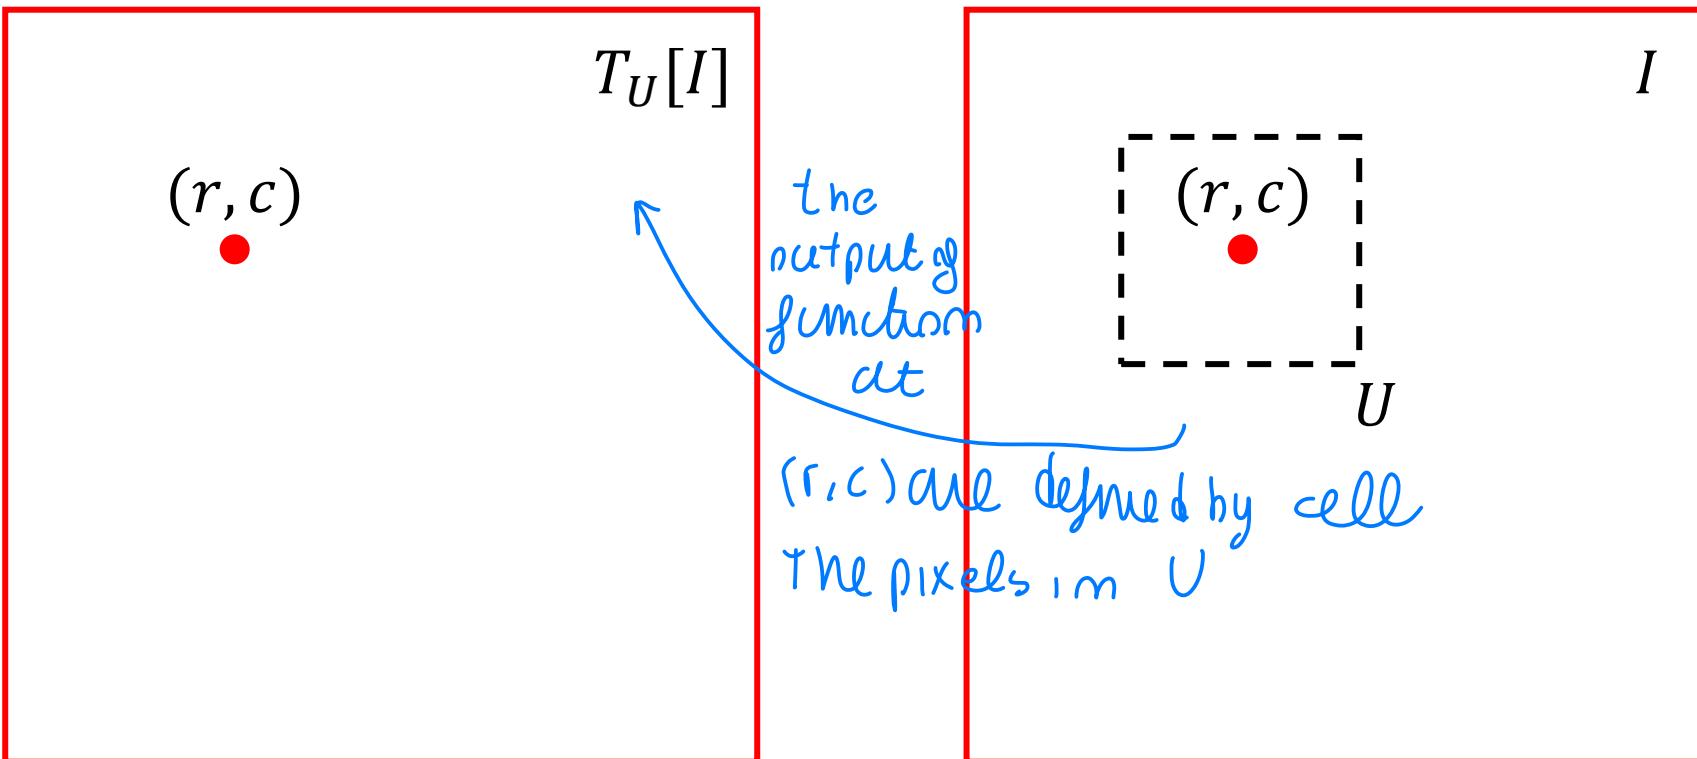
- I is the input image to be transformed
- G is the output
- U is a **neighbourhood**, identifies a region of the image that will concur in the output definition, e.g. $U = \{(-1, -1), (-1, 0), \dots, (0, 0), \dots, (1, 0), (1, 1)\}$
- $T_U: \mathbb{R}^3 \rightarrow \mathbb{R}^3$ or $T_U: \mathbb{R}^3 \rightarrow \mathbb{R}$ is a function

T operates on I “around” the pixel (r, c) on a specific neighborhood U . In particular, $T_U[I](r, c)$ is computed from all the intensity values: $\{I(u, v), (u + r, v + c) \in U\}$

the output in a pixel
is the transformation
of the entire image,
evaluated at that
pixel! before we
see the result directly
on pixel

Local (Spatial) Filters

The dashed square represents $\{I(u, v), (u + r, v + c) \in U\}$



U defines, for each pixel I , which are the neighbourhood of (r, c)

to define neighborhood \mathcal{U}



to define all surrounding pixels

a convenient way is to define a set of

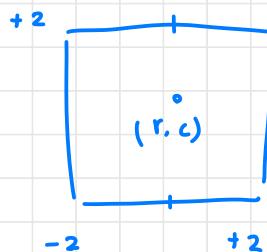
displacements $(r + ir, c + ic)$



including $0,0 \dots$

this displacements
are used to
write down

$$\forall (r, c) \in \mathcal{U} \\ (r+u, c+v), (u, v) \in \mathcal{U} \\ \dots$$



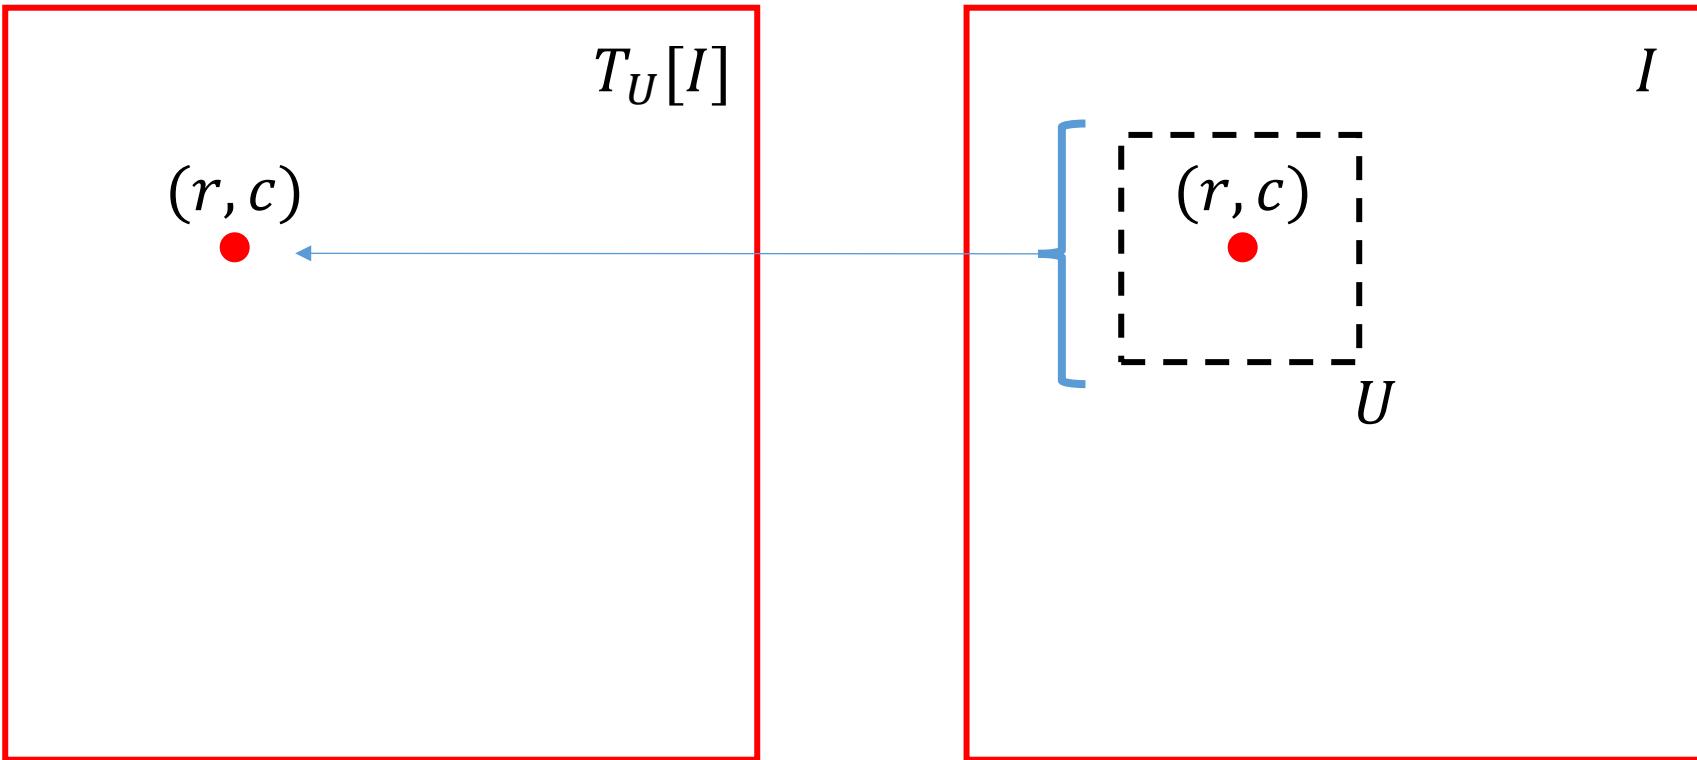
with a
specific
granularity

← used for mathematical reasons

$\forall (r, c)$ I can take pixels around

Local (Spatial) Filters

The dashed square represents $\{I(u, v), (u + r, v + c) \in U\}$

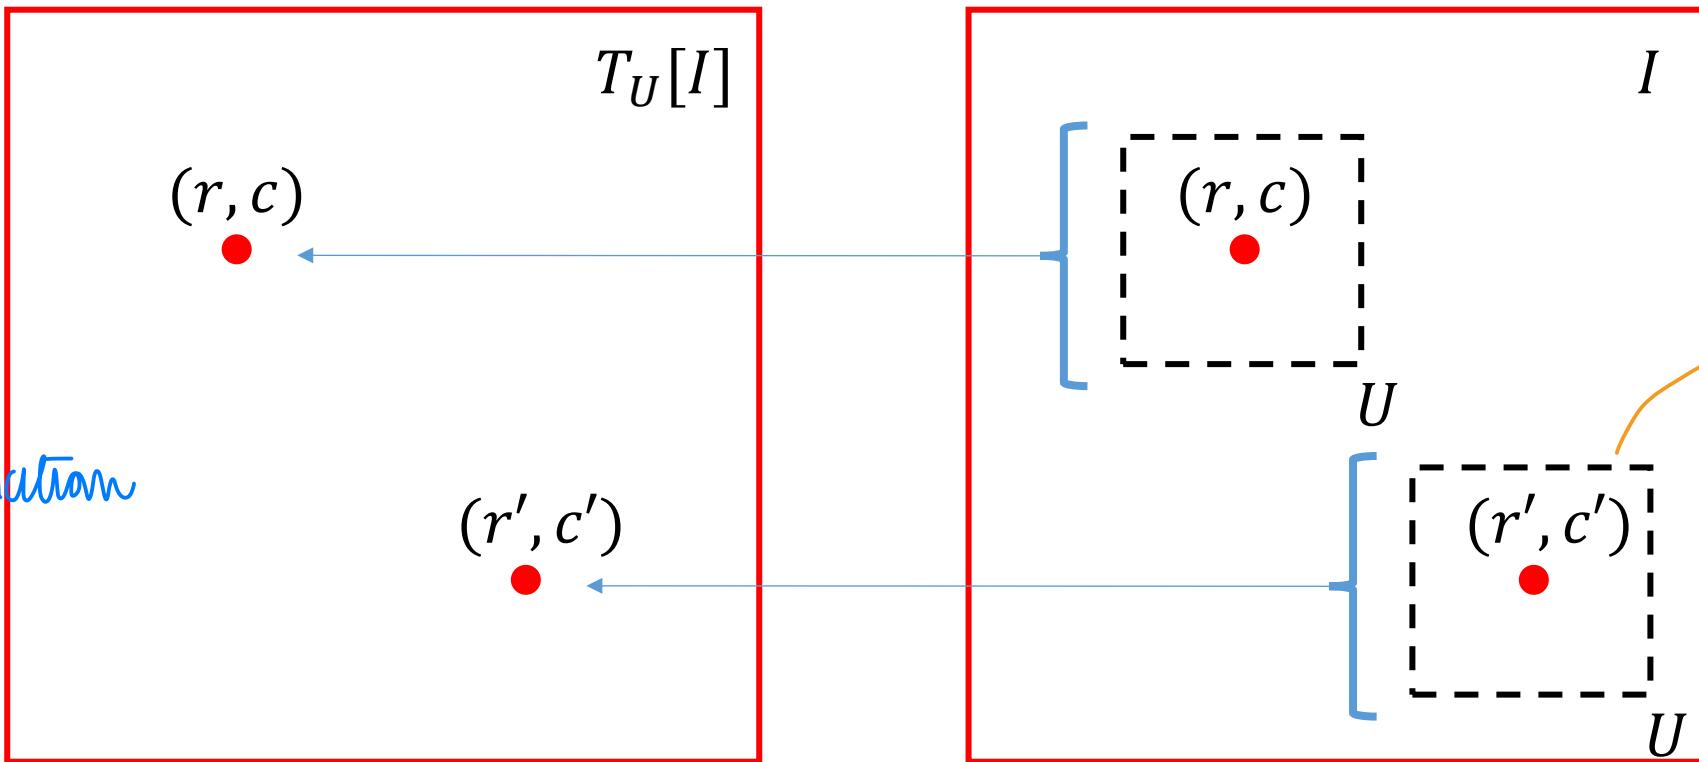


Local (Spatial) Filters

AND the same function is used in each location of the image, output map point is computed by selecting neighborhood and computing an expression over cell U

The dashed square represents $\{I(u, v), (u + r, v + c) \in U\}$

from set
of values
I get an
output value,
mainly
the transformation
escal!



- The location of the output does not change
- **Space invariant transformations** are repeated for each pixel
- T can be either linear or nonlinear

← the TRANSFORMATION is

LOCAL : because depends on neighborhood
and

SPATIAL : because it have same
INVARIANT expression every where

Local Linear Filters

↑ taken all values around (r,c) summing r+u, c+v and
↑ then 1/|U| average

↓ simplest transformation is a kind of average in 1 ORD approx

Linear Transformation: Linearity implies that

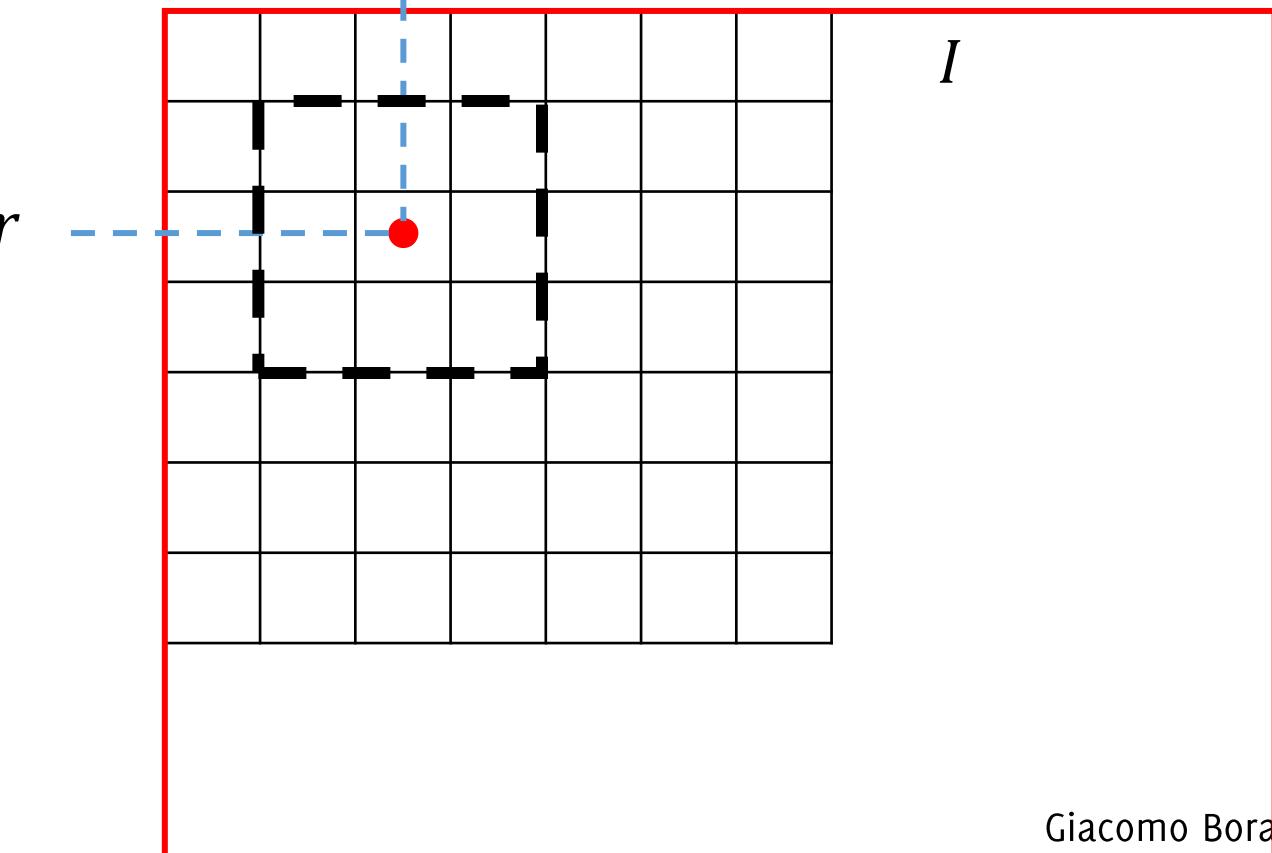
↳ replace each point by avg over neighborhood

"
MAKING AVERAGE FILTER

$$T[I](r, c) = \sum_{(u,v) \in U} w(u, v) * I(r + u, c + v)$$

Considering some weights $\{w_i\}$
and (u, v) has to be interpreted as
a "displacement vector" w.r.t. the
neighborhood center (r, c) , e.g.,
 $(u, v) \in \{(1, -1), (1, 0), (1, -1) \dots\}$

We can now think the neighborhood
 U as set of displacement vectors.

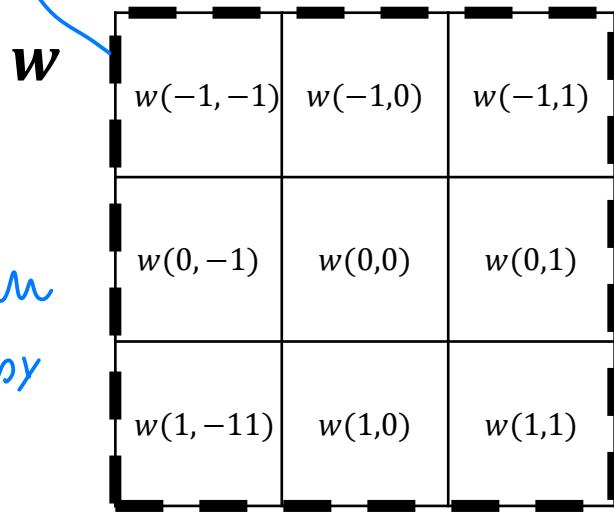


Local Linear Filters

correlation formula

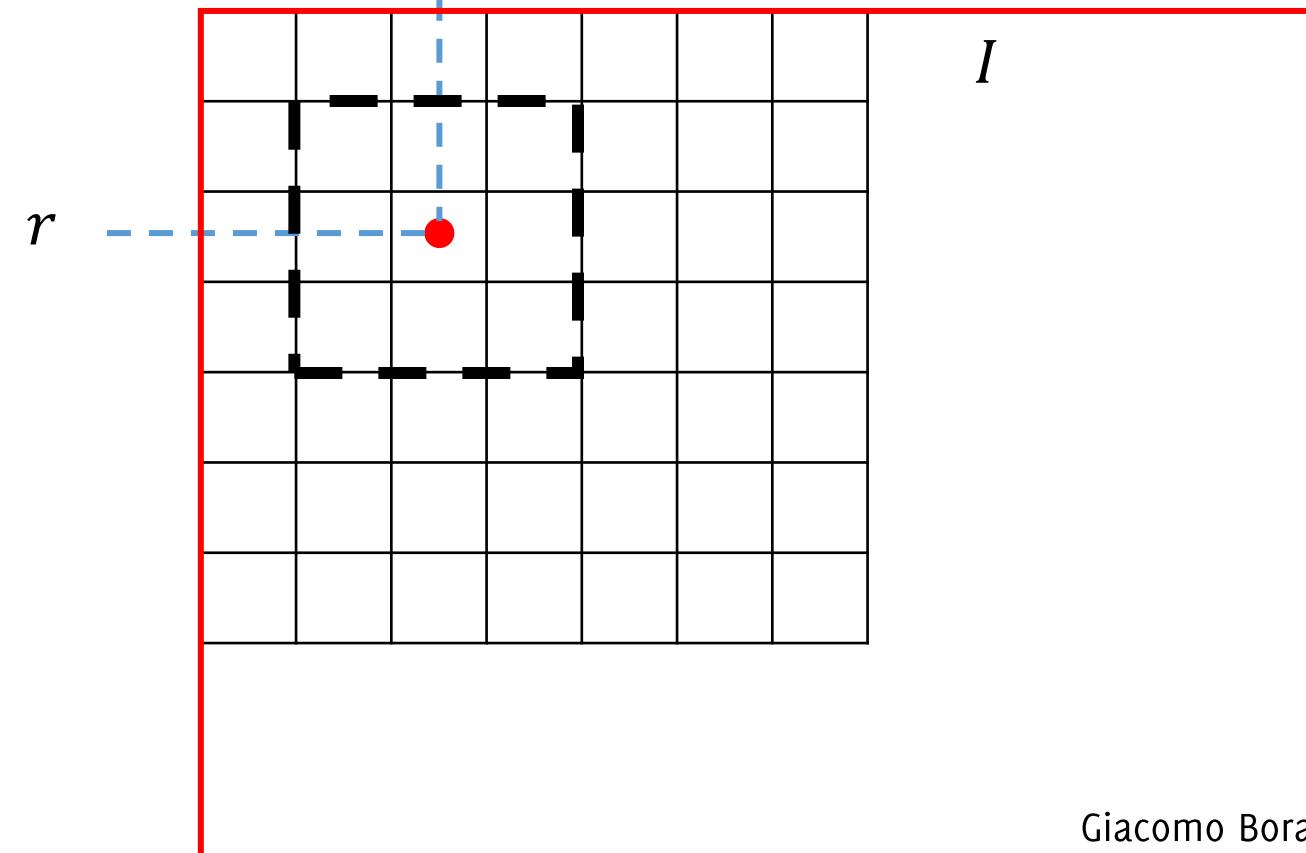
Linear Transformation: the filter weights can be associated to a matrix w

by using weight we have more DOF in the allage
↓
weights themselves can be described by MATRIX of same dimension of neighbourhood set
↓
CORRELATION



We can consider weights as an image, or a filter h
The filter h entirely defines this operation

$$T[I](r, c) = \sum_{(u,v) \in U} w(u, v) * I(r + u, c + v)$$



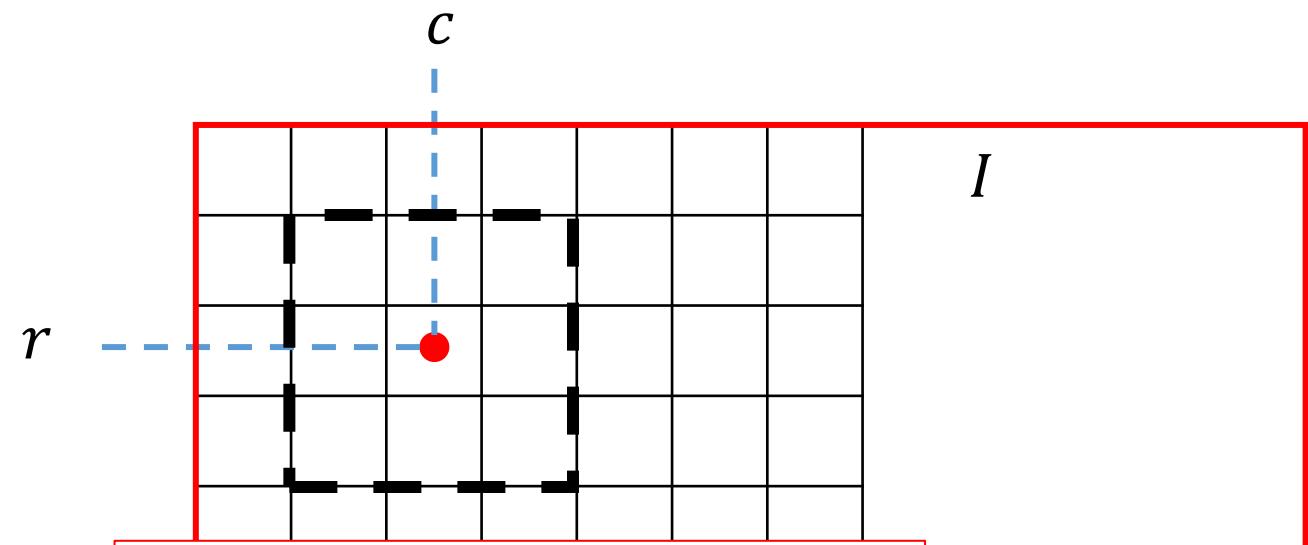
Local Linear Filters

Linear Transformation: the filter weights can be associated to a matrix w

$$T[I](r, c) = \sum_{(u,v) \in U} w(u, v) * I(r + u, c + v)$$

The diagram shows a 3x3 grid of squares representing filter weights. The squares are labeled with their coordinates: top row has w(-1,-1), w(-1,0), w(-1,1); middle row has w(0,-1), w(0,0), w(0,1); bottom row has w(1,-1), w(1,0), w(1,1). The entire grid is enclosed in a double-lined rectangular border.

We can consider weights as an image, or a filter h
The filter h entirely defines this operation



This operation is repeated for each pixel in the input image

Local Linear Filters

this is called CORRELATION = $\begin{cases} \text{convolution} \\ \text{in ANN} \end{cases}$

Linear Transformation: Linearity implies that the output $T[I](r, c)$ is a linear combination of the pixels in U :

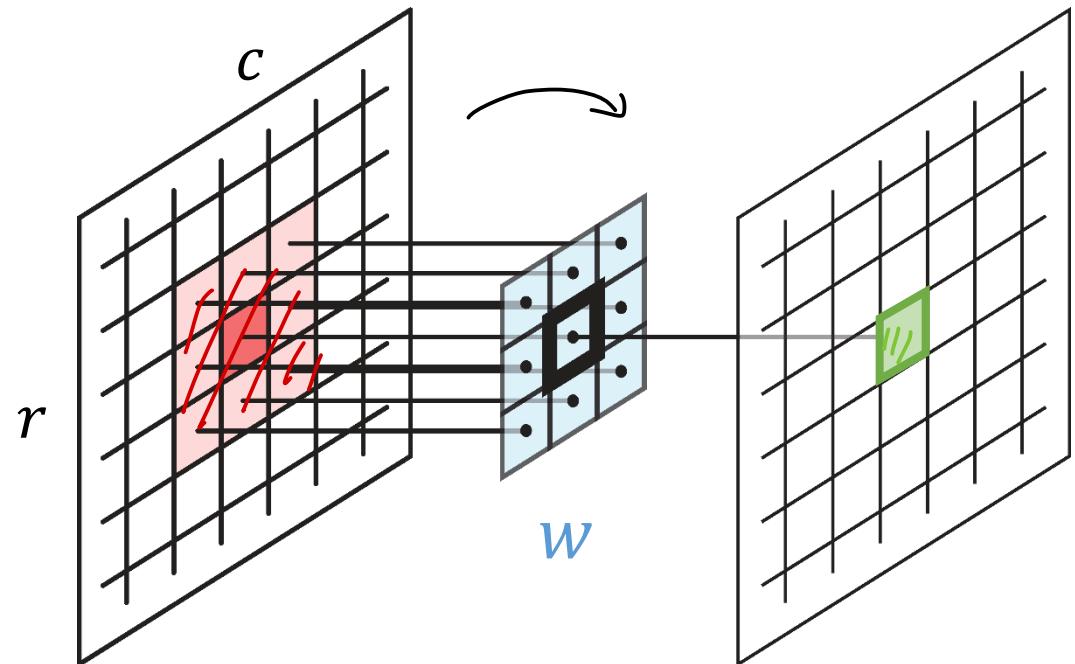
$$T[I](r, c) = \sum_{(u,v) \in U} w(u, v) * I(r + u, c + v)$$

elementwise
product by FILTERING
and summing

Considering some weights $\{w_i\}$

We can consider weights as an image, or a filter w

The filter h entirely defines this operation



in ANN this concept is related to

CONVOLUTIONAL NEURAL
NETWORKS

Local Linear Filters

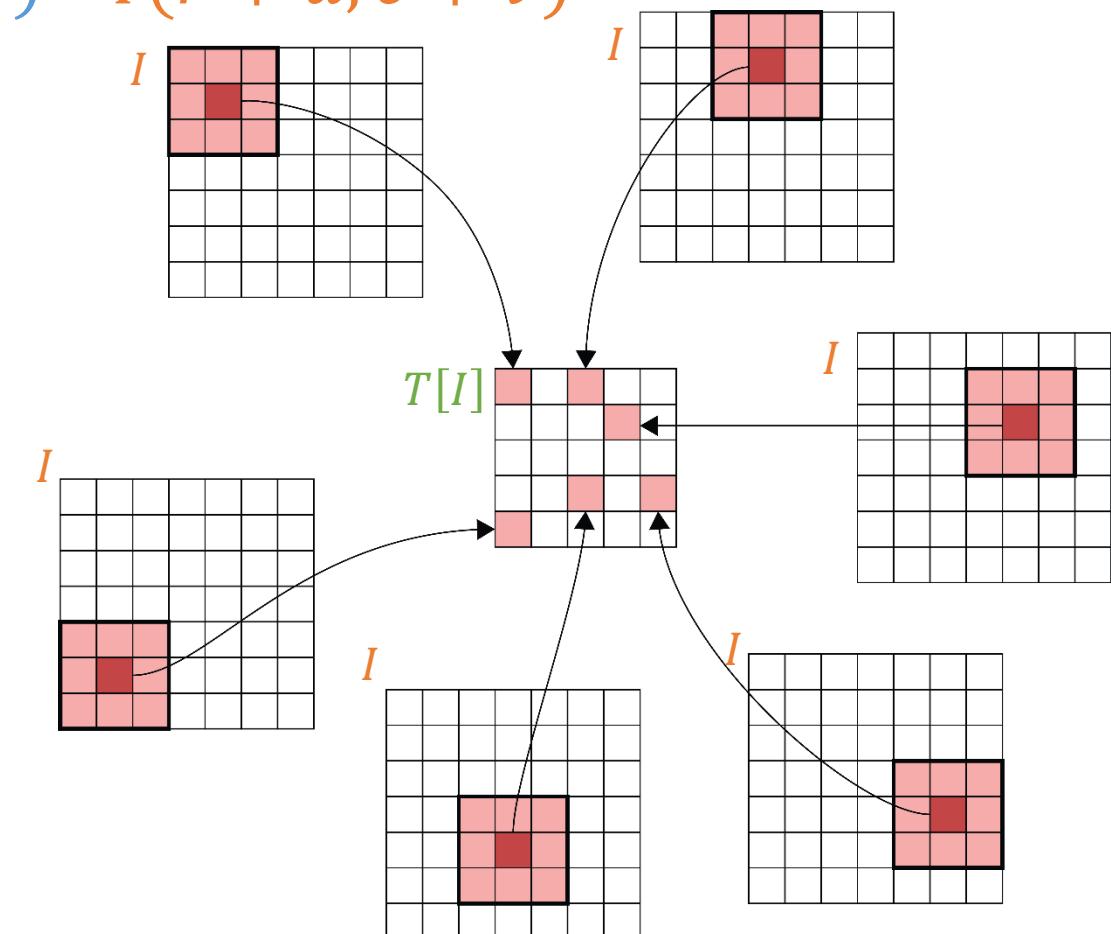
Linear Transformation: the filter weights can be associated to a matrix w

$$T[I](r, c) = \sum_{(u,v) \in U} w(u, v) * I(r + u, c + v)$$

w

$w(-1, -1)$	$w(-1, 0)$	$w(-1, 1)$
$w(0, -1)$	$w(0, 0)$	$w(0, 1)$
$w(1, -1)$	$w(1, 0)$	$w(1, 1)$

This operation is repeated for each pixel in the input image

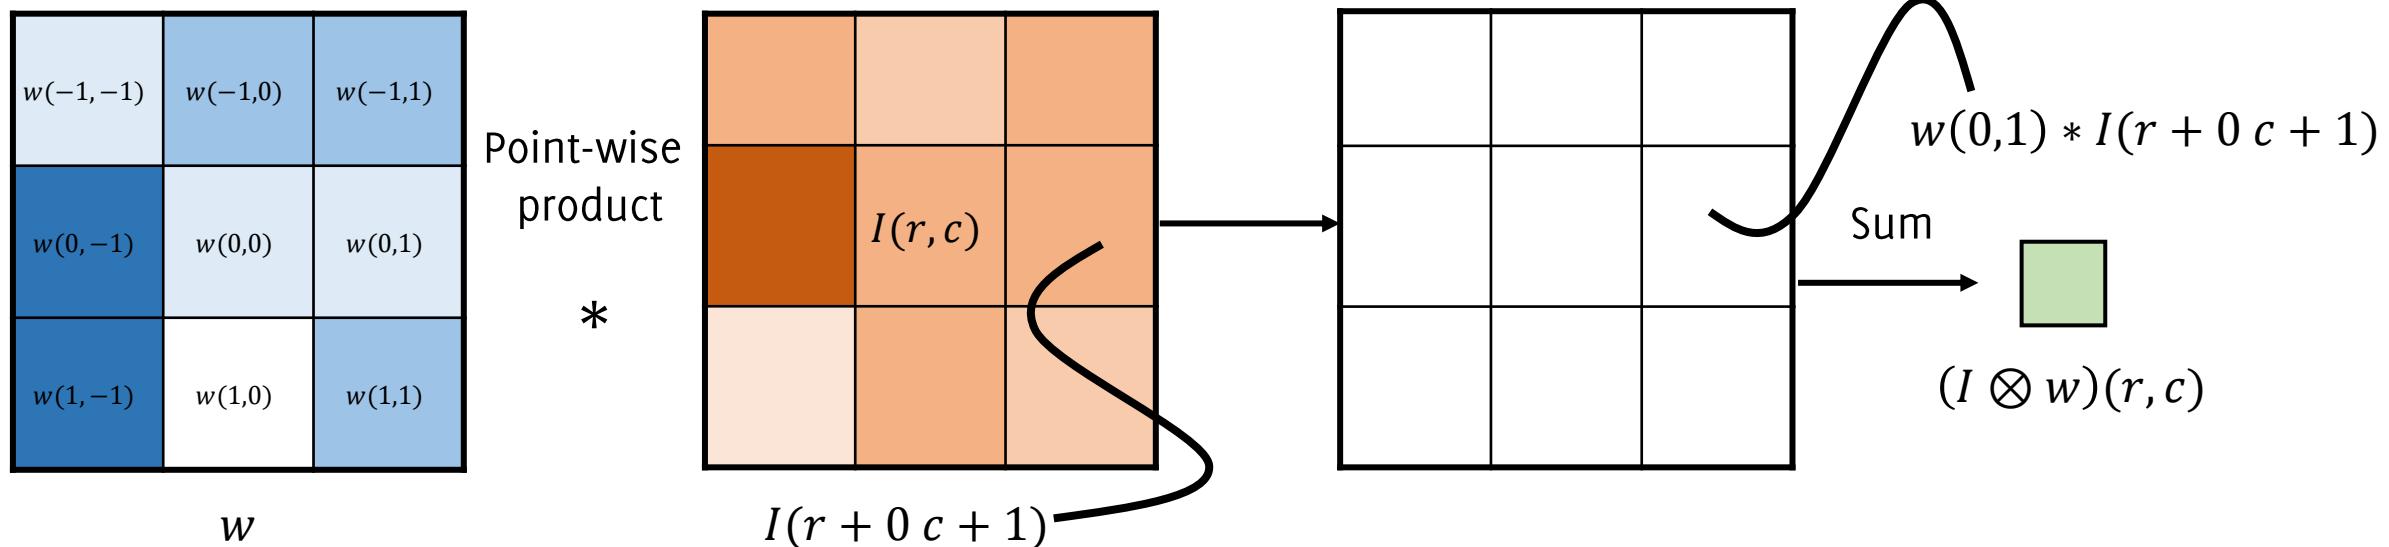


Correlation

The **correlation** among a filter $w = \{w_{ij}\}$ and an image is defined as

$$(I \otimes w)(r, c) = \sum_{u=-L}^L \sum_{v=-L}^L w(u, v) * I(r + u, c + v)$$

where the filter h is of size $(2L + 1) \times (2L + 1)$ and contains the weights defined before as w . The filter w is also sometimes called “kernel”



Correlation

→ example of implementation by nested loops
on image pixels

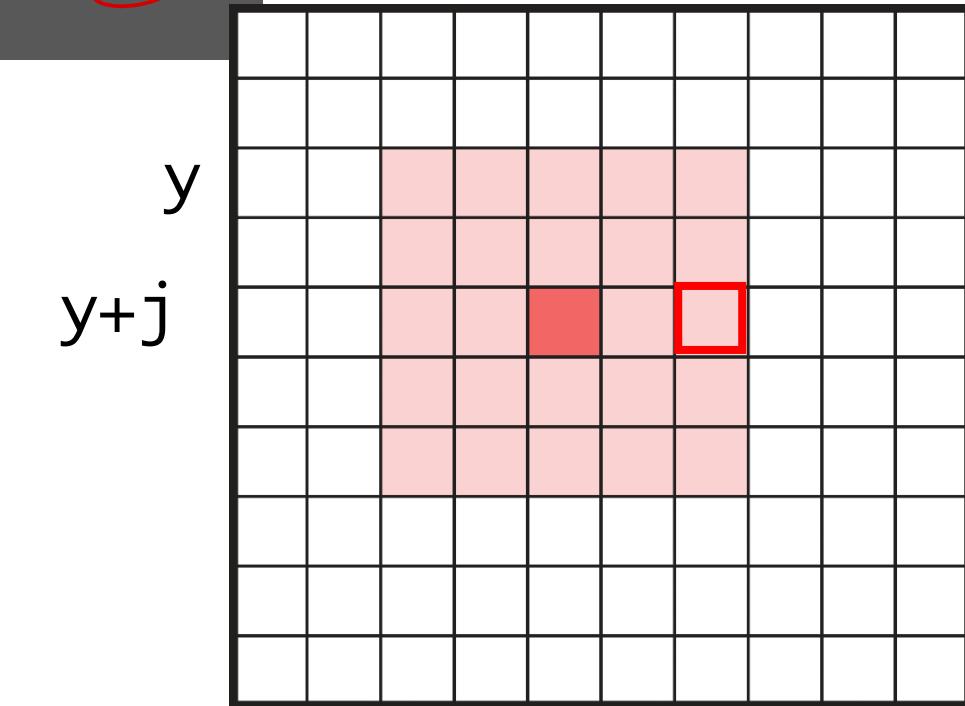
```
acc = 0;  
  
for i in np.arange(template_height)  
    for j in np.arange(template_width)  
        acc = acc + image[y + i, x + j]*template[i,j]  
  
image[x+ template_height//2, y + template_width//2] = acc
```

+ can apply filtering to image pixel

And the two nested loops are repeated for
each and every pixel in the image

expand
on cell
image

cumulate values
x x+i



the weights $w(r, c)$ are defined in standard way
for local average, like $1/|U|$
↓

but it is solution of ZERO ORDER
fit over U...

you can define weights

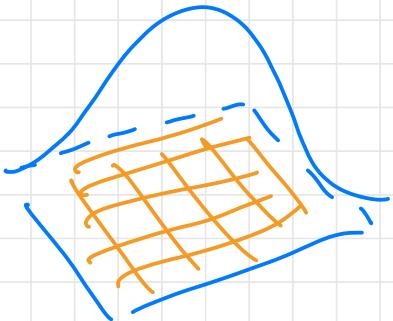
by algebraic manipulation → giving you filters as
First ORDER polynomial
fit returning
center value...



polynomial fitting or
other techniques for

FILTER DESIGN,

typically weighting more the center,
or some form of GAUSSIAN function



here filters are used
for average, in
general for
specific purpose
like edge detection!

Correlation for **BINARY** target matching

IQRM1	DIF1	Det1	#FA1
0.201	0.145	NO	2.000
0.794	0.142	NO	2.000
0.765	0.409	NO	6.000

this is a binary B/W image

\otimes $=$

filter take a portion of the image

Easy to understand with binary images

Target used as a filter

IQRM1	DIF1	Det1	#FA1	
0.201	0.145	NO	2.000	NO
0.794	0.142	NO	2.000	⊗ NO
0.765	0.409	NO	6.000	NO



RESULT..

If filters applied here, BINARY overlapping on
binary_image... \Rightarrow all image below filters

$$Q * \text{FILTER} = Q$$

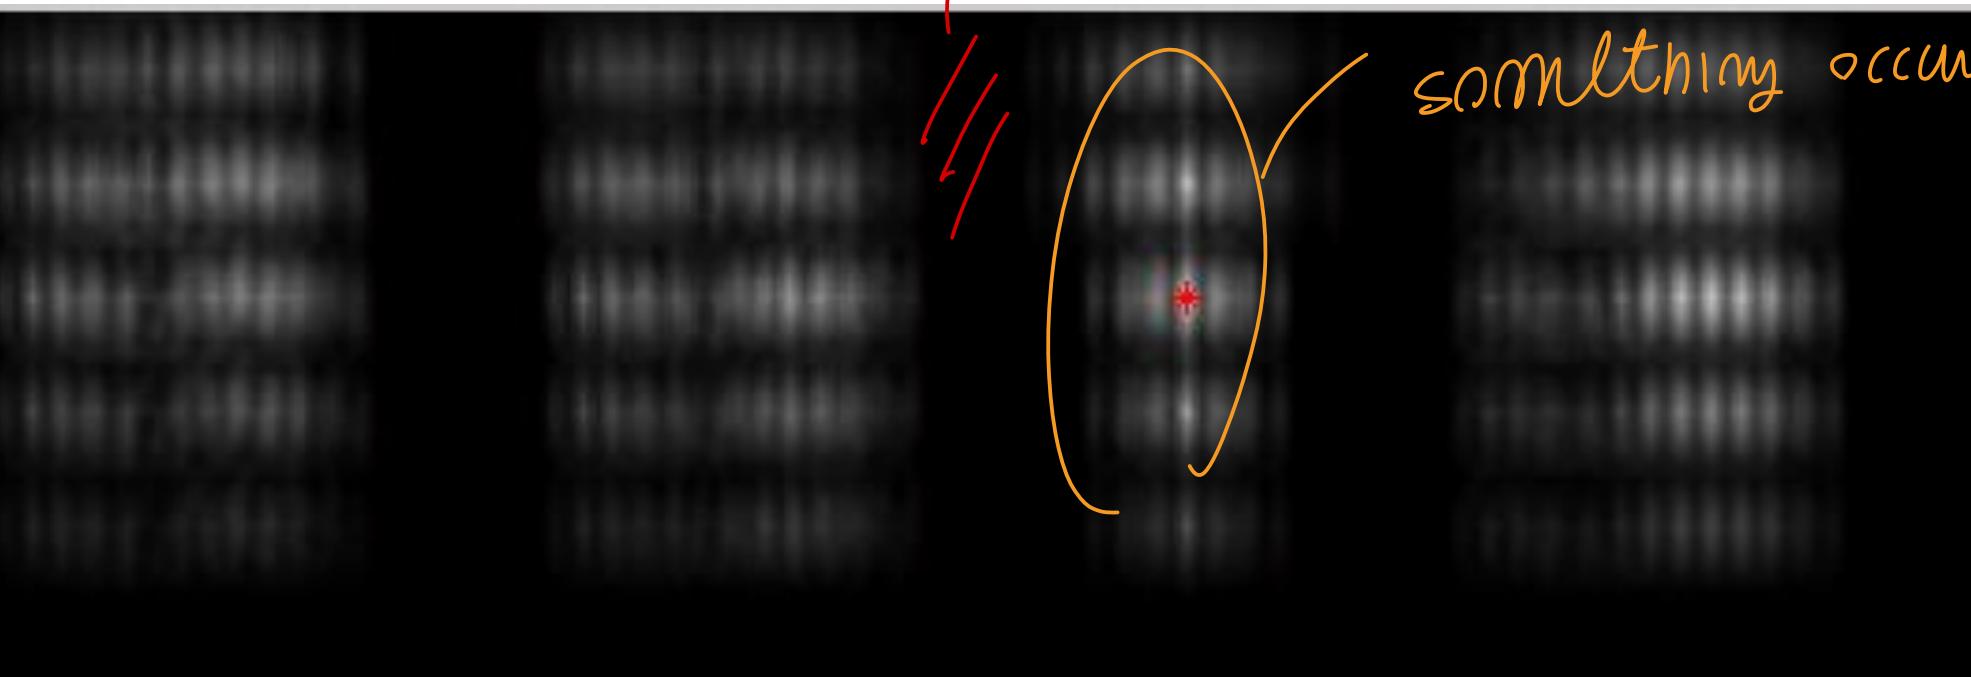
QRM1	DIF1	Det1	#FA1	
0.201	0.145	NO	2.000	NO
0.794	0.142	NO	2.000	NO
0.765	0.409	NO	6.000	NO

CORRELATION filtering, slide the filter through the cell image and compute elementwise product, then sum all cell to get single number

IQRM1	DIF1	Det1	#FA1	
0.201	0.145	NO	2.000	NO
0.794	0.142	NO	2.000	⊗ NO
0.765	0.409	NO	6.000	NO

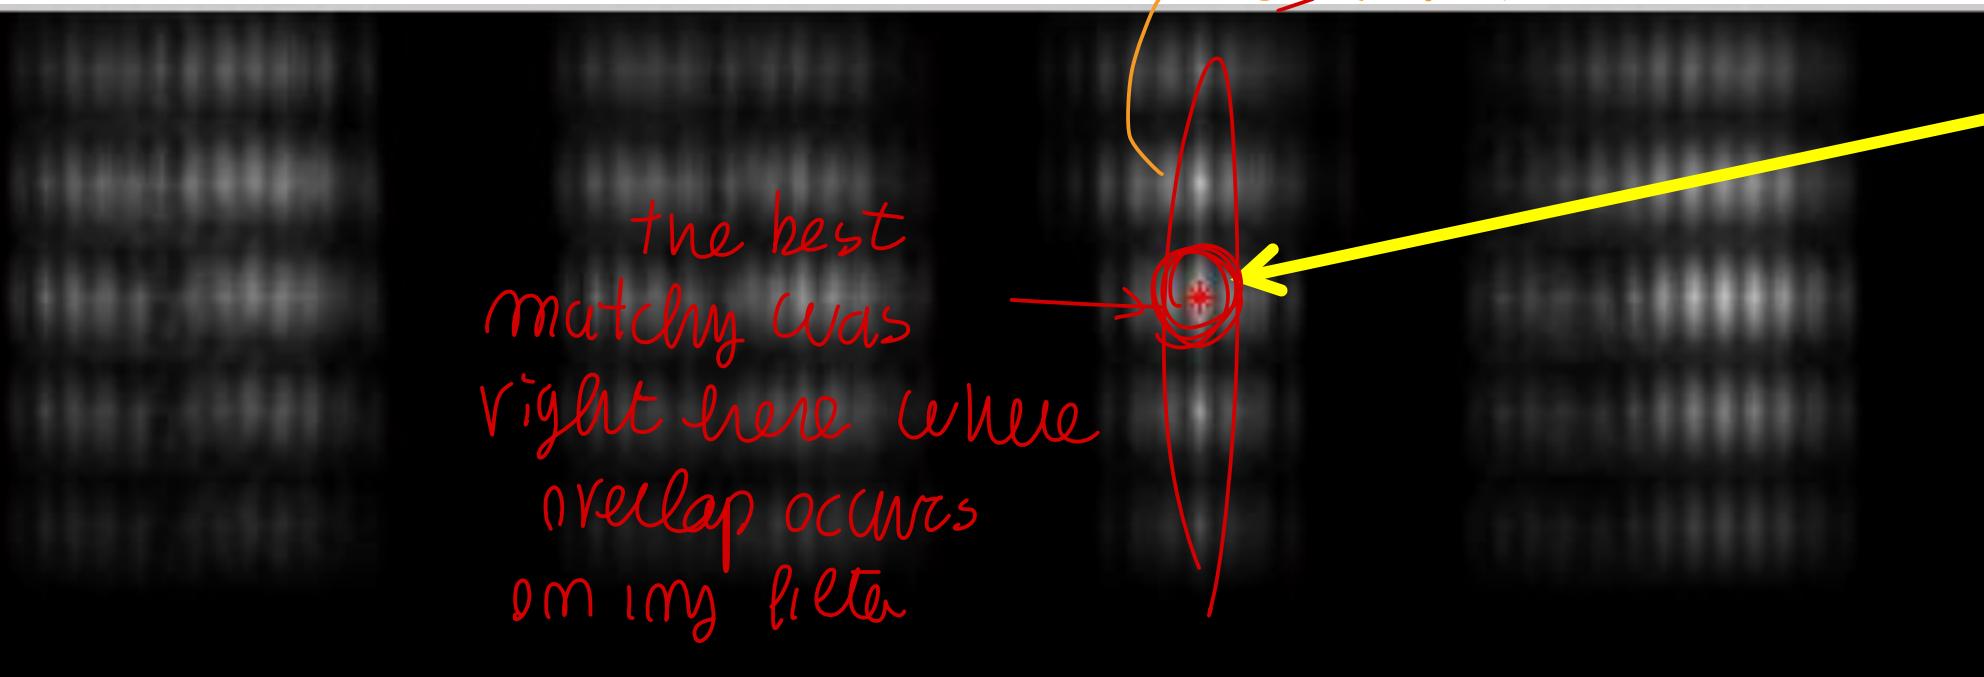
↓(something in the
sum count)

↑ where NO white REGION



something occurs brighter
where we
have
white region
on image

IQRM1	DIF1	Det1	#FA1	
0.201	0.145	NO	2.000	NO
0.794	0.142	NO	2.000	⊗ NO
0.765	0.409	NO	6.000	NO



However...

by flipping B/W image and running

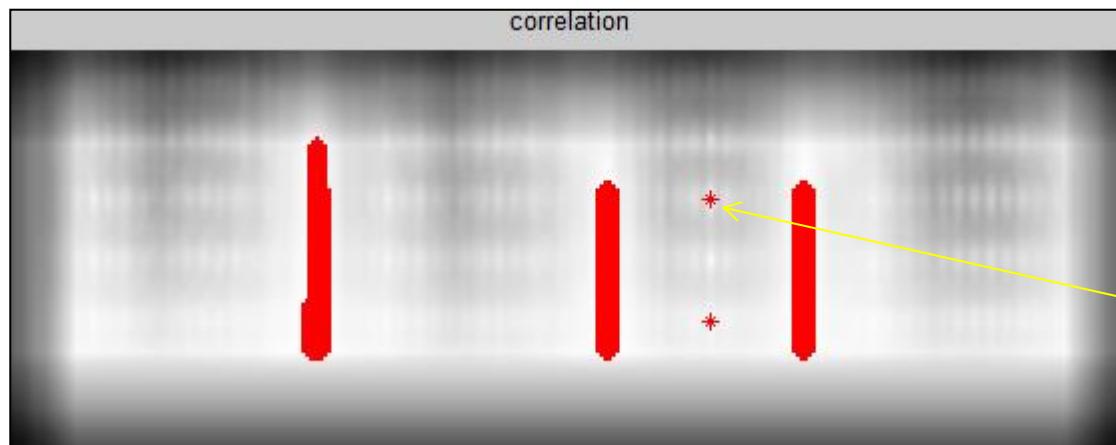
correlation, the max is achieved

in other regions also,

where white
= overlap!

y, original image			
IQRM1	DIF1	Det1	#FA1
0.201	0.145	NO	2.000
0.794	0.142	NO	2.000
0.765	0.409	NO	6.000

also
max where
bright
region
overlapping...



Each point in a white area is
as big as the template
achieve the maximum value
(together with the perfect
match)

However...

To solve this issue, NORMALIZATION helps

y, original image			
IQRM1	DIF1	Det1	#FA1
0.201	0.145	NO	2.000
0.794	0.142	NO	2.000
0.765	0.409	NO	6.000

template
NO
NO
NO

*

=

Normalization is needed when using correlation for template matching!



Each point in a white area is as big as the template achieve the maximum value (together with the perfect match)

Normalized Cross Correlation

A very straightforward approach to template matching

Normalized Cross Correlation

Normalized Cross Correlation $NCC(A, B) \in [-1, 1]$ is defined as

$$NCC(A, B) = \frac{N(A, B)}{\sqrt{N(A, A)N(B, B)}}$$

↑
in filter applied

where

image
neighborhood
region

$$N(A, B) = \iint_{W} (A(x, y) - \bar{A})(B(x, y) - \bar{B}) dx dy$$

take each A and remove $\text{avg } \bar{A}$ in all region
W take each patch of the filter in $dxdy$

and \bar{A} represents the average image value on patch A , similarly \bar{B} . W is the support of A or B .

by \otimes means...

{ brights \otimes you count
black \otimes overlapping

it returns a number in $[-1, 1]..$

Where in our case,

- A is the region in the image,
 - B is the filter
- and they are comparable in size

↳ this actually brings to \otimes means,
this prevent the previous issue of overlapping
parts

Normalized Cross Correlation

Normalized Cross Correlation $NCC(A, B) \in [-1, 1]$ is defined as

$$NCC(A, B) = \frac{N(A, B)}{\sqrt{N(A, A)N(B, B)}}$$

↑ normalized

where

$$N(A, B) = \iint_W (A(x, y) - \bar{A})(B(x, y) - \bar{B}) dx dy$$

subtract avg and compute product...

and \bar{A} represents the average image value on patch A , similarly \bar{B} . W is the support of A or B .

Where in our case,

- A is the region in the image,
 - B is the filter
- and they are comparable in size

Normalized Cross Correlation

Remarks:

- NCC yields a measure in the range $[-1,1]$,
- NCC is invariant to changes in the average intensity.
- While this seems quite computationally demanding, there exists fast implementations where local averages are computed by running sums: the integral image

← in convolution usually we just subtract avg...

While in image you subtract mean
proper normalization
 $\hookrightarrow \frac{\bar{I}(r,c)}{N(r,c)}$

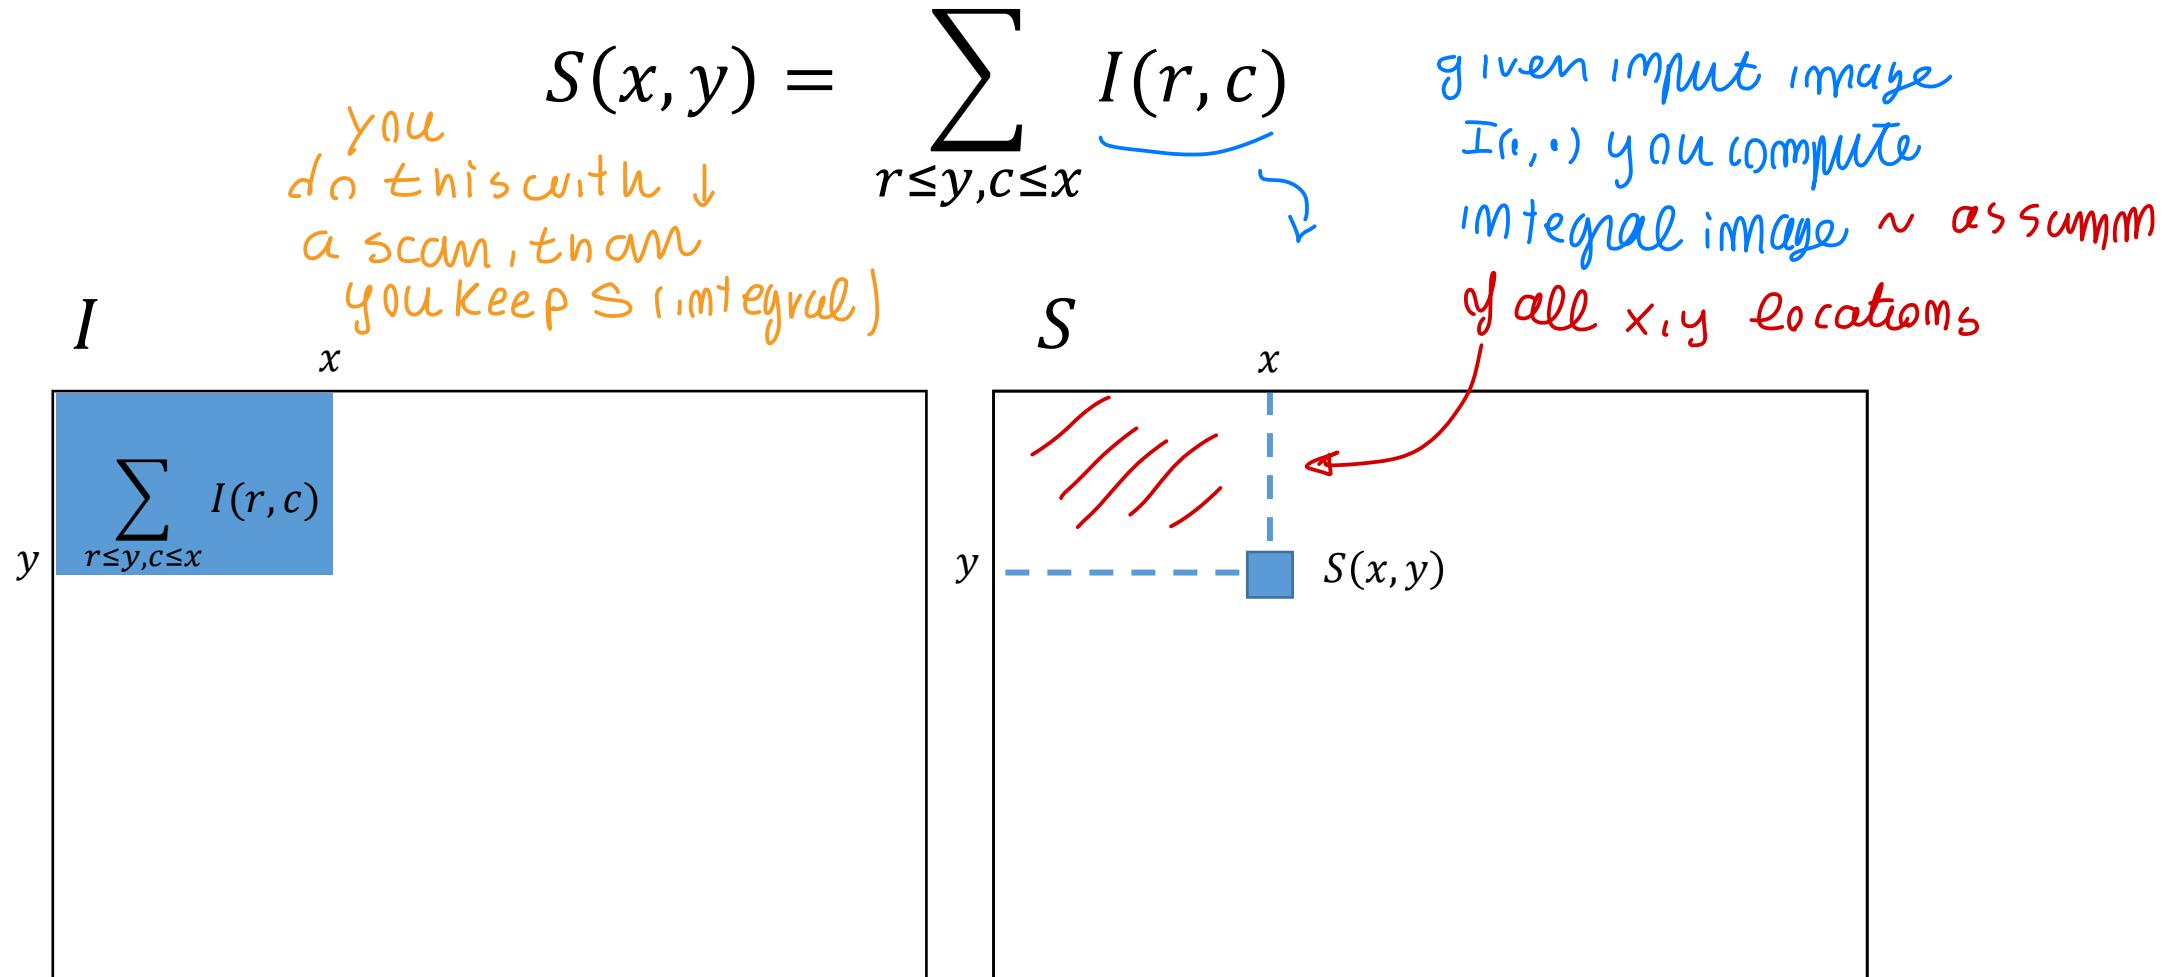
+ to do it in efficient way \Rightarrow
to obtain a
commonly to do
for each point!

Integral Image

→ To normalize in a simpler way...
doing it efficiently:

(really useful
for descriptors)

The integral image S is defined from an image I as follows



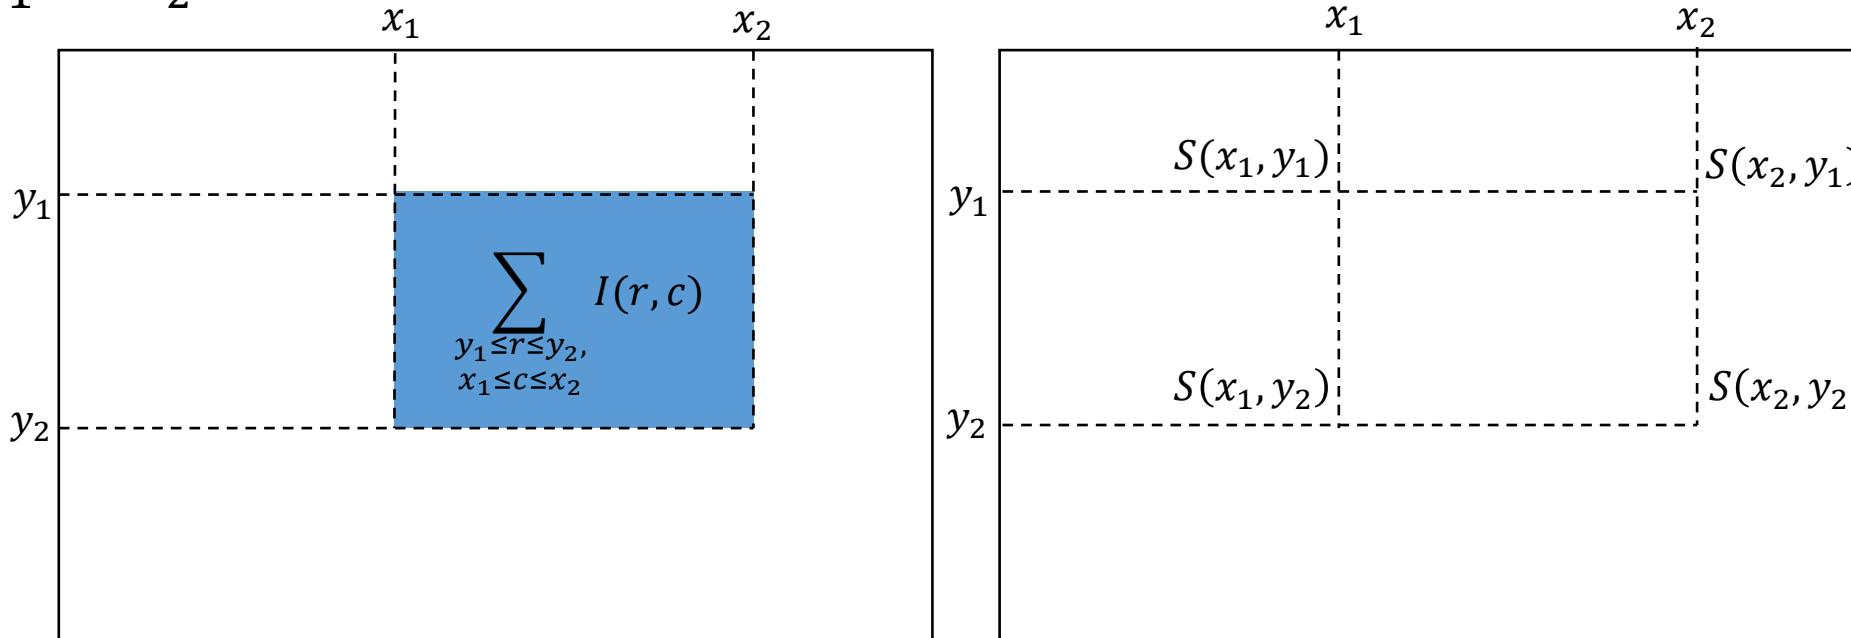
Using the Integral Image

Given the integral image $S(x, y)$
to compute the AREA of any square:

The integral image allows fast computation of the sum (average) of any rectangular region in the image

$$\sum_{\substack{y_1 \leq r \leq y_2, \\ x_1 \leq c \leq x_2}} I(r, c) = S(x_2, y_2) - S(x_2, y_1) - S(x_1, y_2) + S(x_1, y_1)$$

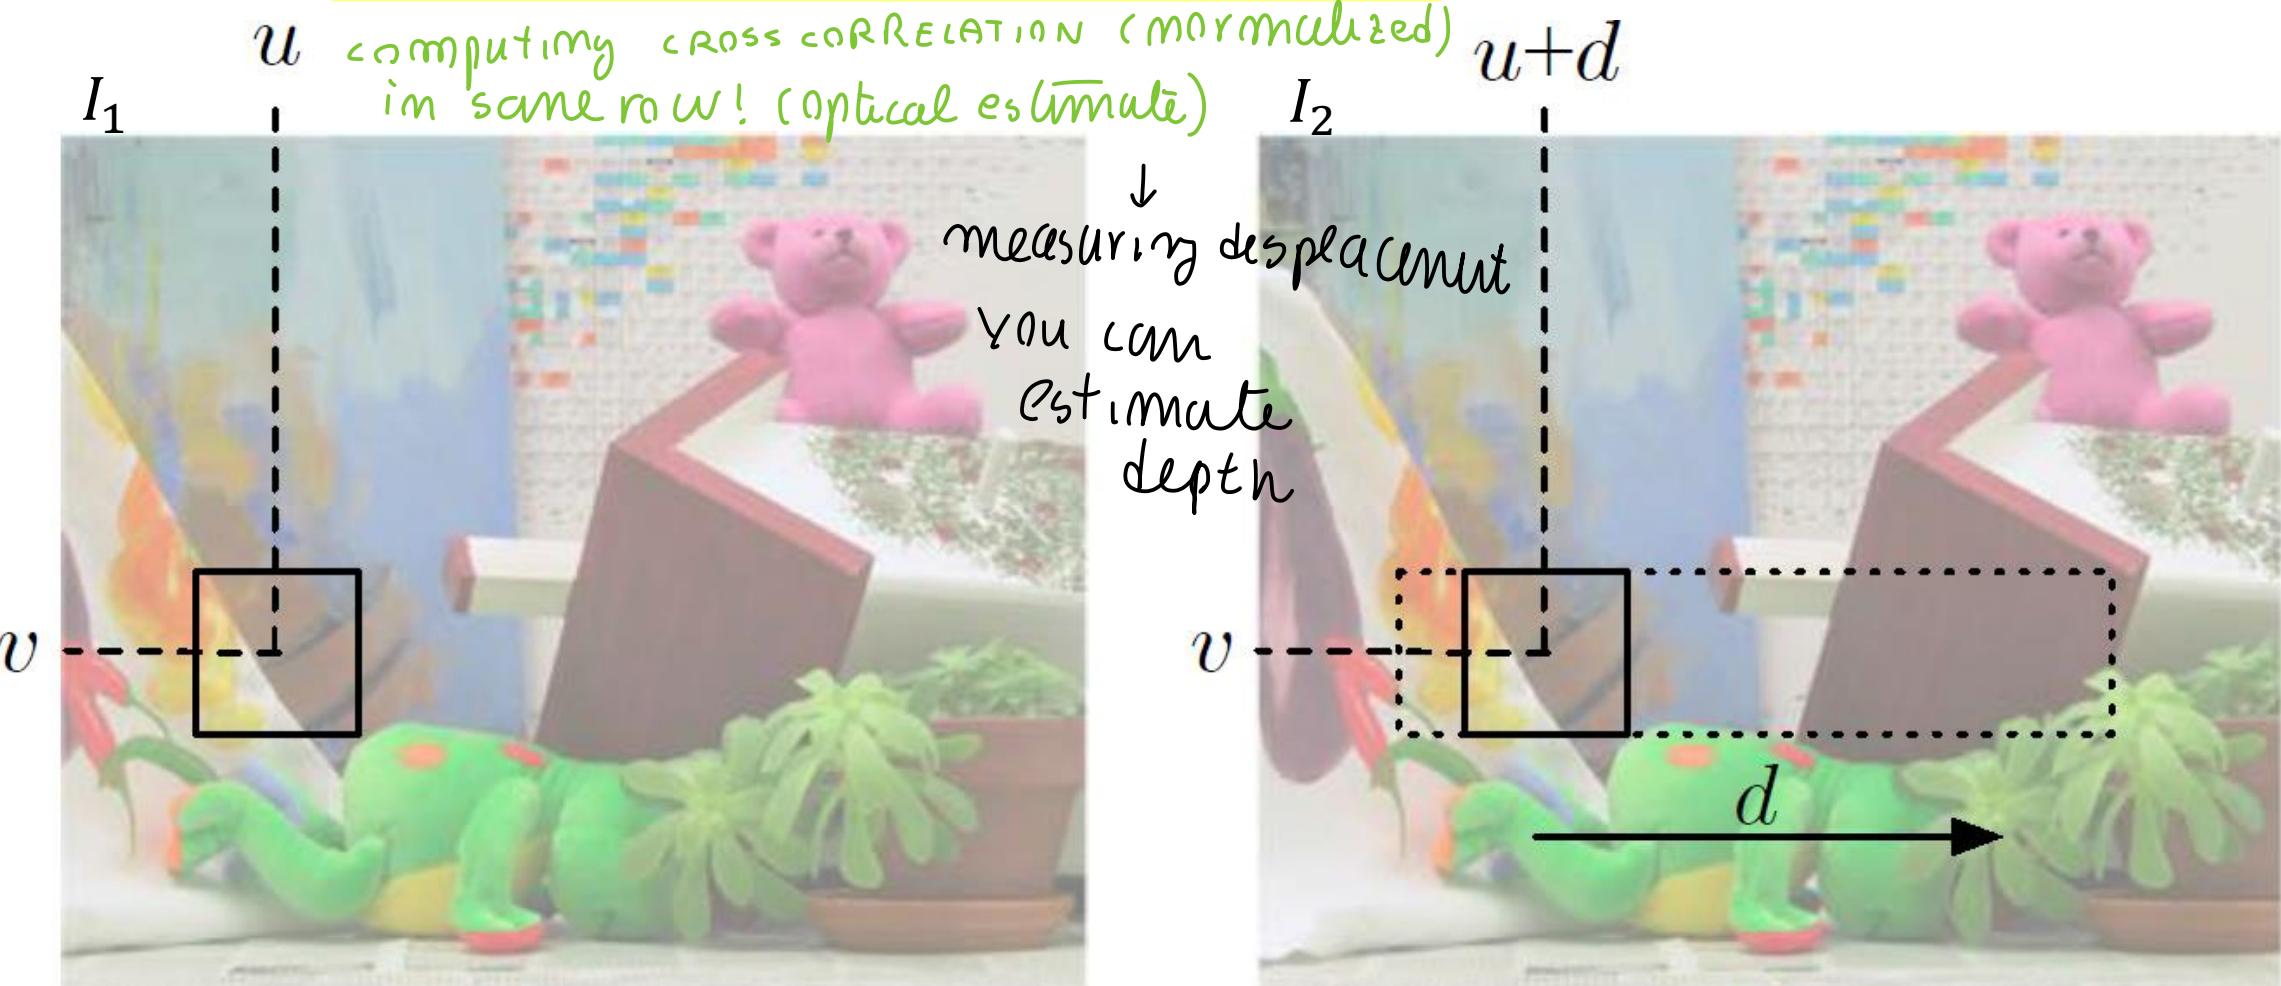
by composition of different areas



Disparity Map Estimation

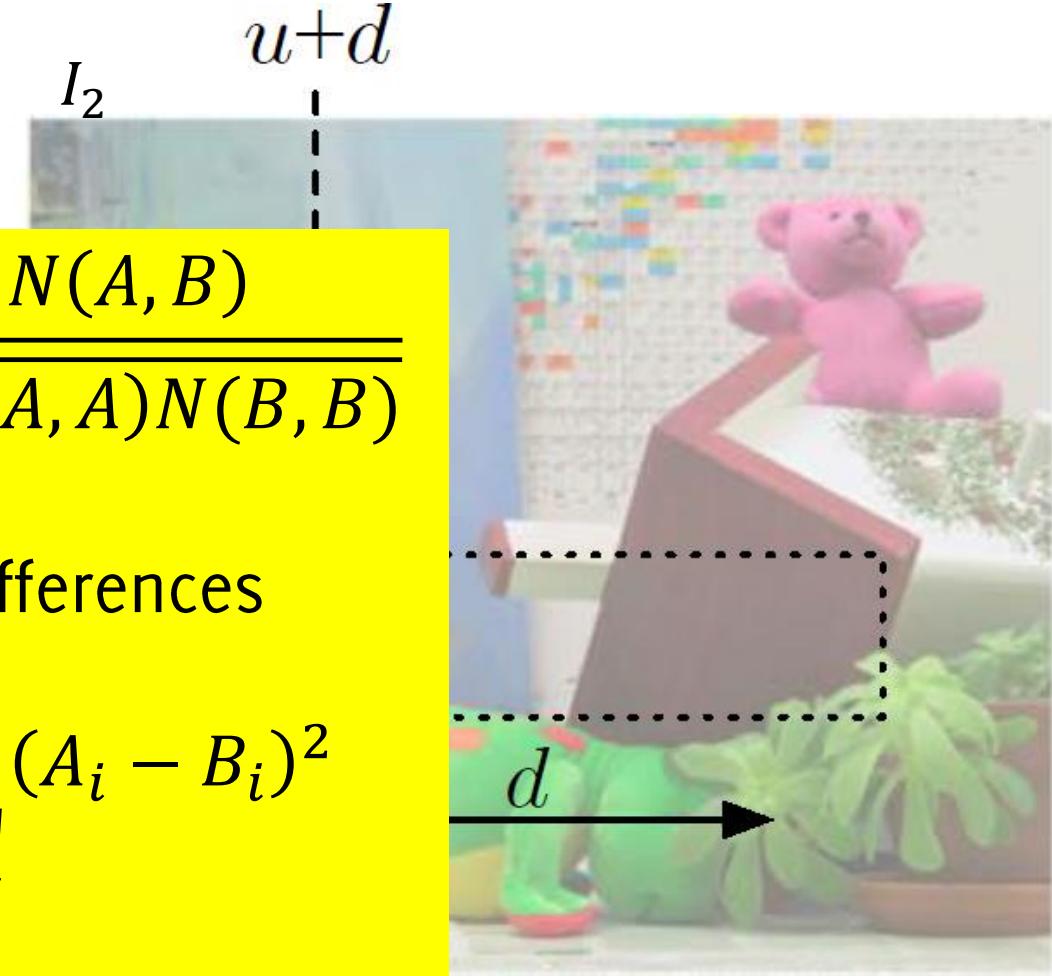
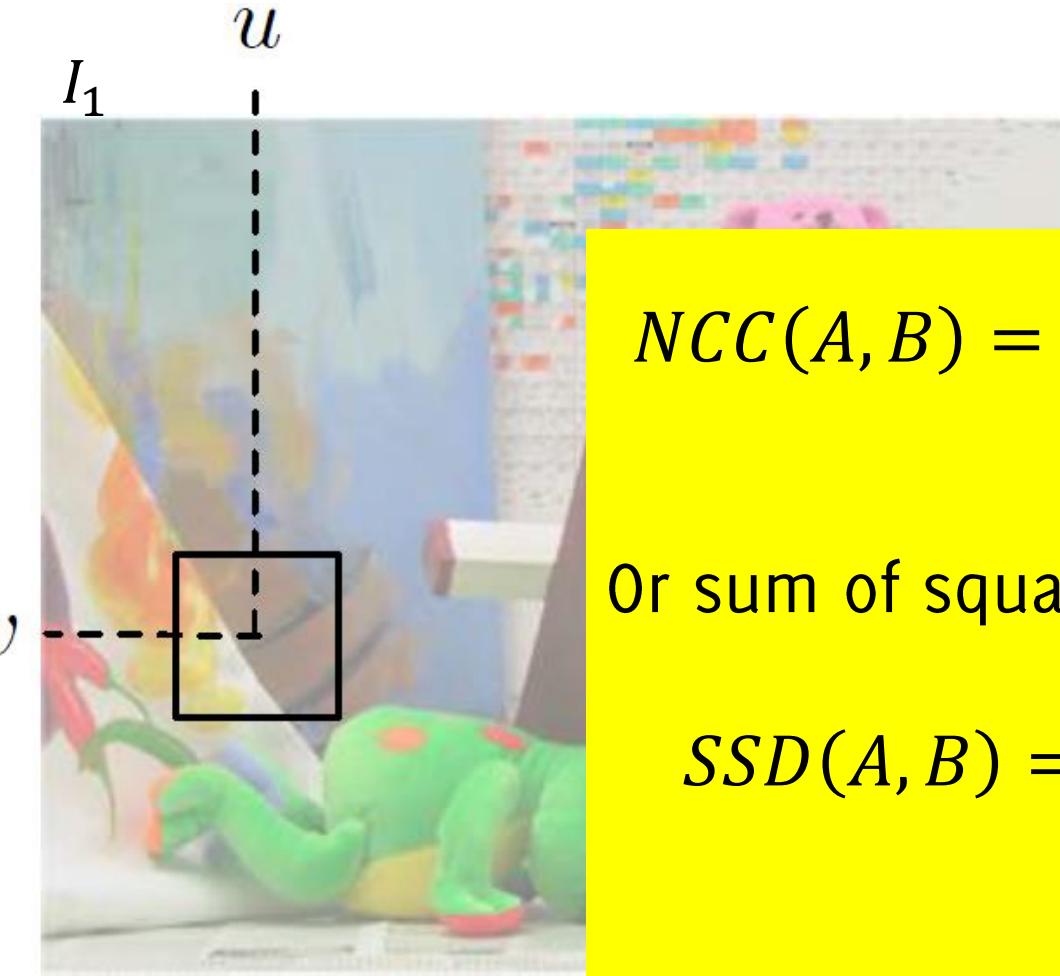
↑ that integral image is used for example:

There are different measures to compare a patch in I_1 with all the candidate matches in I_2



Disparity Map Estimation

There are different measures to compare a patch in I_1 with all the candidate matches in I_2



$$NCC(A, B) = \frac{N(A, B)}{\sqrt{N(A, A)N(B, B)}}$$

Or sum of square differences

$$SSD(A, B) = \sum_{i=1}^N (A_i - B_i)^2$$

Stereo Pairs <http://vision.middlebury.edu/stereo/data/>

(*)
things
moving
more cue
closer,
what
move less
→ further
↓
depth
estimate



you have two
views of
the same
scene...

↳ any point
lie on same
horizontal line

↓
by measuring
two pixels how
they displace,
you can estimate
depth

(*)

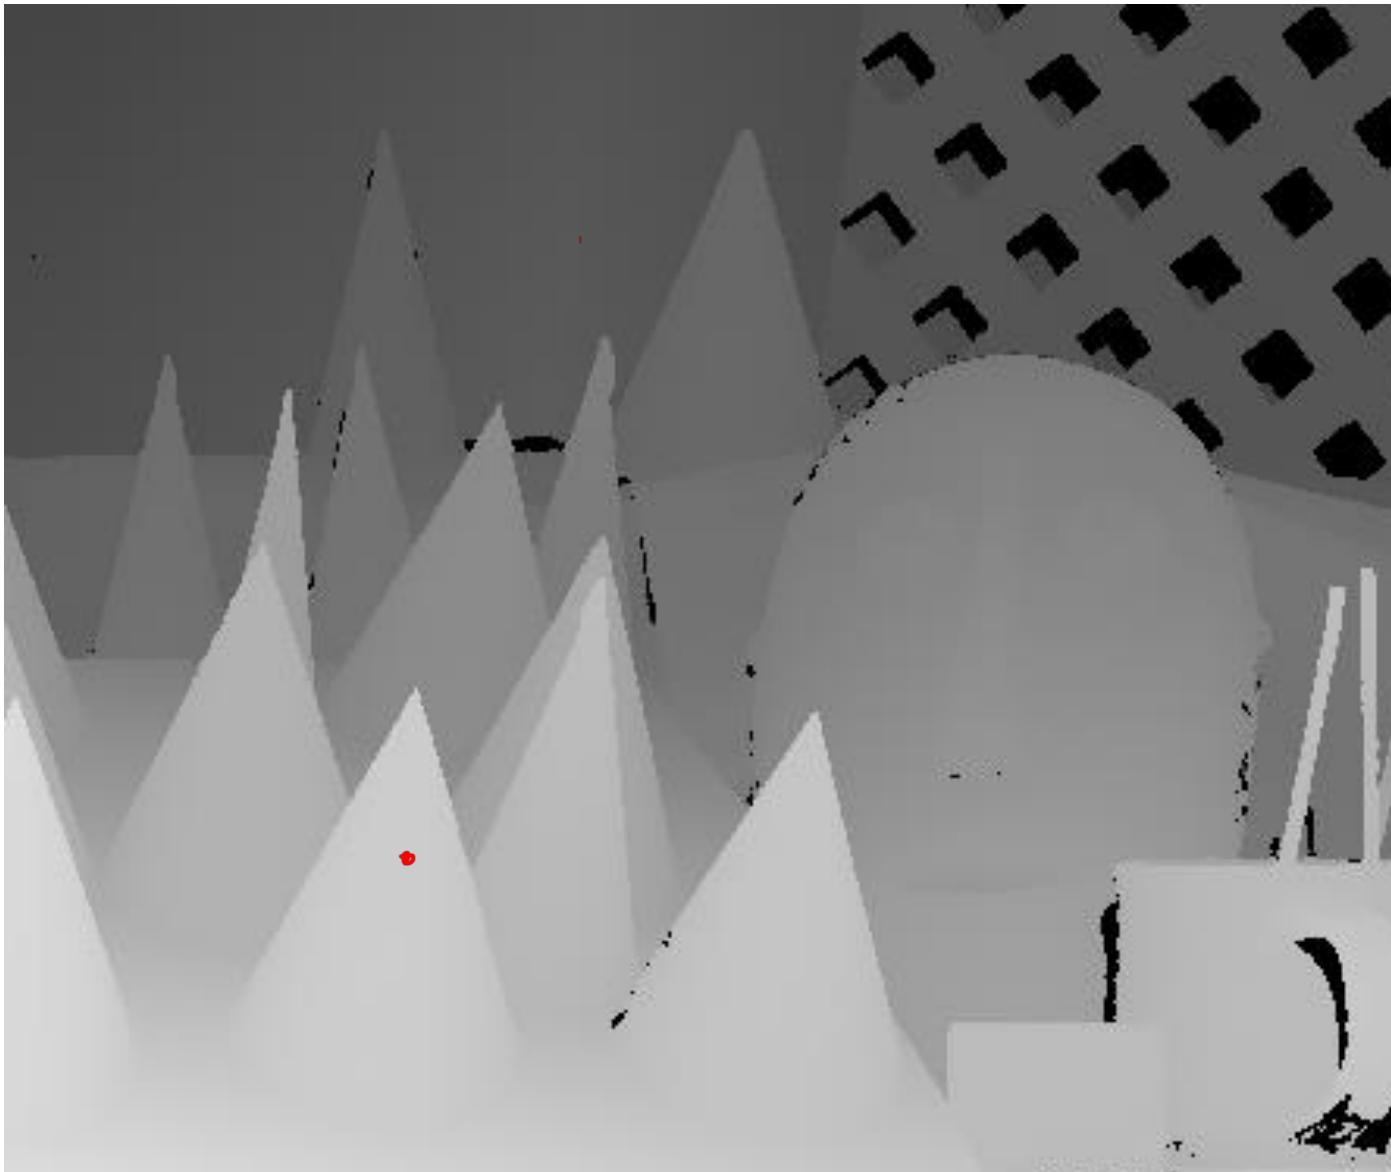
Stereo Pairs <http://vision.middlebury.edu/stereo/data/>



Giacomo Boracchi

Stereo Pairs <http://vision.middlebury.edu/stereo/data/>

by computing
(NORMALIZED) CROSS
CORRELATION
of each
region in
one image



Convolution

Correlation and Convolution



The **correlation** among a filter w and an image is defined as

$$(I \otimes w)(r, c) = \sum_{u=-L}^L \sum_{v=-L}^L w(u, v) * I(r + u, c + v)$$

where the filter w is of size $(2L + 1) \times (2L + 1)$

The **convolution** among a filter w and an image is defined as

$$(I \circledast w)(r, c) = \sum_{u=-L}^L \sum_{v=-L}^L w(u, v) * I(r - u, c - v)$$

where the filter w is of size $(2L + 1) \times (2L + 1)$

we need (-) sign, convolution
almost never with nested
for loops... BUT
with Fourier
transform...

↑
simpler to
compute in
fourier domain

in fourier
transform,
convolution is
just product

There is just a swap in the filter before computing correlation!

Correlation and Convolution

in deep learning, the filter is learned,
so you don't care about the sign

↑ when you design it,
you flip the sign

The **correlation** among a filter w and an image is defined as

$$(I \otimes w)(r, c) = \sum_{u=-L}^L \sum_{v=-L}^L w(u, v) * I(r + u, c + v)$$

where the filter w is of size $(2L + 1) \times (2L + 1)$

The **convolution** among a filter w and an image is defined as

$$(I \circledast w)(r, c) = \sum_{u=-L}^L \sum_{v=-L}^L w(u, v) * I(r - u, c - v)$$

where the filter w is of size $(2L + 1) \times (2L + 1)$

There is just a swap in the filter before computing correlation!

Convolution - and filter flip

morning the
sig in that way...

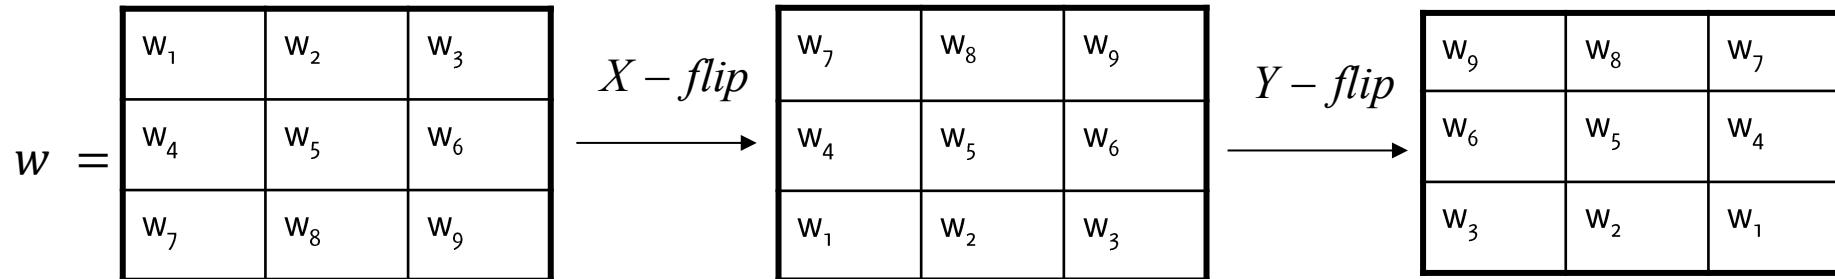
move (-) to w

Let I, w be two discrete 2D signals of $(2L + 1) \times (2L + 1)$

$$G(r, c) = (I \odot w)(r, c) = \sum_{u=-L}^L \sum_{v=-L}^L I(r + u, c + v) w(-u, -v)$$

like flipping
the filter
and use as
before

using
filter
like
correlation
to look
same
text exactly



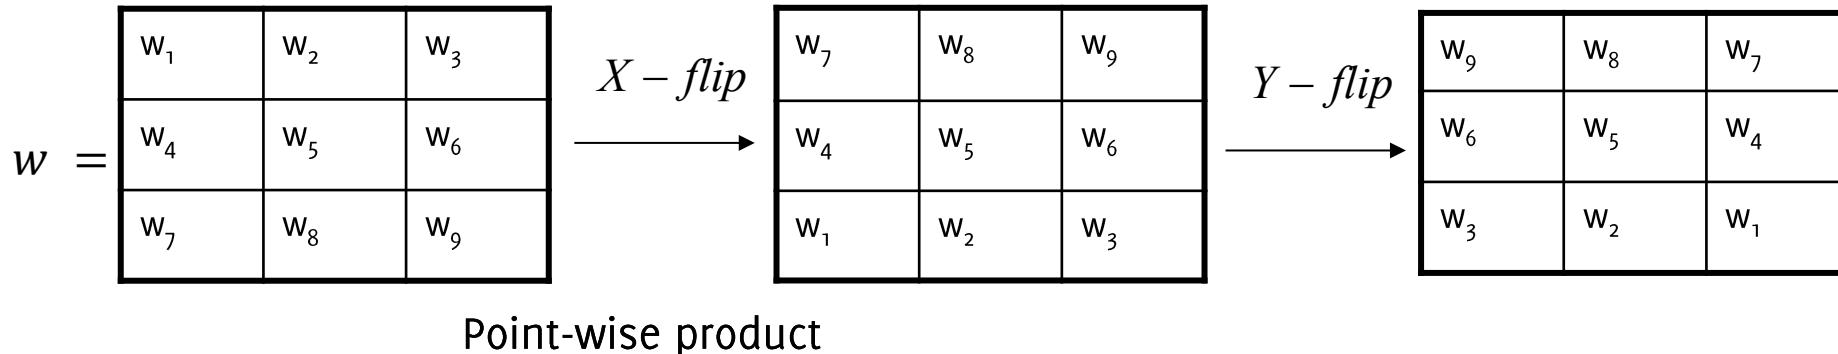
In this particular case $L = 1$ and both the image and
the filter have size 3×3

The convolution is evaluated at $(r, c) = (0, 0)$

Convolution – and filter flip

Let I, h be two discrete 2D signals of $(2L + 1) \times (2L + 1)$

$$G(r, c) = (I \odot w)(r, c) = \sum_{u=-L}^L \sum_{v=-L}^L I(r + u, c + v)w(-u, -v)$$

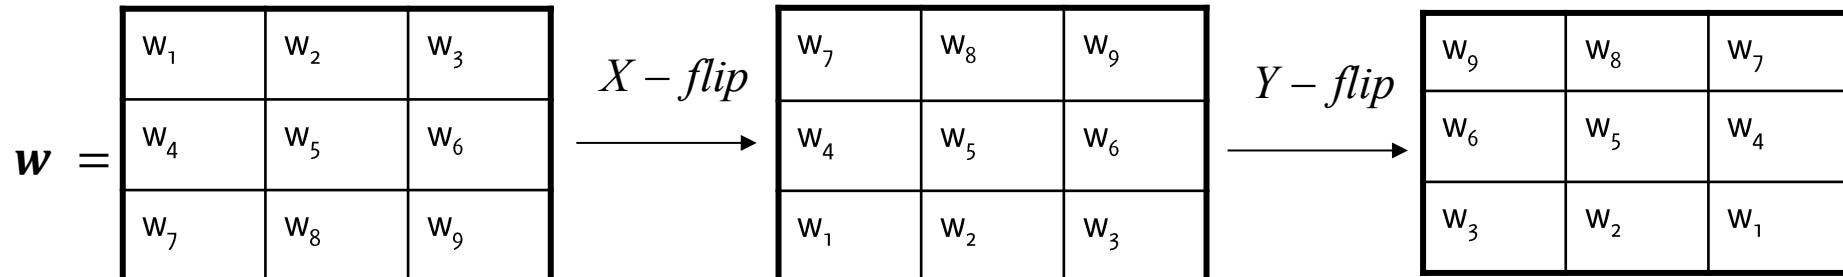


$$\begin{array}{c} \begin{array}{|c|c|c|} \hline w_9 & w_8 & w_7 \\ \hline w_6 & w_5 & w_4 \\ \hline w_3 & w_2 & w_1 \\ \hline \end{array} \cdot * \begin{array}{|c|c|c|} \hline l_1 & l_2 & l_3 \\ \hline l_4 & l_5 & l_6 \\ \hline l_7 & l_8 & l_9 \\ \hline \end{array} \rightarrow \begin{array}{|c|c|c|} \hline w_9l_1 & w_8l_2 & w_7l_3 \\ \hline w_6l_4 & w_5l_5 & w_4l_6 \\ \hline w_3l_7 & w_2l_8 & w_1l_9 \\ \hline \end{array} \end{array}$$

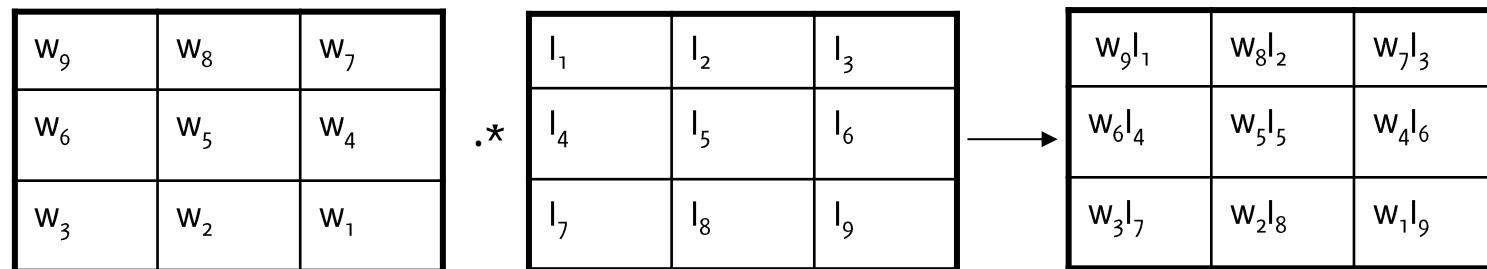
Convolution

Let I, \mathbf{w} be two discrete 2D signals of $(2L + 1) \times (2L + 1)$

$$G(r, c) = (I \circledast \mathbf{w})(r, c) = \sum_{u=-L}^L \sum_{v=-L}^L I(r + u, c + v)w(-u, -v)$$



Point-wise product



$$G_5 = w_9I_1 + w_8I_2 + w_7I_3 + w_6I_4 + w_5I_5 + w_5I_6 + w_3I_7 + w_2I_8 + w_1I_9$$

Question

The filter (a.k.a. the kernel) yields the coefficients used to compute the linear combination of the input to obtain the output

1	3	0
2	10	2
4	1	1

*

1	0	-1
1	0.1	-1
1	0	-1

=

	?	

Image

Kernel

Filter Output

Let's have a look at 1D
convolution

Let's have a look at 1D Convolution

Let us consider a 1d signal y and a filter \mathbf{w} .

- Their convolution is also a signal $z = y \otimes \mathbf{w}$.
- For continuous-domain 1D signals and filters

$$z(\tau) = (y \otimes \mathbf{w})(\tau) = \int_{\mathbb{R}} y(t) \mathbf{w}(\tau - t) dt$$

that is equivalent to

$$z(\tau) = (h \otimes \mathbf{w})(\tau) = \int_{\mathbb{R}} y(\tau - t) \mathbf{w}(t) dt$$

- For discrete signals and filters

$$z(n) = (y \otimes \mathbf{w})(n) = \sum_{m=-L}^L y(n-m) \mathbf{w}(m)$$

where the filter has $(2L + 1)$ samples

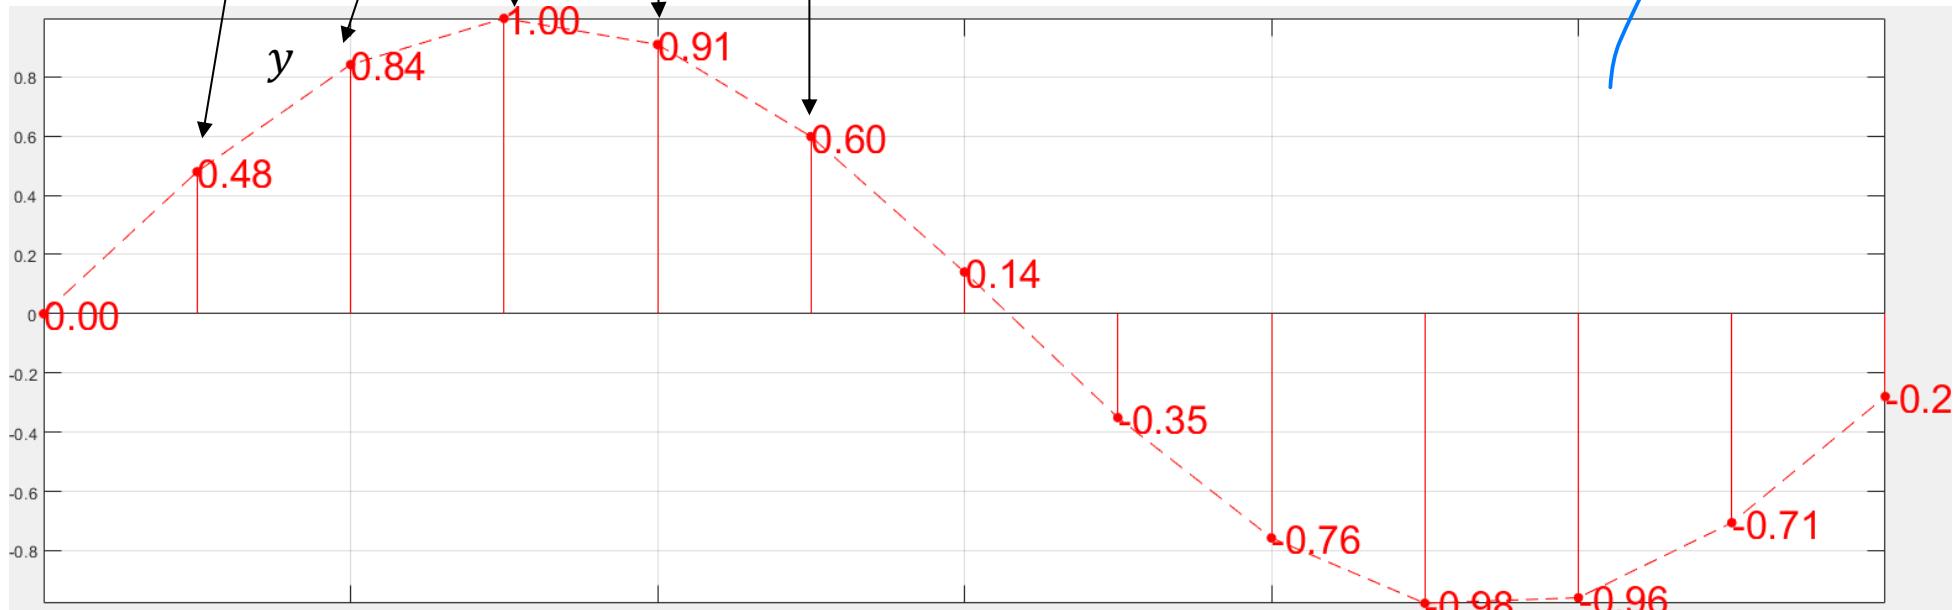
1D Convolution - example

How it behaves...

$$z(n) = (y \otimes w)(n) = \sum_{m=-L}^L y(n-m)w(m)$$

$$y = \sin(x), w = \left[\frac{1}{5}, \frac{1}{5}, \frac{1}{5}, \frac{1}{5}, \frac{1}{5} \right], L = 2$$

$$\left[\frac{1}{5}, \frac{1}{5}, \frac{1}{5}, \frac{1}{5}, \frac{1}{5} \right]$$

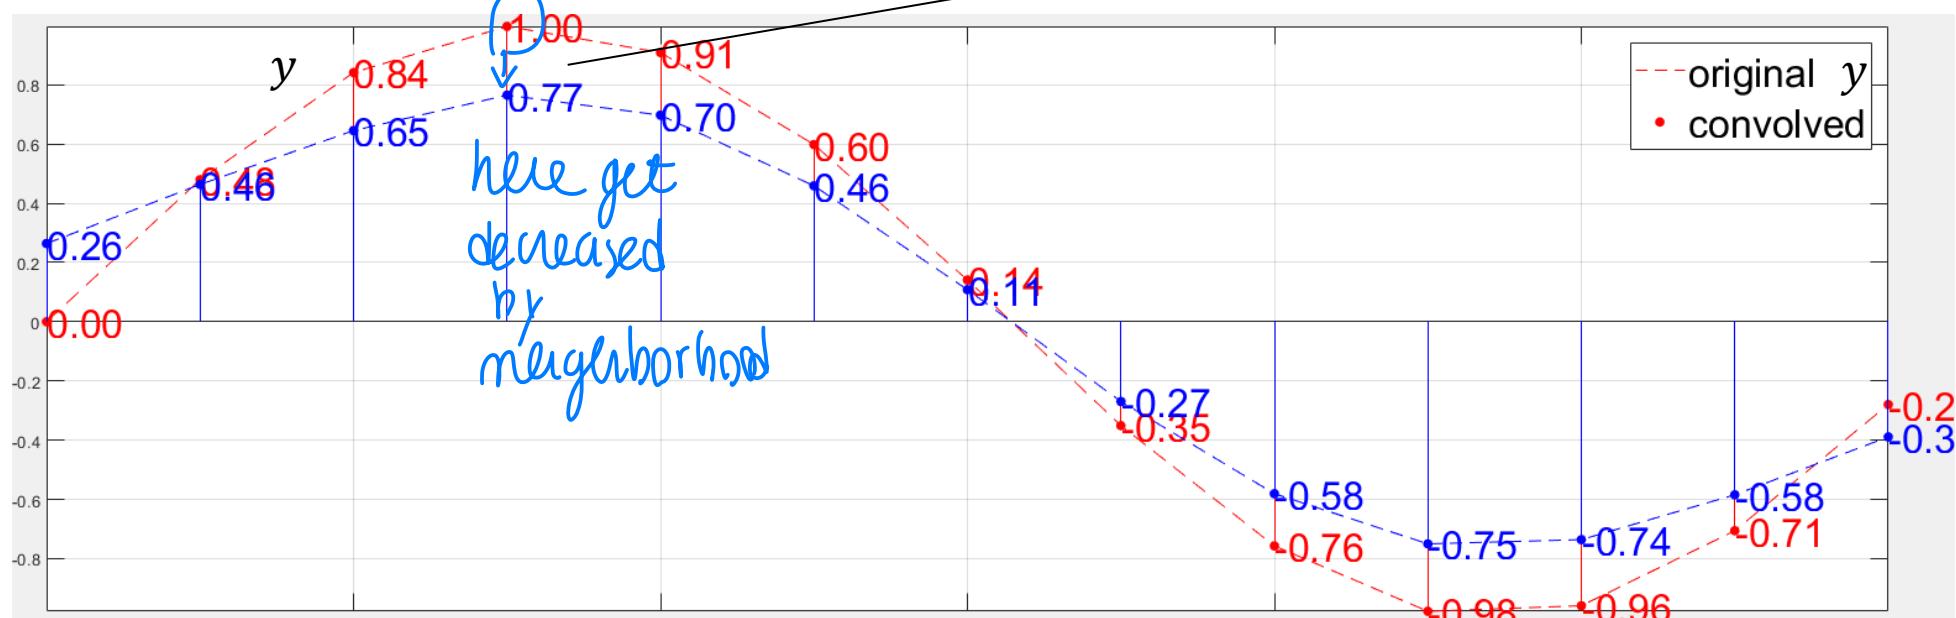


like a
sampled
sine signal

1D Convolution - example

filter is
like the average
over a region
of 5... each
pixel is replaced
by avg of 5
neighborhood

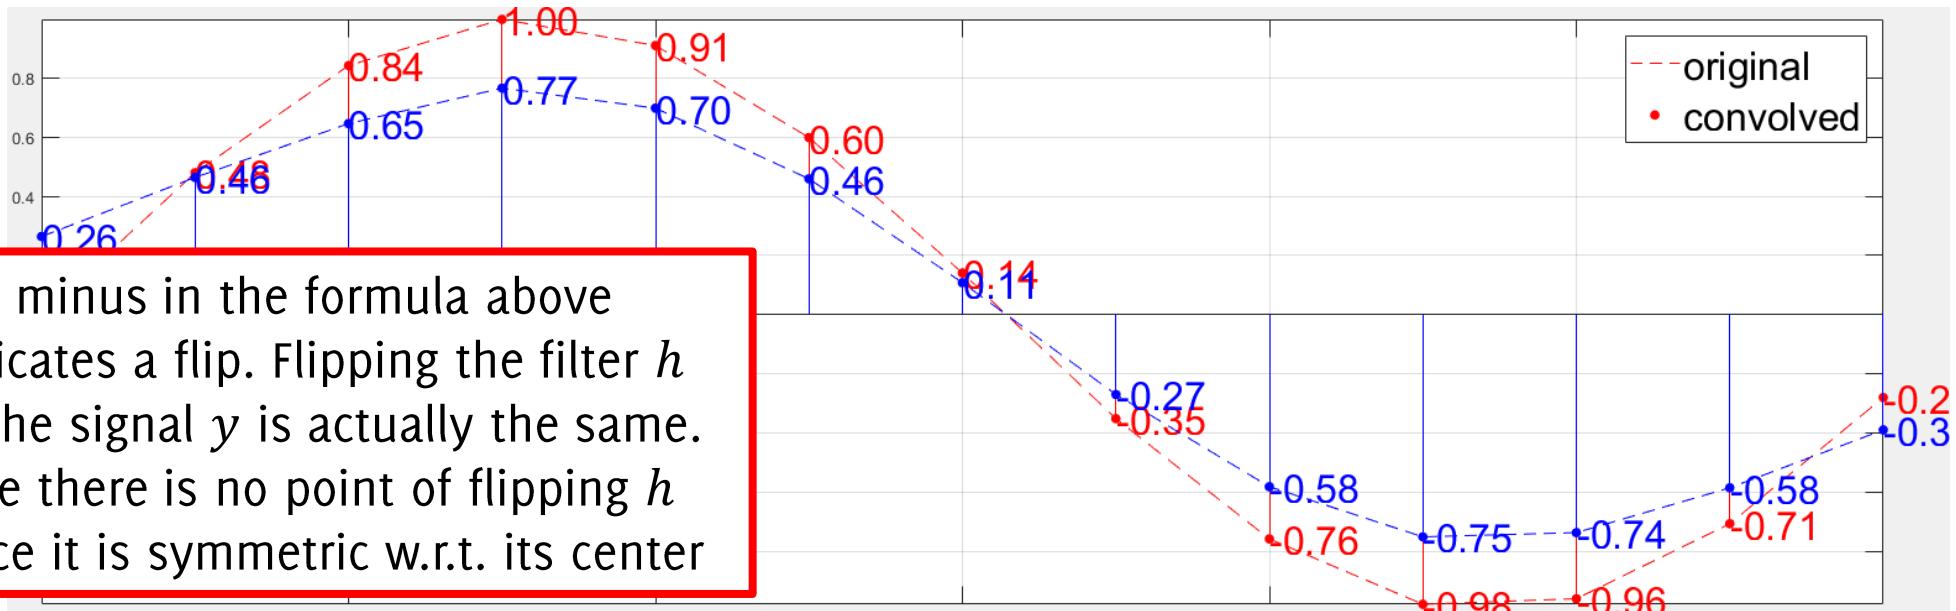
$$z(n) = (y \otimes w)(n) = \sum_{m=-L}^L y(n-m)w(m)$$
$$y = \sin(x), w = \left[\frac{1}{5}, \frac{1}{5}, \frac{1}{5}, \frac{1}{5}, \frac{1}{5} \right], L = 2$$
$$0.766 \approx \frac{1}{5} * 0.48 + \frac{1}{5} * 0.84 + \frac{1}{5} * 1 + \frac{1}{5} * 0.91 + \frac{1}{5} * 0.60$$



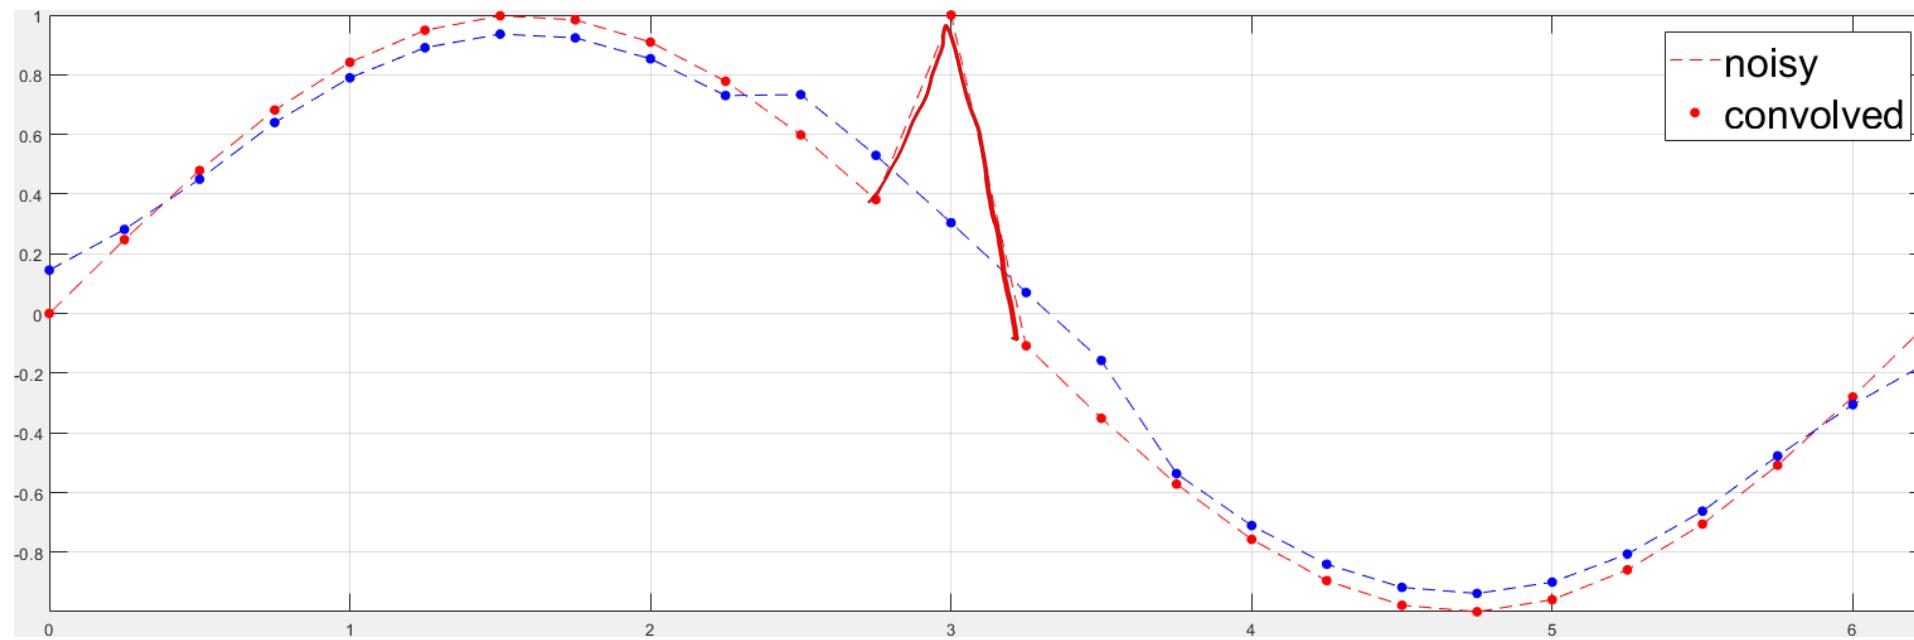
1D Convolution - example

$$\begin{aligned} z(n) = (y \otimes w)(n) &= \sum_{m=-L}^L y(n-m)w(m) \\ &= \sum_{m=-L}^L y(n+m)w(-m) \end{aligned}$$

↙ flip



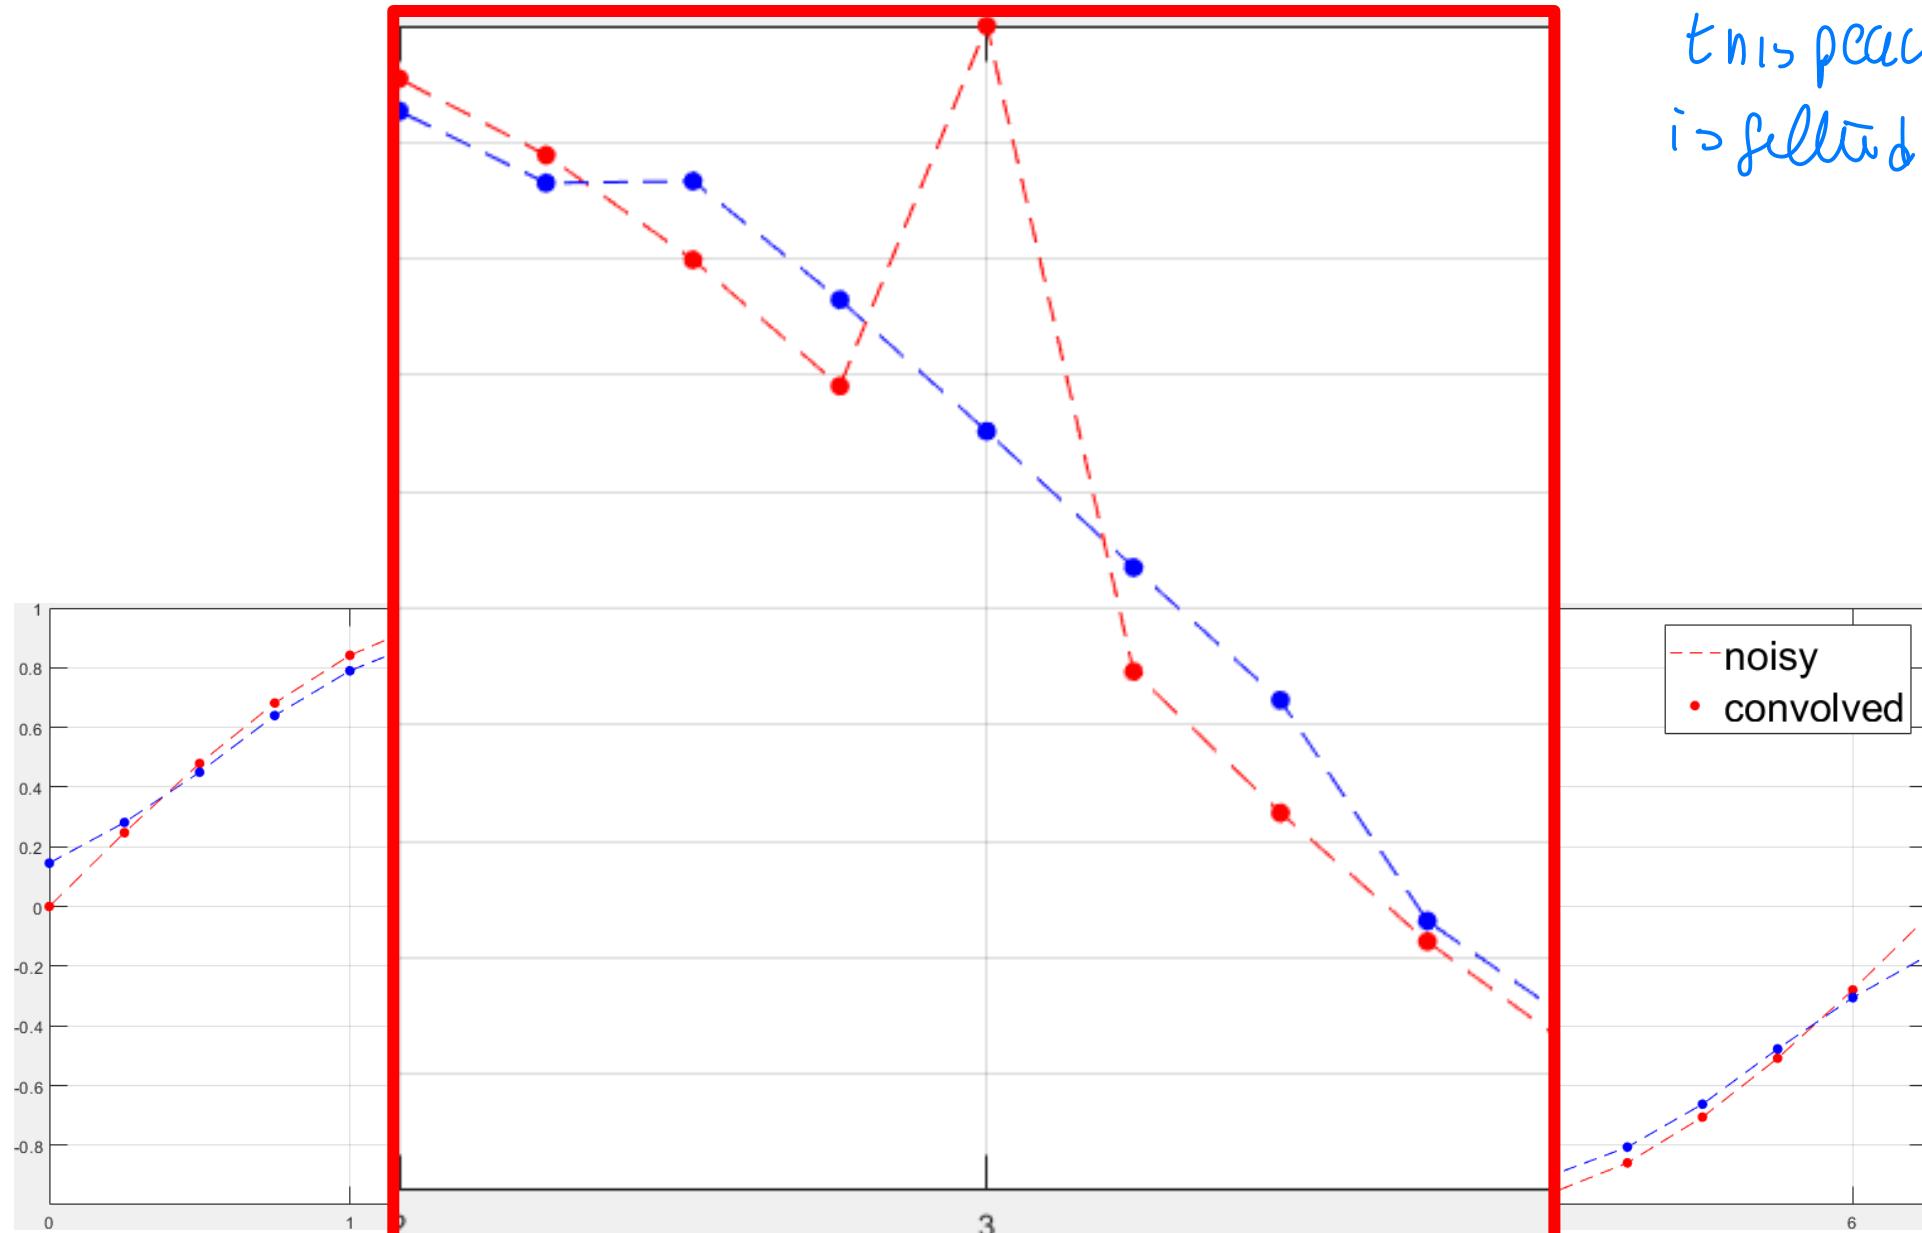
What about an impulse?



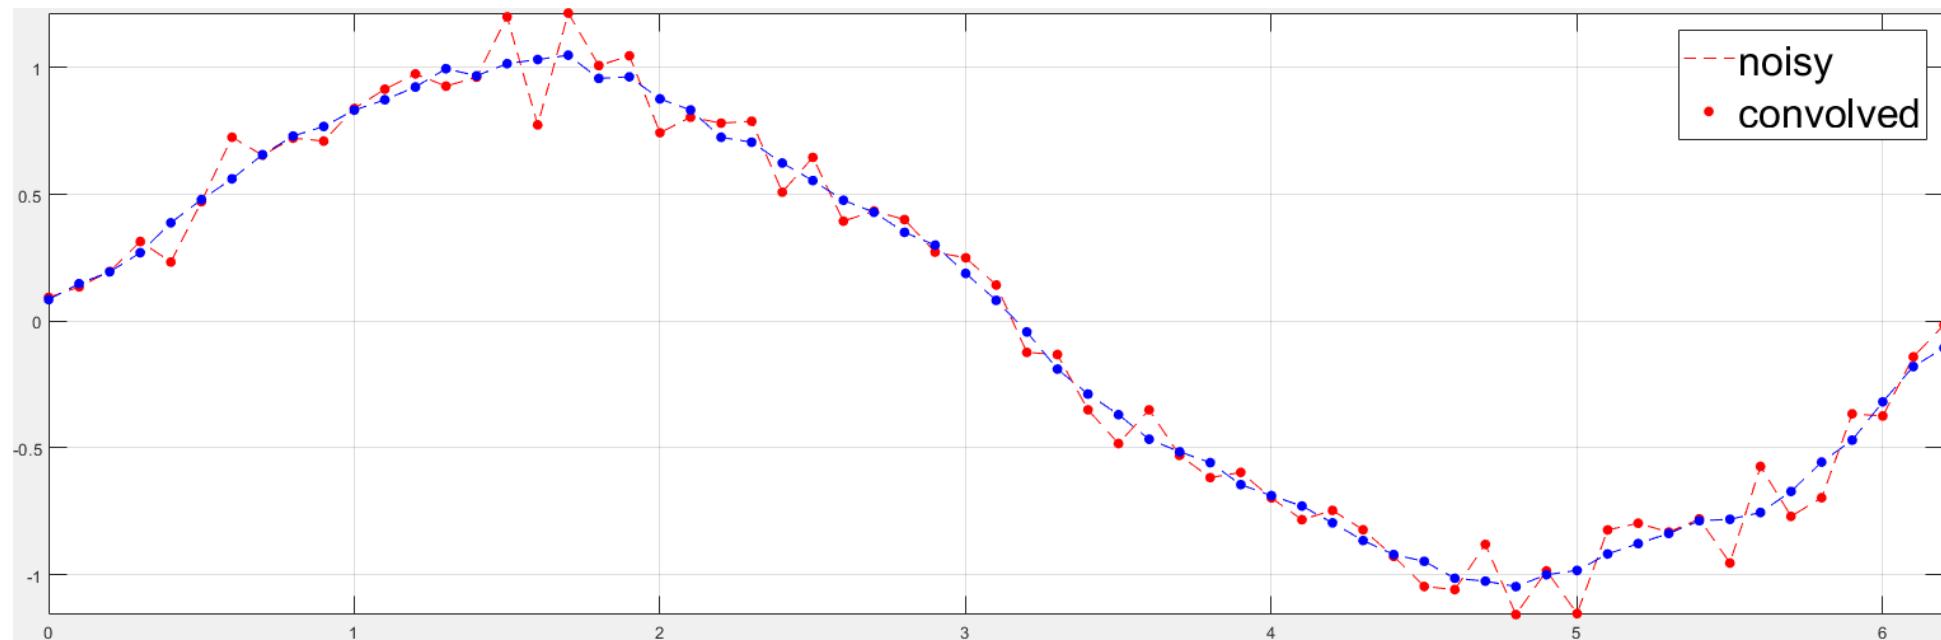
What about an impulse?

convolution when there is impulse
↓ filtered in 5 neighborhood,

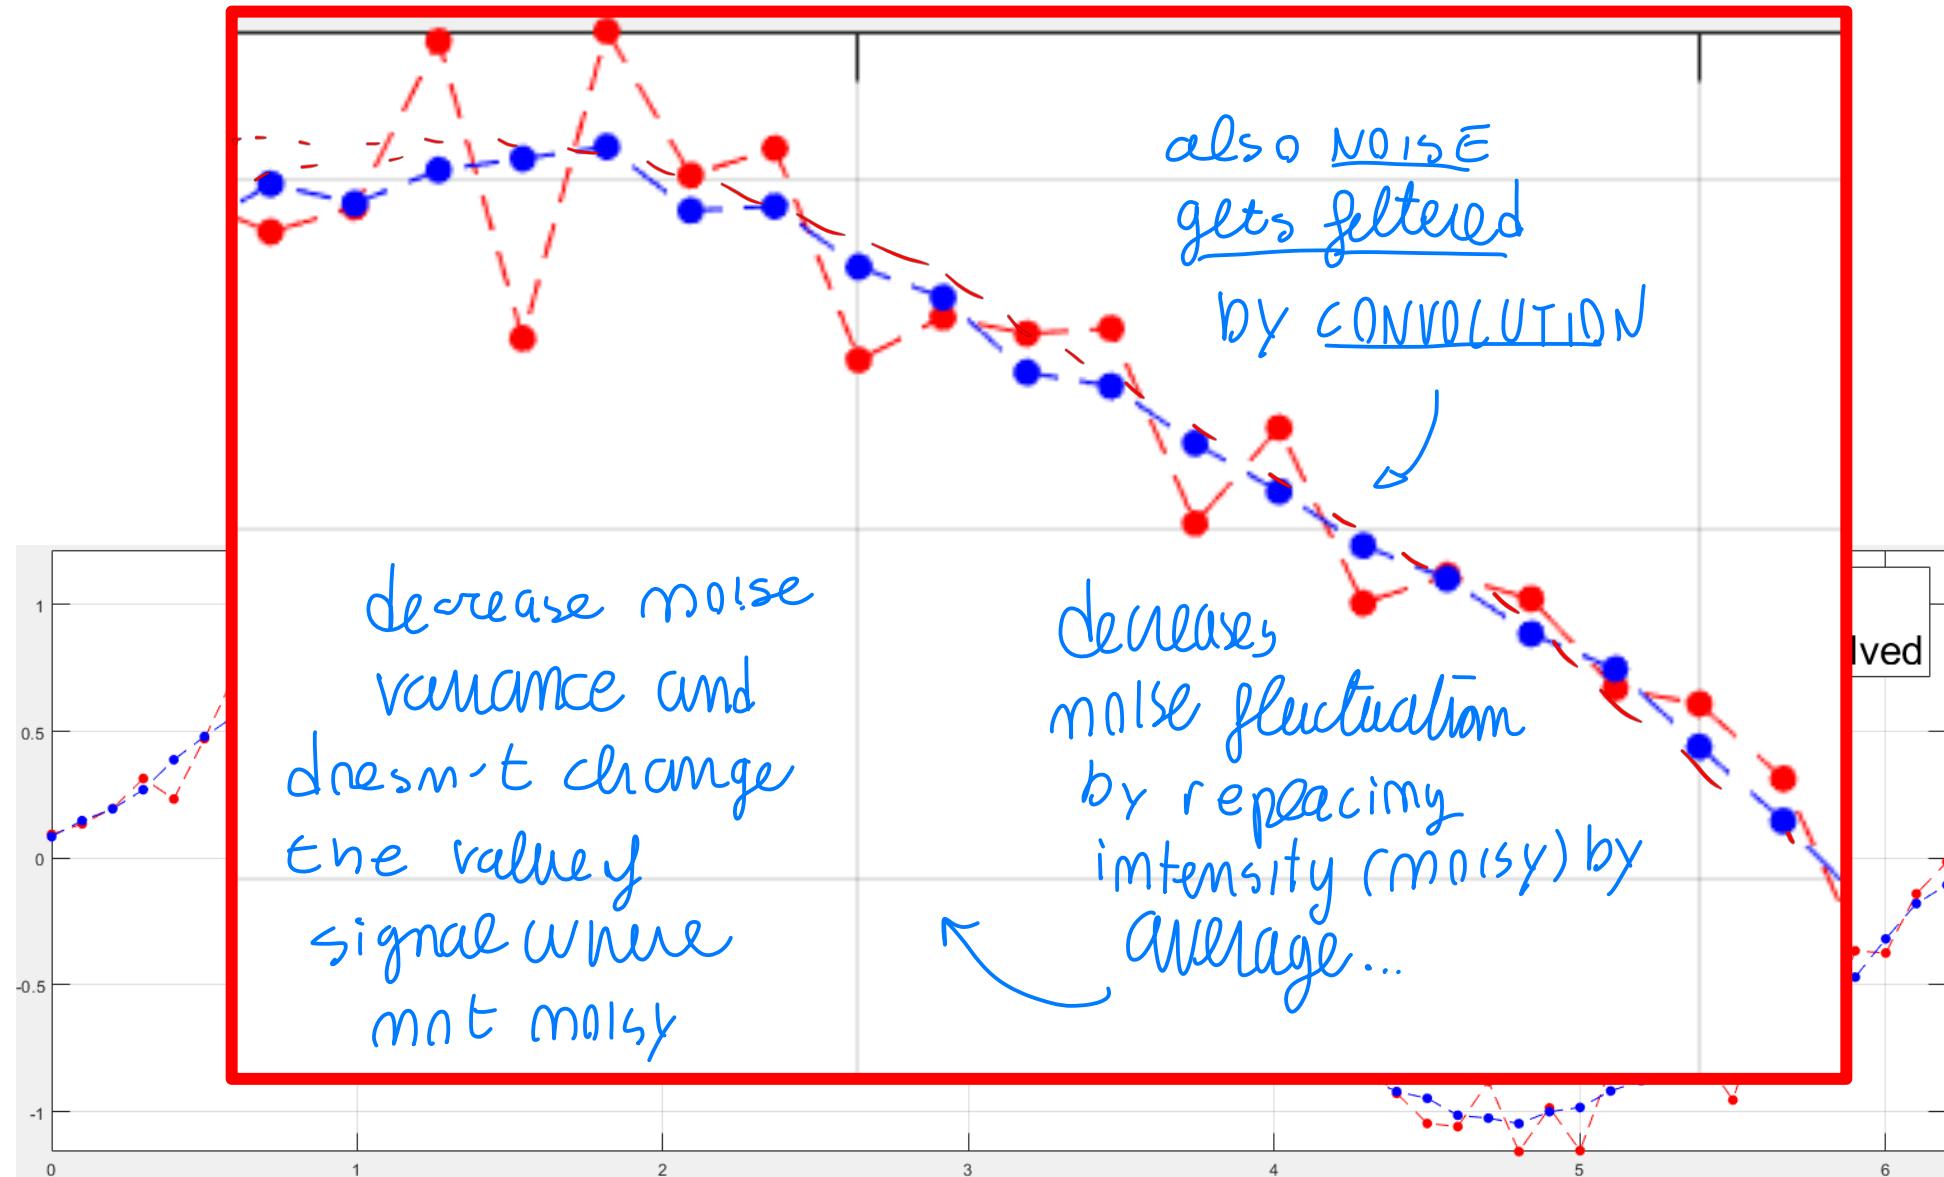
this peak
is filtered



What about noise?



What about noise?



Let's go back to
2D convolution now

A well-known Test Image - Lena



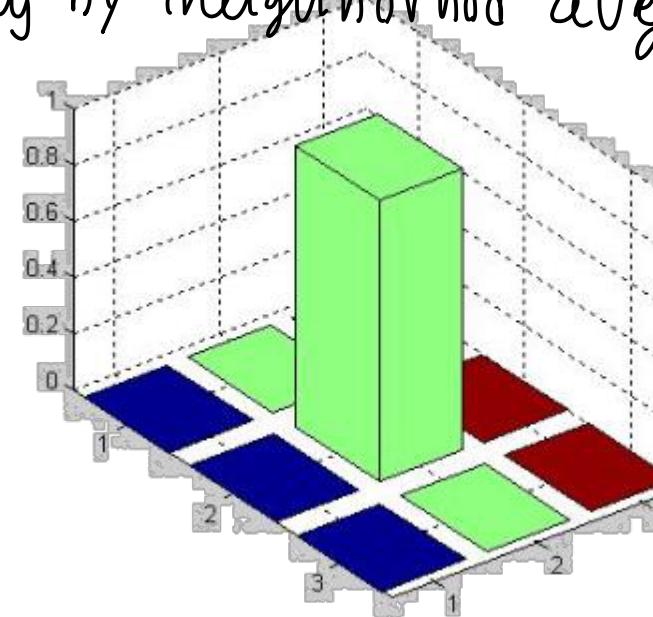
A Trivial example

replacing intensity by neighborhood avg accounted
With \mathbb{Q}



↑
filter of \mathbb{Q} around
1 in center is
the same image
as input I,

$$\begin{matrix} 0 & 0 & 0 \\ * & 0 & 1 & 0 \\ 0 & 0 & 0 \end{matrix} =$$

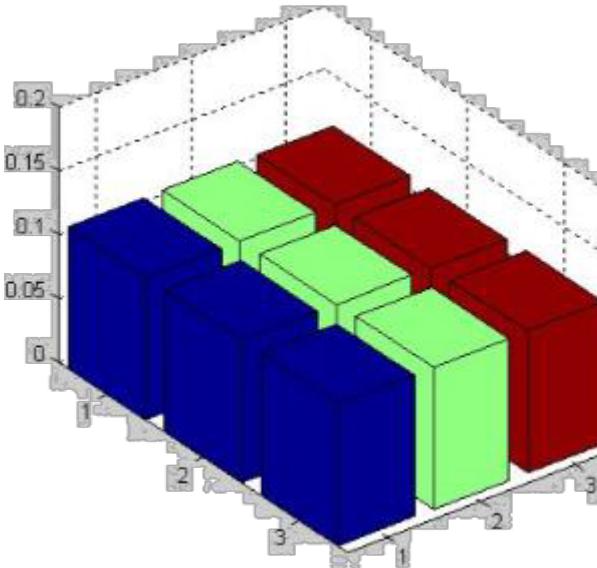


Linear Filtering



$$*\frac{1}{9}$$

1	1	1
1	1	1
1	1	1



= ?

*you average!
so you
lose details*

The original Lena image



Filtered Lena Image

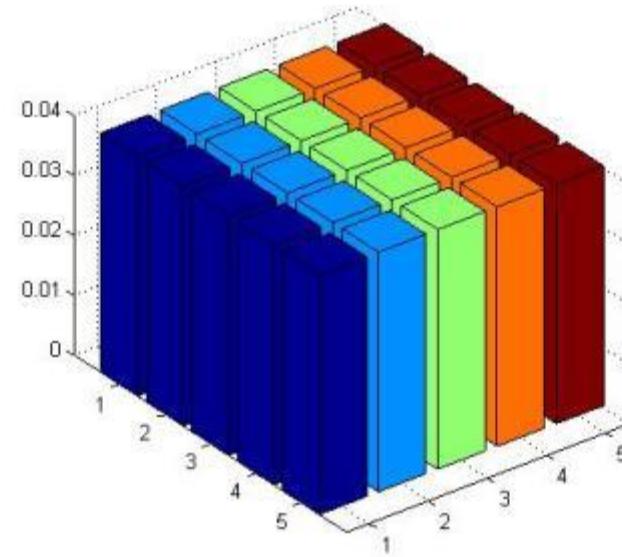




$$*\frac{1}{25}$$

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

=



The original Lena image



The filtered Lena image

the bigger the
filter, the
more details
you miss



What about normalization?

...what about



when NORMALIZATION
occurs on filtering

$$\otimes \frac{2}{25} \cdot \begin{array}{|c|c|c|c|c|} \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline \end{array} = \underbrace{\hspace{1cm}}_{5}$$

... convolution is linear

by
convolution
we blur image
and increase
intensity
(rescaling)



...what about

$$\frac{2}{25} \cdot$$



It is the same as rescaling before
and then
convolute

$$\otimes \begin{array}{|c|c|c|c|c|} \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline 1 & 1 & 1 & 1 & 1 \\ \hline \end{array} =$$

... convolution is linear

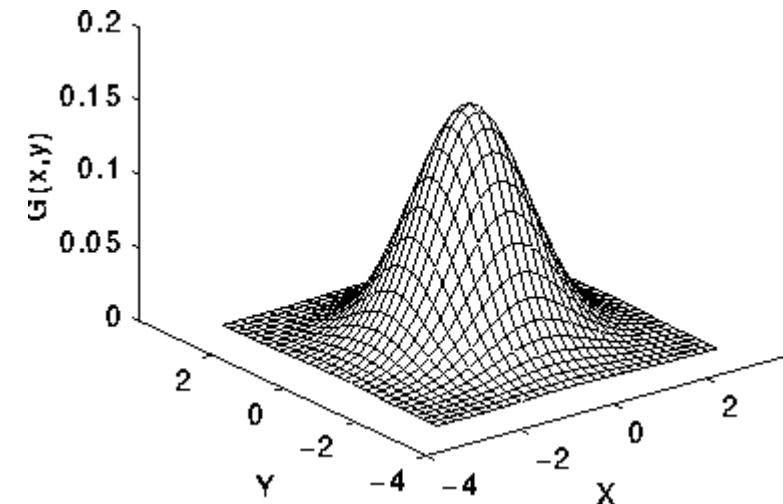


2D Gaussian Filter

design a filter in Gaussian Way

Continuous Function

$$H_{\sigma}(x, y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{(x^2 + y^2)}{2\sigma^2}\right)$$



Discrete kernel: assuming G is a $(2k + 1) \times (2k + 1)$ filter

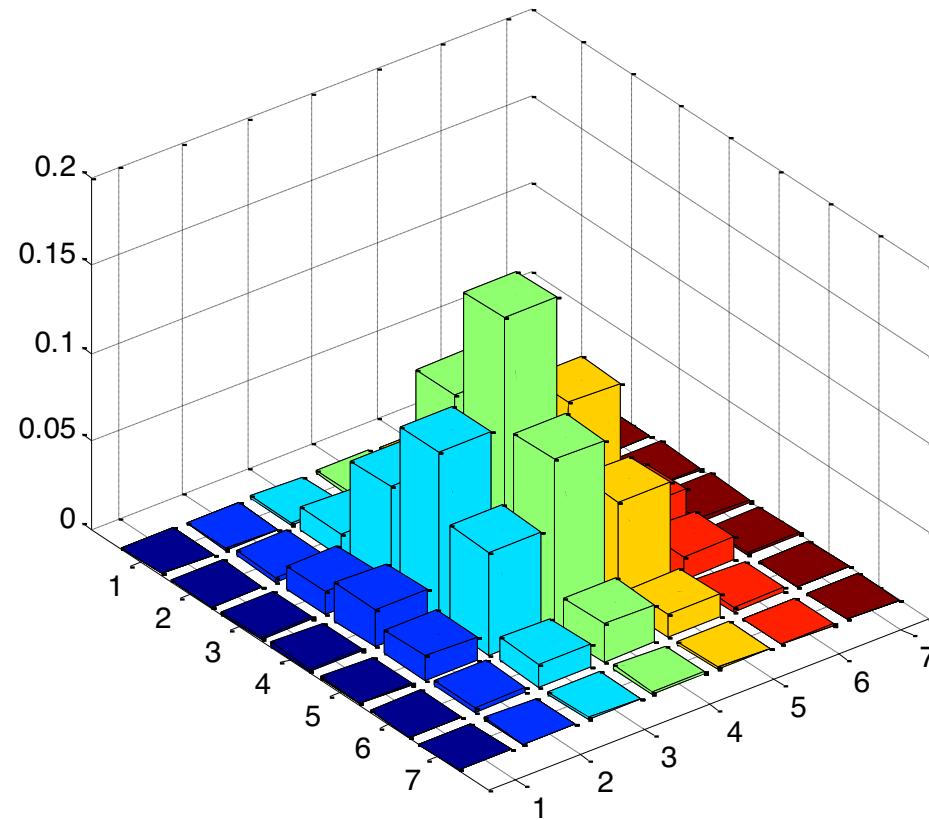
$$G(i, j) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{(i^2 + j^2)}{2\sigma^2}\right)$$

That is then normalized such that $\sum_{i=-k}^k \sum_{j=-k}^k G(i, j) = 1$

Weighted local averaging filters: Gaussian Filter



*



Weighted local averaging filters: Gaussian Filter



↔ blur
image,
less than
before

Gaussian Smoothing vs Averaging Filters



Gaussian Smoothing
Support 7x7



Smoothing by Averaging
On 7x7 window

Convolution Properties

Properties of Convolution: Linearity

It is a **linear operator**

(any linear combination of images convolved) *is equivalent to convolve each than sum*

$$((\lambda I_1 + \mu I_2) \odot w)(r, c) = \lambda(I_1 \odot w)(r, c) + \mu(I_2 \odot w)(r, c)$$

where $\lambda, \mu \in \mathbb{R}$

Obviously, when the filter is center-symmetric, convolution and correlation are equivalent

Properties of Convolution (and Padding)

It is commutative (in principle)

$$I_1 \circledast I_2 = I_2 \circledast I_1$$

However, in discrete signals it depends on the padding criteria In continuous domain
it holds as well as on periodic signals

to define im boundary, you
PAD the image to convolute
at any location, fully
cess a mystery point

it is possible to
define padding s.t.
your output stay
also outside of the
image \rightarrow as long as
overlap filter you
can define ... full padding
get layer output an
gets commutative convolution

Input image I

0	0	0	0	0	0	0	0	0	0	1	-1
0	1	0	2	1	0	0	-1	-1	0		
0	1	1	1	0	1	0	1	-1	-1		
0	1	2	1	0	1	0					
0	1	2	0	2	2	0					
0	1	0	1	0	0	0					
0	0	0	0	0	0	0					

filter w

Original image is in violet,
grey values are padded to
zero to enable convolution
at image boundaries

Is Convolution Commutative?

Gaussian
filtering

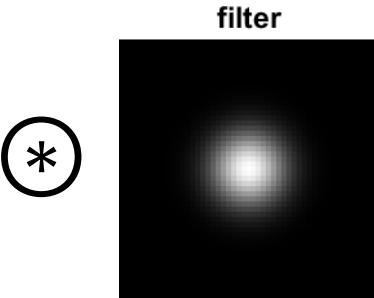


$$\text{input} \otimes \text{filter} = \text{output}$$

A diagram illustrating convolution. On the left is the input image of the woman. In the center is a square "filter" kernel labeled "filter" with a red arrow pointing to its top edge. To the left of the filter is a circled asterisk (*) symbol, indicating convolution. To the right of the filter is an equals sign (=). On the right is the output image, which is the input image convolved with the Gaussian filter, resulting in a blurred version of the woman's face.



Is Convolution Commutative?



=



boundary padding with \mathbb{Q}
outside.. may \mathbb{Q} in the border
which gets
darker

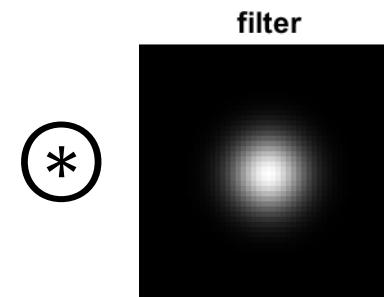
Is Convolution Commutative?

NO, because OUTPUT has same size of image input... (?)

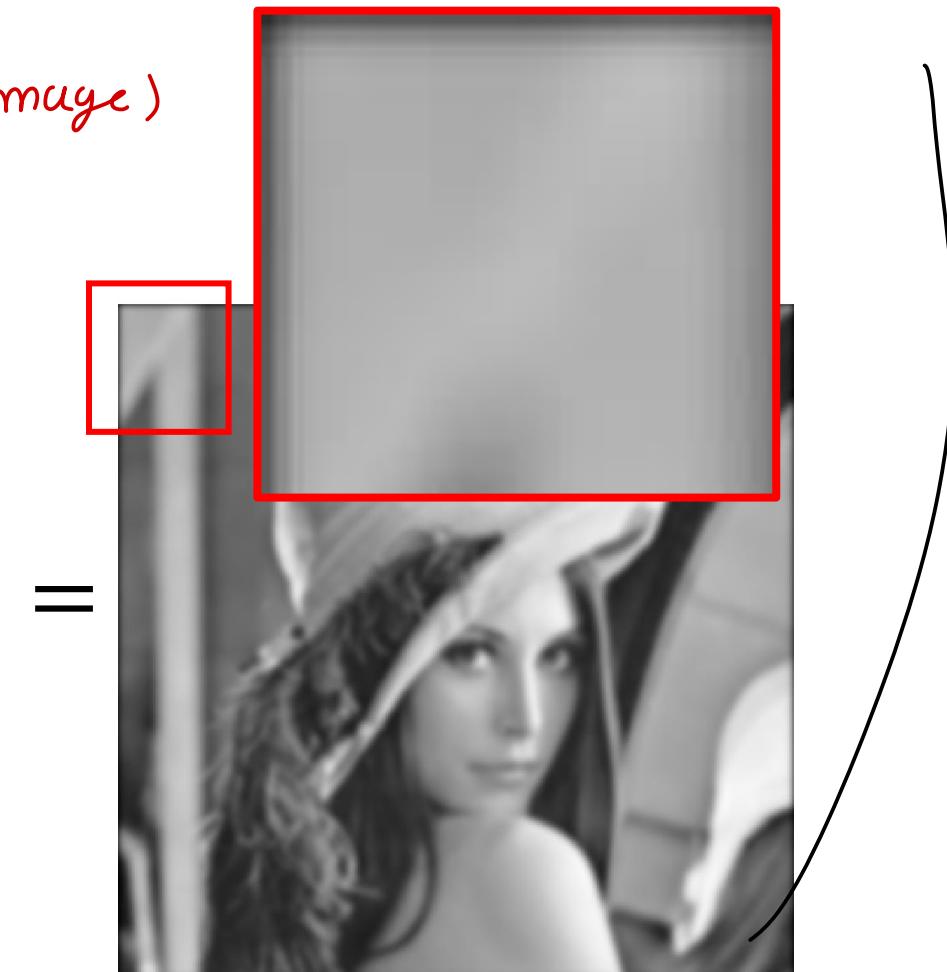
what if I flip the filter?

convolution of the filter... (NOT of image)

↓
don't get the same...



*



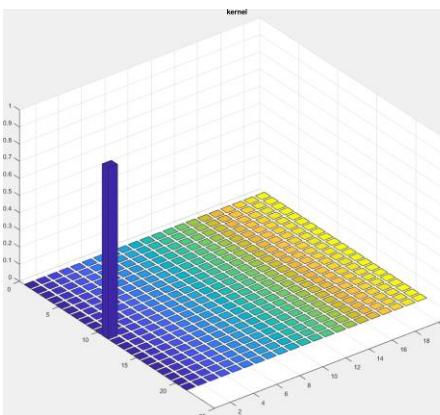
by using image as filter

commutativeness depends on how you define padding options

Translation

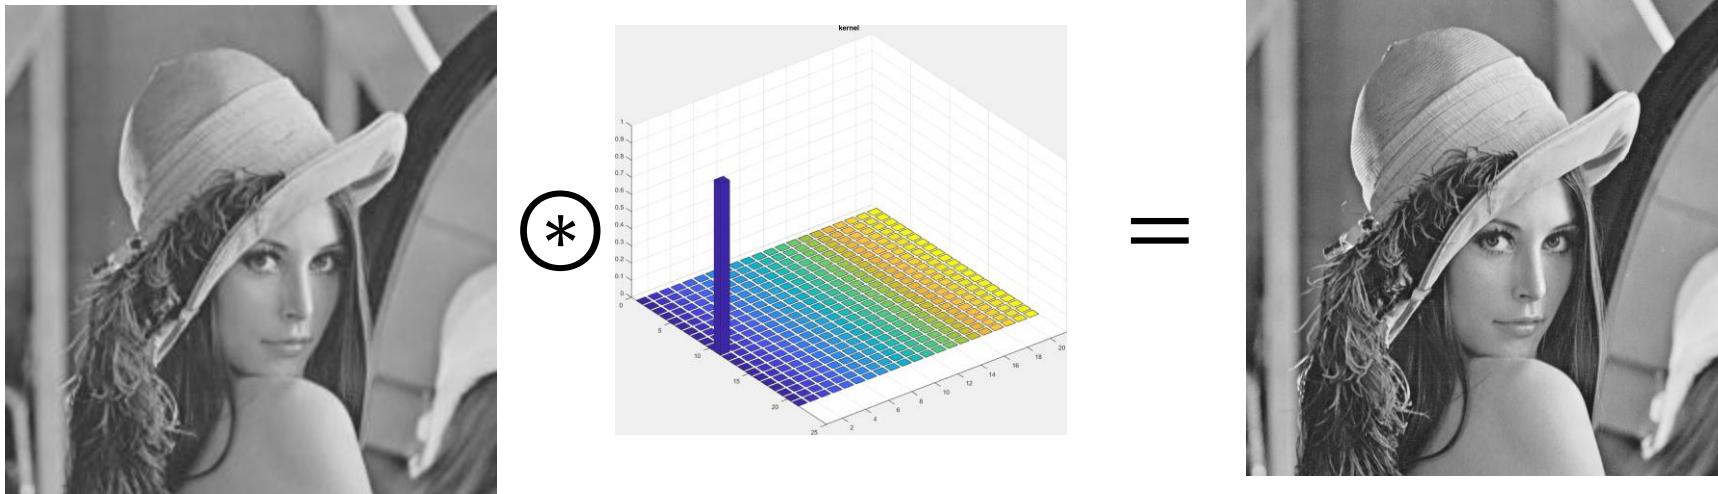


⊗



=

Translation



Remember the filter has to be flipped before convolution

Is Convolution Commutative?

by using image as filter



This holds for the «full convolution» modality, not the «same» or «valid»

Properties of Convolution: Associative

It is also **associative**

$$f \odot (g \odot w) = (f \odot g) \odot w = f \odot g \odot w$$

and **dissociative**

$$f \odot (g + w) = f \odot g + f \odot w$$

Properties of Convolution: Shift invariance

It is also **associative**

$$f \odot (g \odot w) = (f \odot g) \odot w = f \odot g \odot w$$

and **dissociative**

$$f \odot (g + w) = f \odot g + f \odot w$$

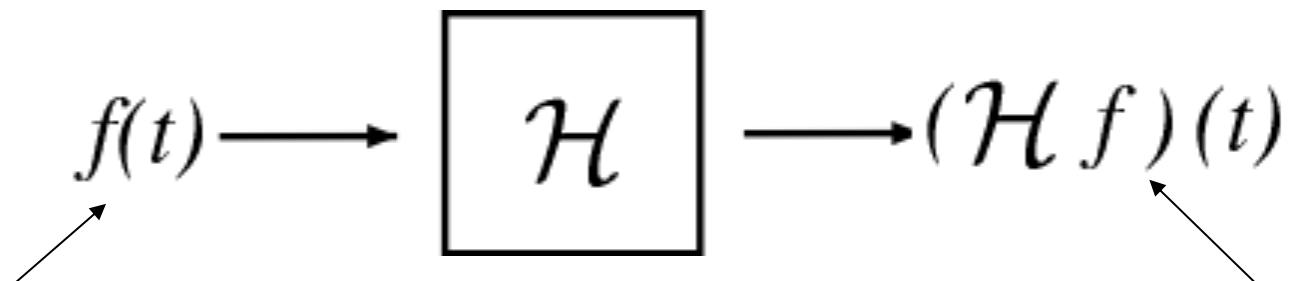
It is **shift-invariant**, namely

$$(I(\cdot - r_0, \cdot - c_0) \odot w)(r, c) = (I \odot w)(r - r_0, c - c_0)$$

Any linear and shift invariant system can be written as a convolution

Systems

Consider a system H as a black box that processes an input signal (f) and gives the output (i.e, $H[f]$)



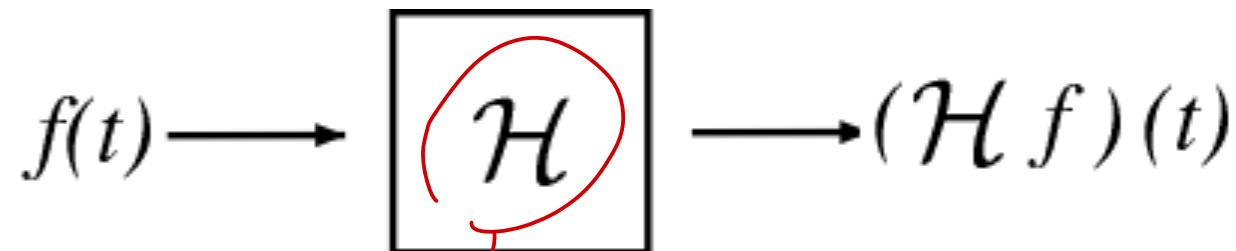
The input is a signal

The output is a signal

Systems

Consider a system H as a black box that processes an input signal (f) and gives the output (i.e., $H[f]$)

it is behaving like linear filter



just defined by filter coefficients

In our case, f is a digital image (a 2D matrix), but in principle could be any (analogic or digital) n-dimensional signal

Linearity and Time Invariance

A system is **linear** if and only if

$$H[\lambda f(t) + \mu g(t)] = \lambda H[f](t) + \mu H[g](t)$$

holds for any $\lambda, \mu \in \mathbb{R}$ and for f, g arbitrary signals (this is the canonical definition of linearity for an operator)

A system is **time (or shift) - invariant** if and only if

$$H[f(t - t_0)] = H[f](t - t_0)$$

holds for any $t_0 \in \mathbb{R}$ and for any signal f

Linear and Time Invariant Systems

All the systems that are Linear and Time Invariant (LTI) have an equivalent convolutional operator

- LTI systems are characterized entirely by a **single function**, the **filter**

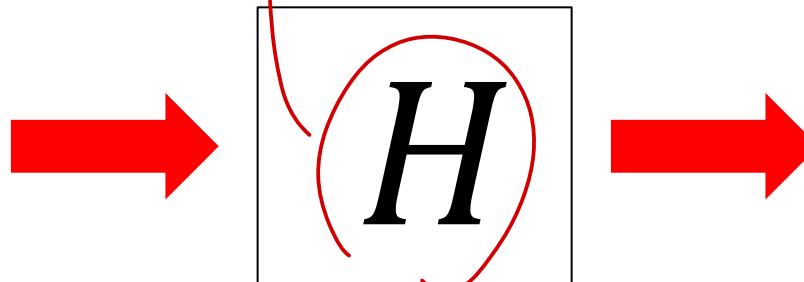
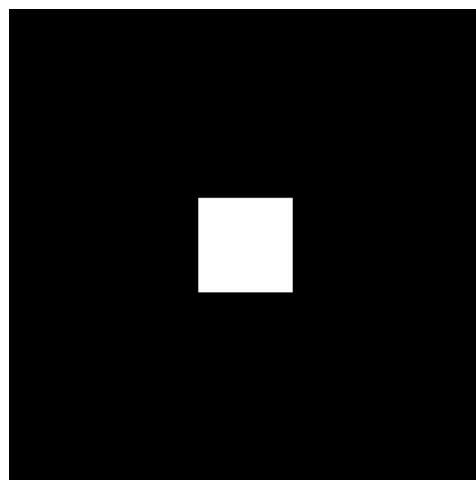
Linear and Time Invariant Systems

↓ every LINEAR and TRANSLATION invariant FILTER can be written as convolution!

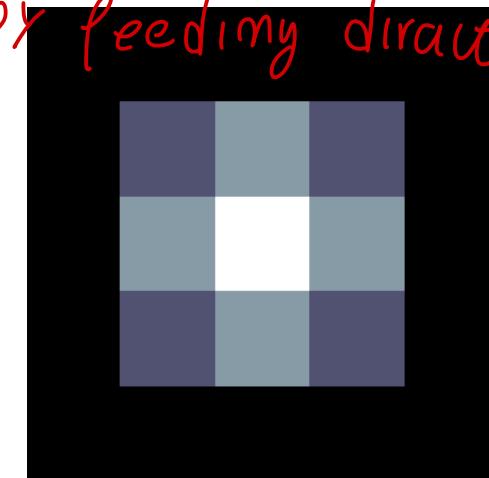
All the systems that are Linear and Time Invariant (LTI) have an equivalent convolutional operator

- LTI systems are **characterized** entirely by a single function, the filter
- The filter is also called system's the **impulse response** or **point spread function**, as it corresponds to the output of an impulse fed to the system

Uniquely defined by filter coefficients, you can completely understand it by feeding directly



this returns as output the image of filter itself

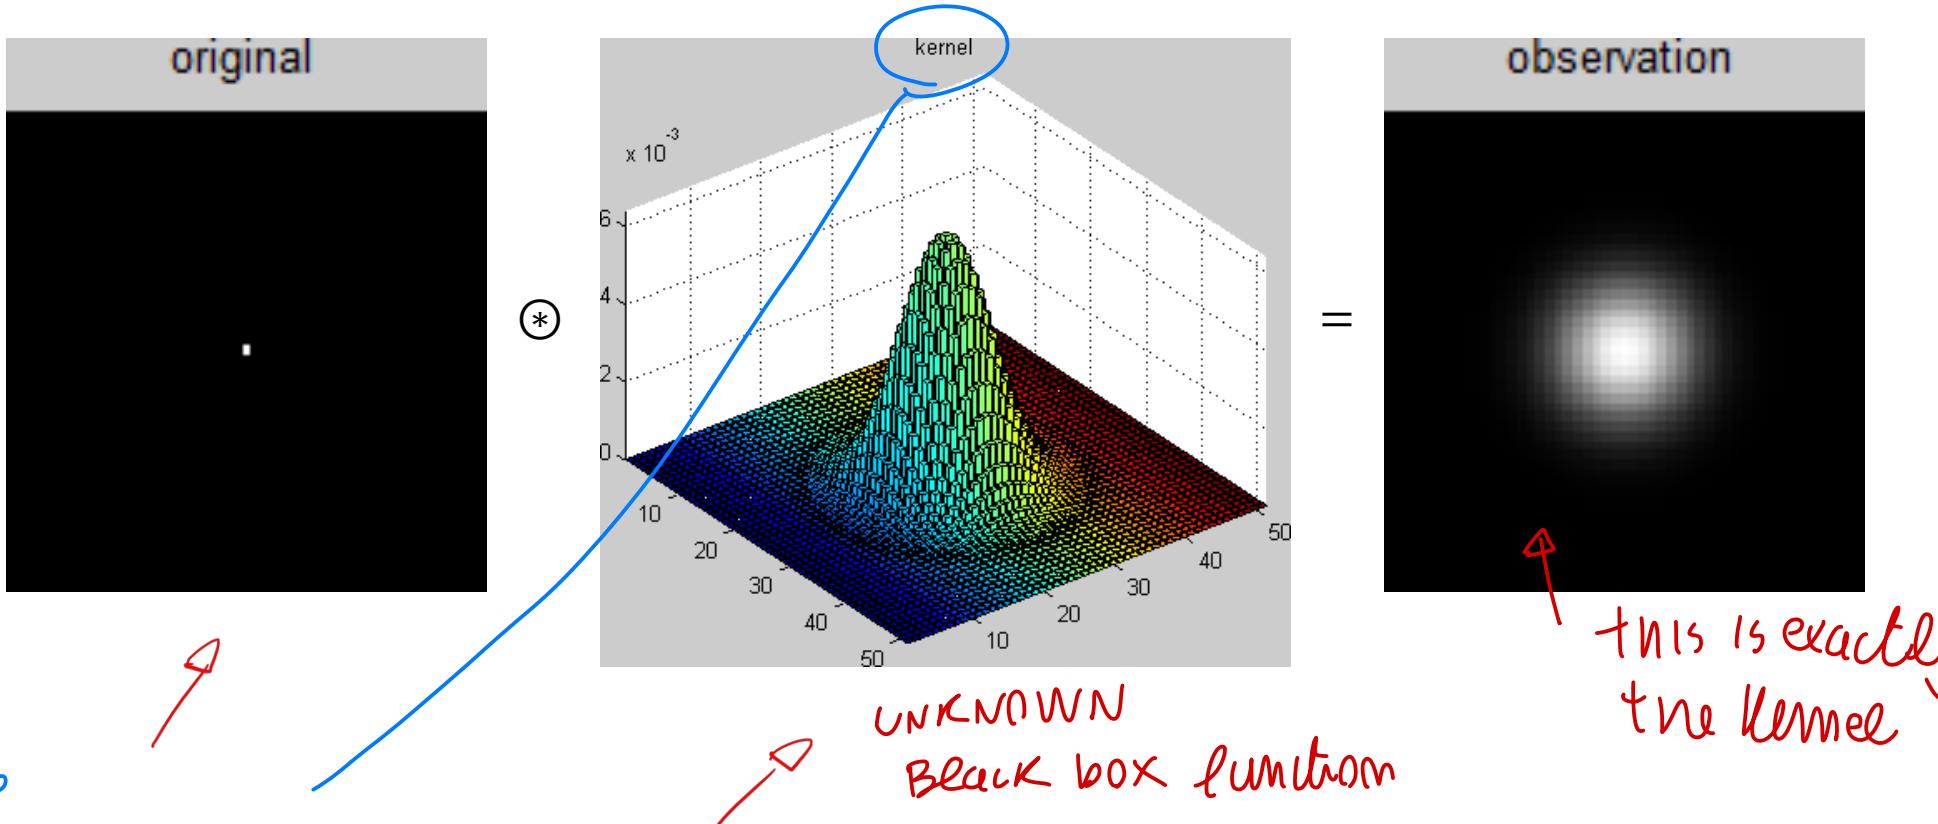


↳ every value
↳ on center

The Impulse Response

Take as input image a discrete Dirac

you can understand what H is doing by feeding in input the direct image to the filter



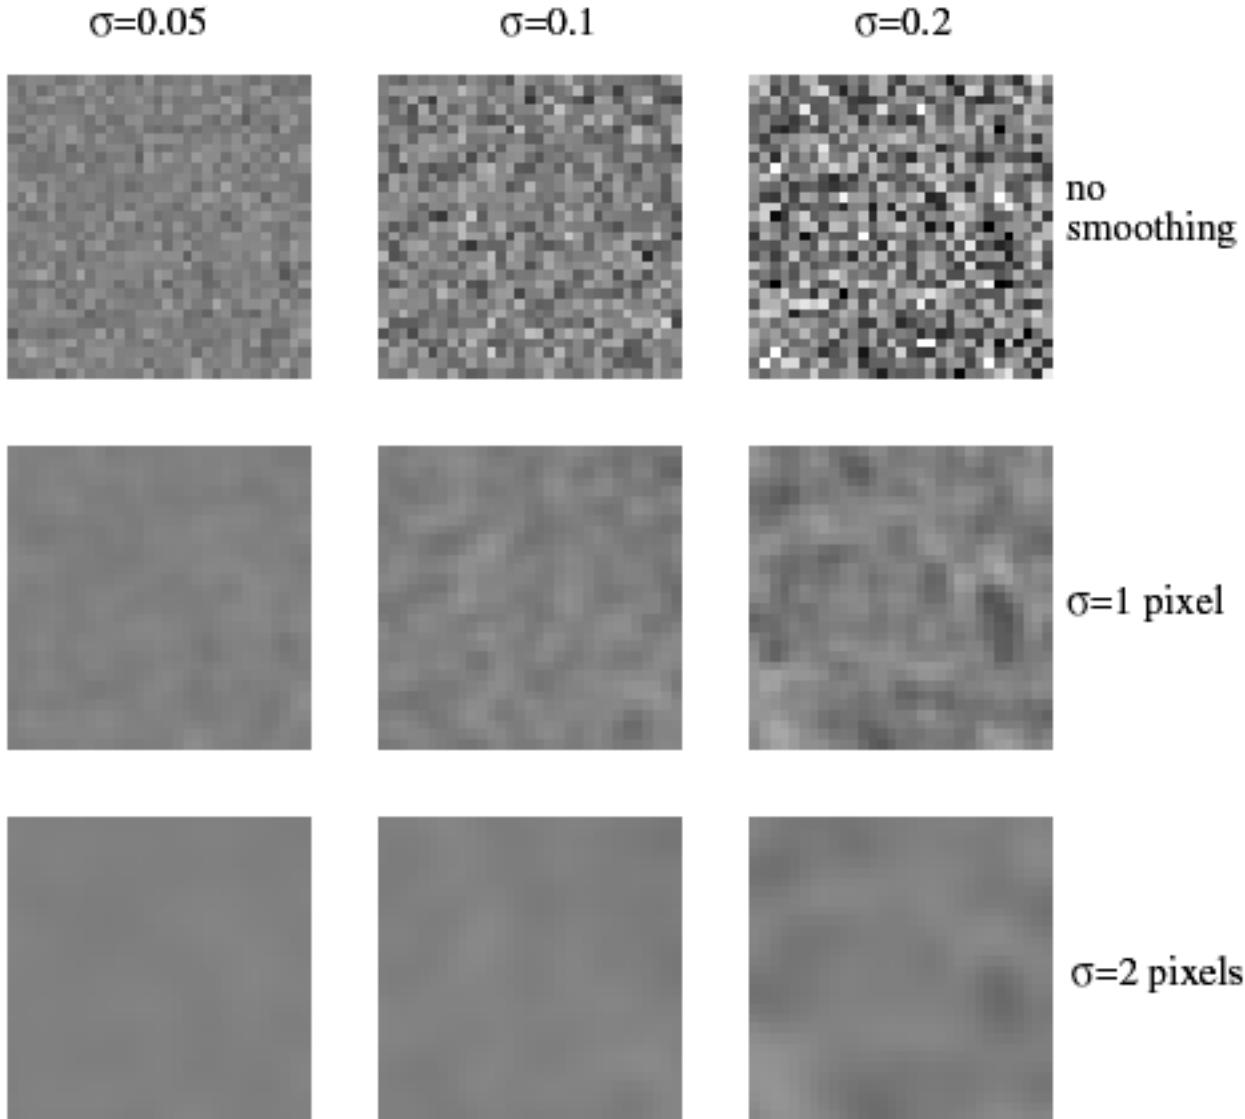
"Kernel" is another way to call the "filter"

This is why h is also called the "Point Spread Function"

Denoising

An application scenario for digital filters

Low - Pass



no
smoothing

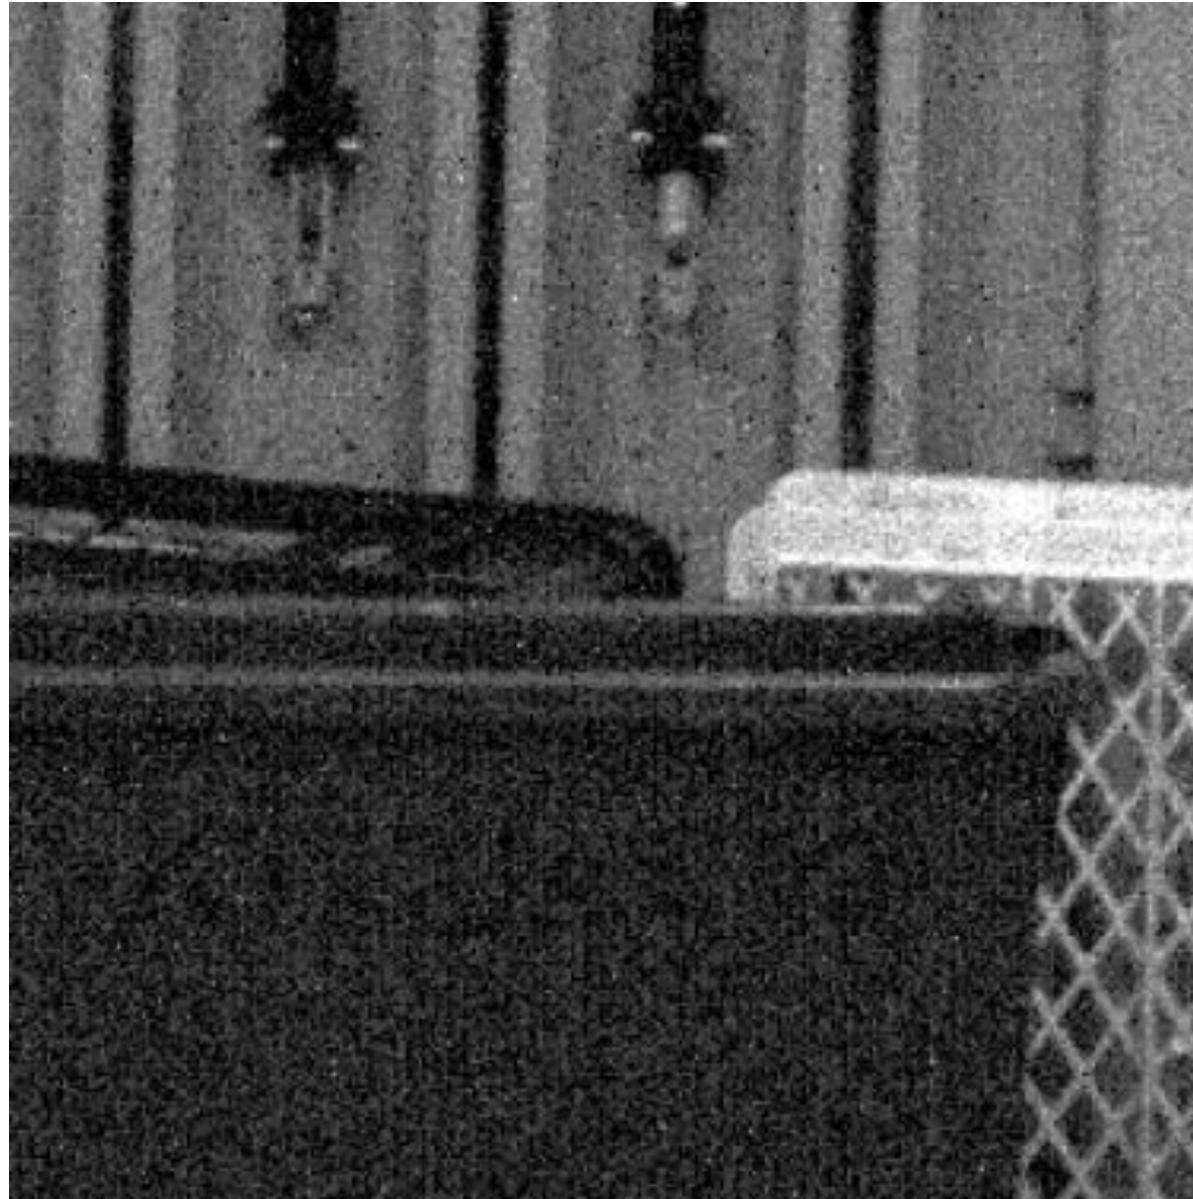
$\sigma=1 \text{ pixel}$

$\sigma=2 \text{ pixels}$

The effects of smoothing
Each row shows smoothing
with gaussians of different
width; each column shows
different realisations of
an image of gaussian noise.

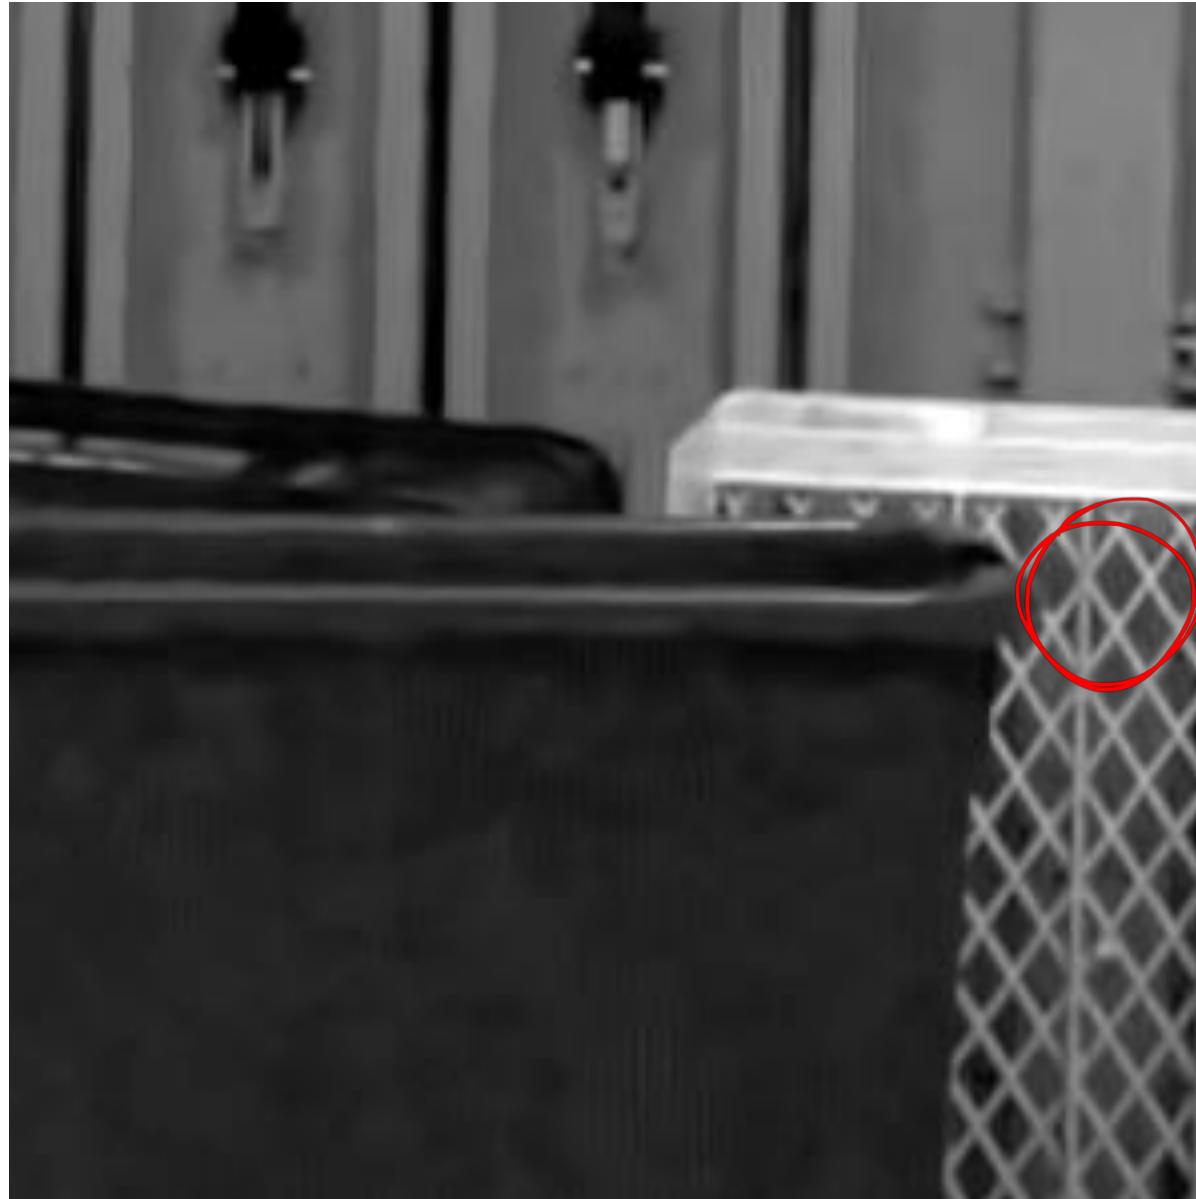
Denoising: The Issue

A Detail in
Camera Raw
Image



Denoising: The Issue

Denoised



Denoising: The Issue

A Detail in Camera
Raw Image



Denoising: The Issue

Denoised



Image Formation Model

Observation model is

$$z(x) = y(x) + \eta(x), \quad x \in \mathcal{X}$$

Where

- x denotes the pixel coordinates in the domain $\mathcal{X} \subset \mathbb{Z}^2$
- y is the original (noise-free and unknown) image
- z is the noisy observation
- η is the noise realization

Image Formation Model

Observation model is

$$z(x) = y(x) + \eta(x), \quad x \in \mathcal{X}$$

The goal is to compute \hat{y} *realistic* estimate of y , given z and the distribution of η .

For the sake of simplicity we assume AWG: $\eta \sim N(0, \sigma^2)$ and $\eta(x)$ independent realizations.

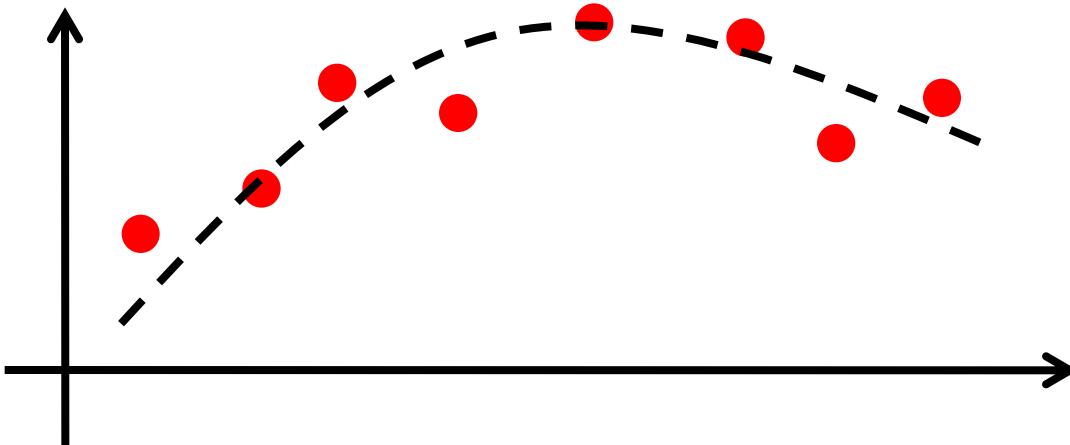
The noise standard deviation σ is also assumed as known.

Convolution and Regression

Observation model is

$$z(x) = y(x) + \eta(x) \quad x \in X$$

Consider a regression problem



Fitting and Convolution

The convolution provides the BLUE (Best Linear Unbiased Estimator) for regression when the image y is constant

The problem: estimating the constant C that minimizes a weighted loss over noisy observations

$$\widehat{y}_h(x_0) = \operatorname{argmin}_C \sum_{x_s \in X} w_h(x_0 - x_s) (z(x_s) - C)^2$$

Where

$$w_h = \{w_h(x)\} \quad s.t. \quad \sum_{x \in X} w_h(x) = 1$$

This problem can be solved by **computing the convolution** of the image z against a **filter whose coefficients are the error weights**

$$\widehat{y}(x_0) = (z \circledast w_h)(x_0)$$

Image Formation Model

Observation model is

$$z(x) = y(x) + \eta(x) \quad x \in X$$

Thus we can pursue a “regression-approach”, but on images it may not be convenient to assume a **parametric expression** of y on X

$z =$



Image Formation Model

Observation model is

$$z(x) = y(x) + \eta(x) \quad x \in X$$

Thus we can pursue a “regression-approach”, but on images it may not be convenient to assume a **parametric expression** of y on X

$z =$



$y =$



Local Smoothing



Additive Gaussian
White Noise

$$\eta \approx N(\mu, \sigma)$$



After Averaging



After Gaussian Smoothing

Denoising Approaches

Parametric Approaches

- Transform Domain Filtering, they assume the noisy-free signal is somehow sparse in a suitable domain (e.g Fourier, DCT, Wavelet) or w.r.t. some dictionary based decomposition)

Denoising Approaches

Parametric Approaches

- Transform Domain Filtering, they assume the noisy-free signal is somehow sparse in a suitable domain (e.g Fourier, DCT, Wavelet) or w.r.t. some dictionary based decomposition)

Non Parametric Approaches

- Local Smoothing / Local Approximation
- Non Local Methods

Denoising Approaches

Parametric Approaches

- Transform Domain Filtering, they assume the noisy-free signal is somehow sparse in a suitable domain (e.g Fourier, DCT, Wavelet) or w.r.t. some dictionary based decomposition)

Non Parametric Approaches

- Local Smoothing / Local Approximation
- Non Local Methods

Estimating $y(x)$ from $z(x)$ can be statistically treated as regression of z given x

$$\hat{y}(x) = E[z | x]$$

Denoising Approaches

Parametric Approaches

- Transform Domain Filtering, they assume the noisy-free signal is somehow sparse in a suitable domain (e.g Fourier, DCT, Wavelet) or w.r.t. some dictionary based decomposition)

Non Parametric Approaches

- Local Smoothing / Local Approximation
- Non Local Methods

Estimating $y(x)$ from $z(x)$ can be statistically treated as regression of z given x

$$\hat{y}(x) = E[z | x]$$

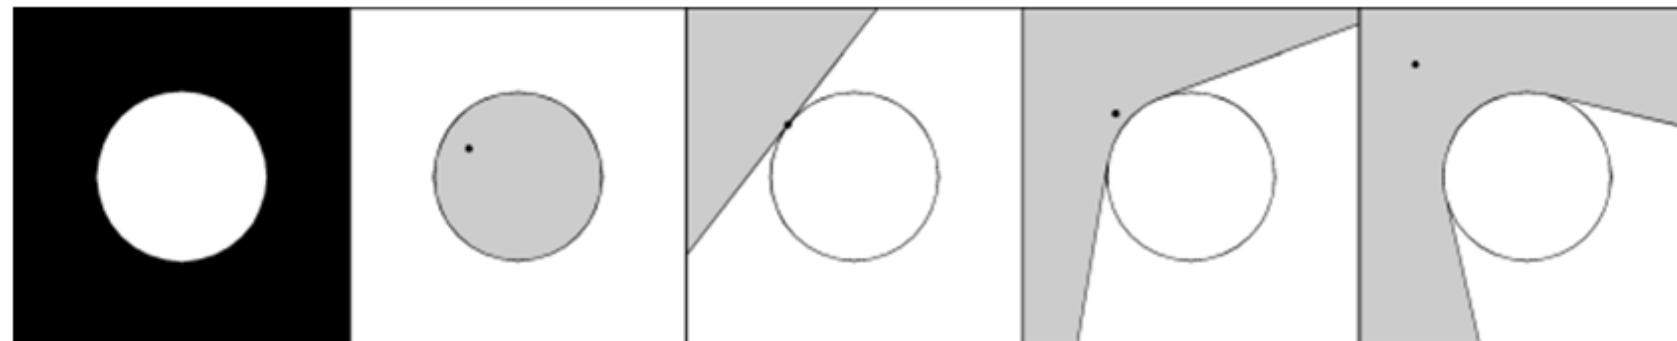
Denoising Approaches

Spatially adaptive methods, The basic principle:

- there are no simple models able to describe the whole image y , thus perform the regression $\hat{y}(x) = E[z | x]$
- Adopt a simple model in small image regions. For instance
$$\forall x \in X, \quad \exists \tilde{U}_x \text{ s.t. } y|_{\tilde{U}_x} \text{ is a polynomial}$$
- Define, in each image pixel, the “**best neighborhood**” where a simple parametric model can be enforced to perform regression.
- For instance, assume that on a suitable pixel-dependent neighborhood, where the image can be described by a polynomial

Ideal neighborhood – an illustrative example

Ideal in the sense that it defines the support of a pointwise Least Square Estimator of the reference point.



Typically, even in simple images, every point has its own different ideal neighborhood.

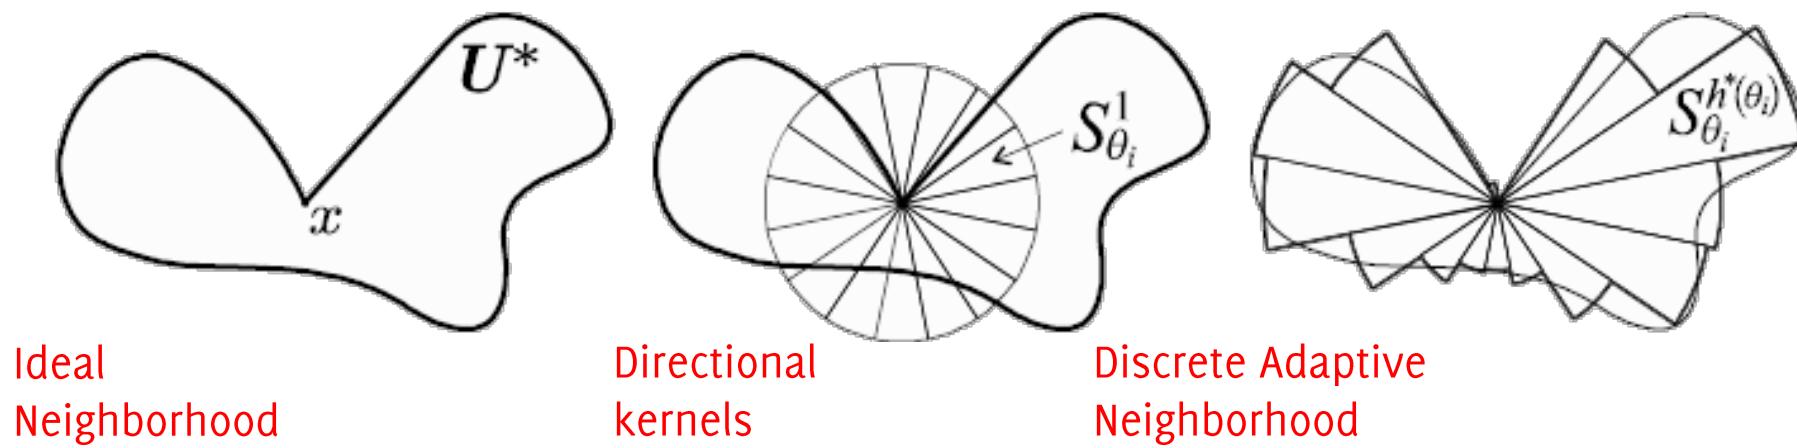
For practical reasons, the ideal neighborhood is assumed starshaped

Further details at LASIP c/o Tampere University of Technology

<http://www.cs.tut.fi/~lasip/>

Neighborhood discretization

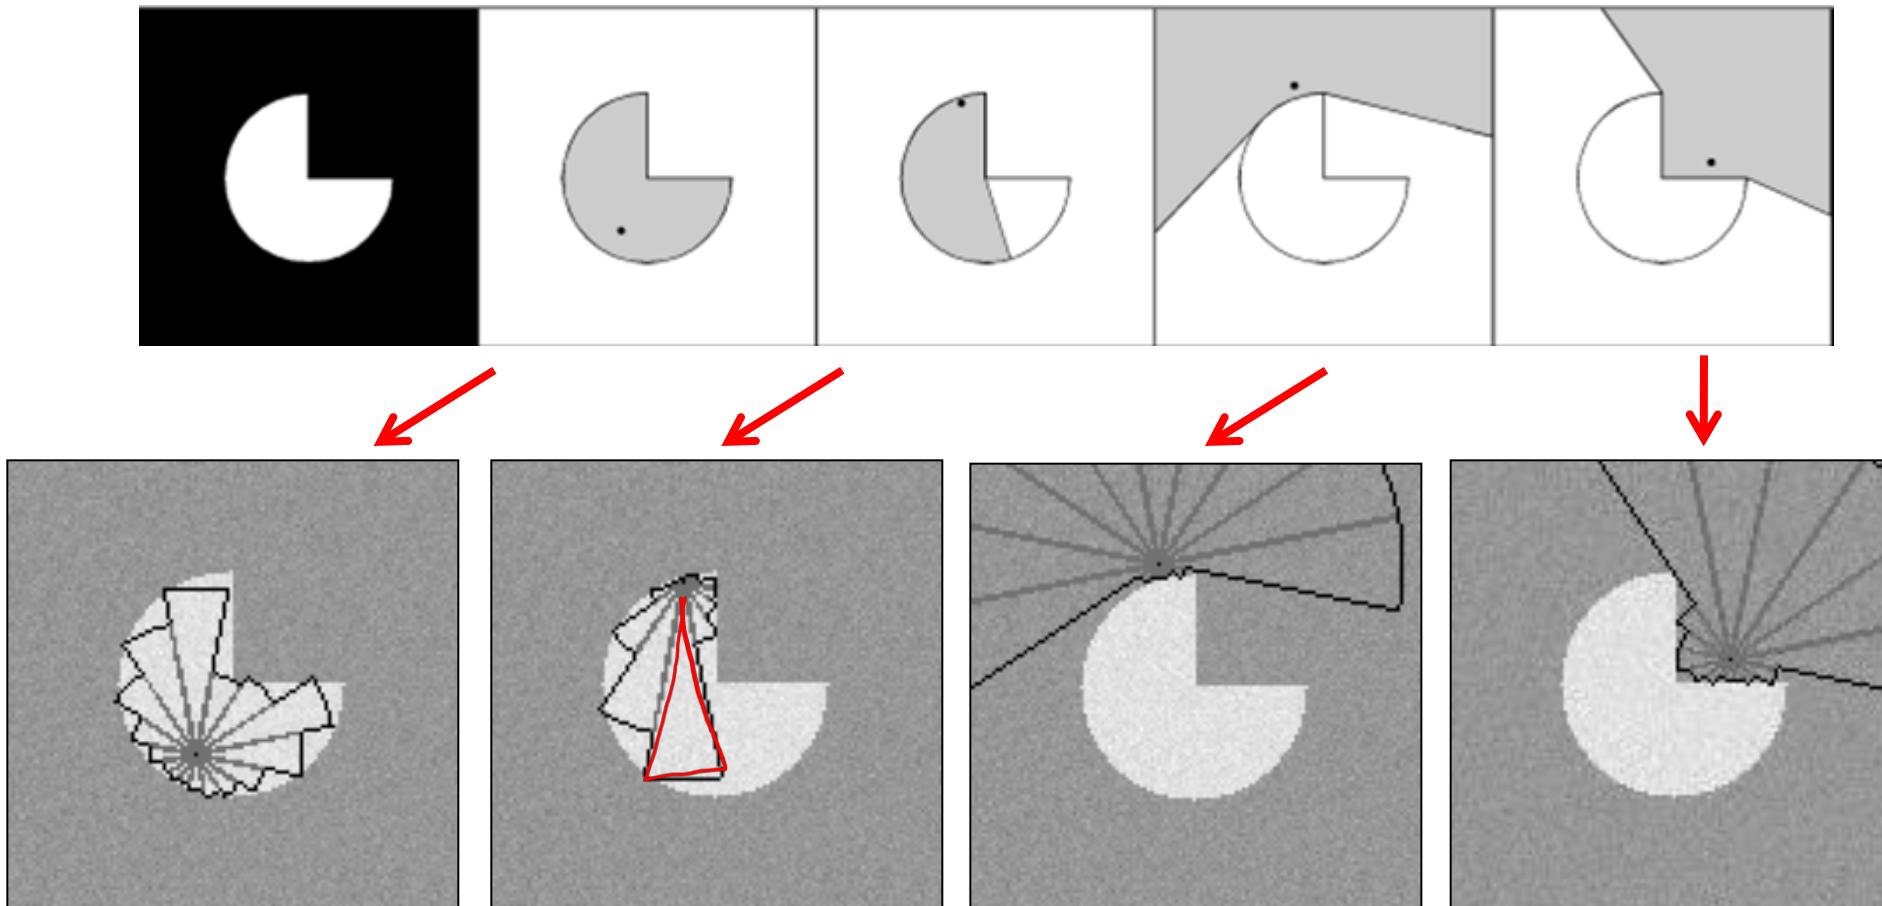
A suitable discretization of this neighborhood is obtained by using a set of directional LPA kernels $\{g_{\theta,h}\}_{\theta,h}$



where θ determines the orientation of the kernel support, and h controls the scale of kernel support.

Ideal neighborhood – an illustrative example

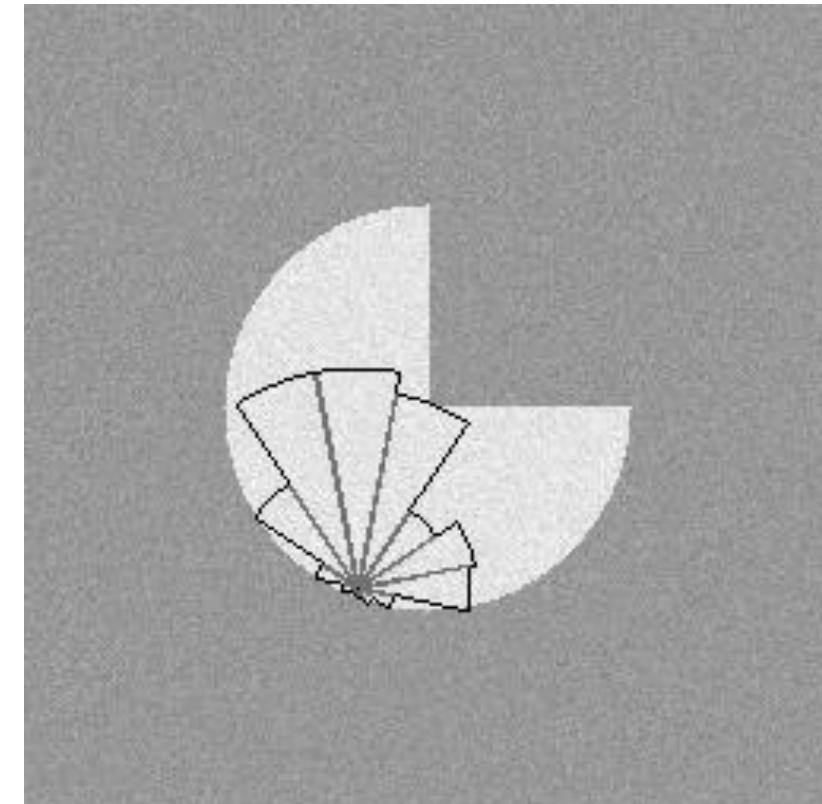
Ideal in the sense that the neighborhood defines the support of pointwise Least Square Estimator of the reference point.



Examples of Adaptively Selected Neighborhoods

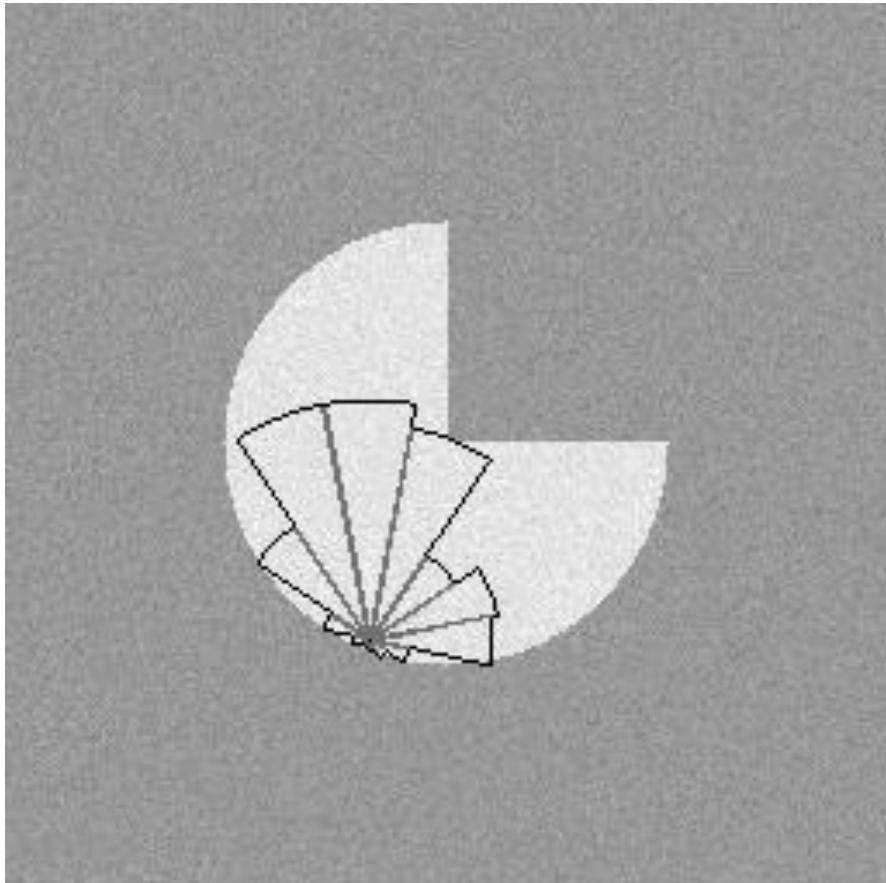
Define, $\forall x \in X$, the “ideal” neighborhood \tilde{U}_x

Compute the denoised estimate at x by “using” only pixels in \tilde{U}_x and a polynomial model to perform regression $\hat{y}(x) = E[z | x, \tilde{U}_x]$



Examples of adaptively selected neighborhoods

Neighborhoods adaptively selected using the LPA-ICI rule



Example of Performance

Original, noisy, denoised using polynomial regression on adaptively defined neighborhoods (LPA-ICI)

original



noisy



denoised

